



Running and Securing Kubernetes Apps



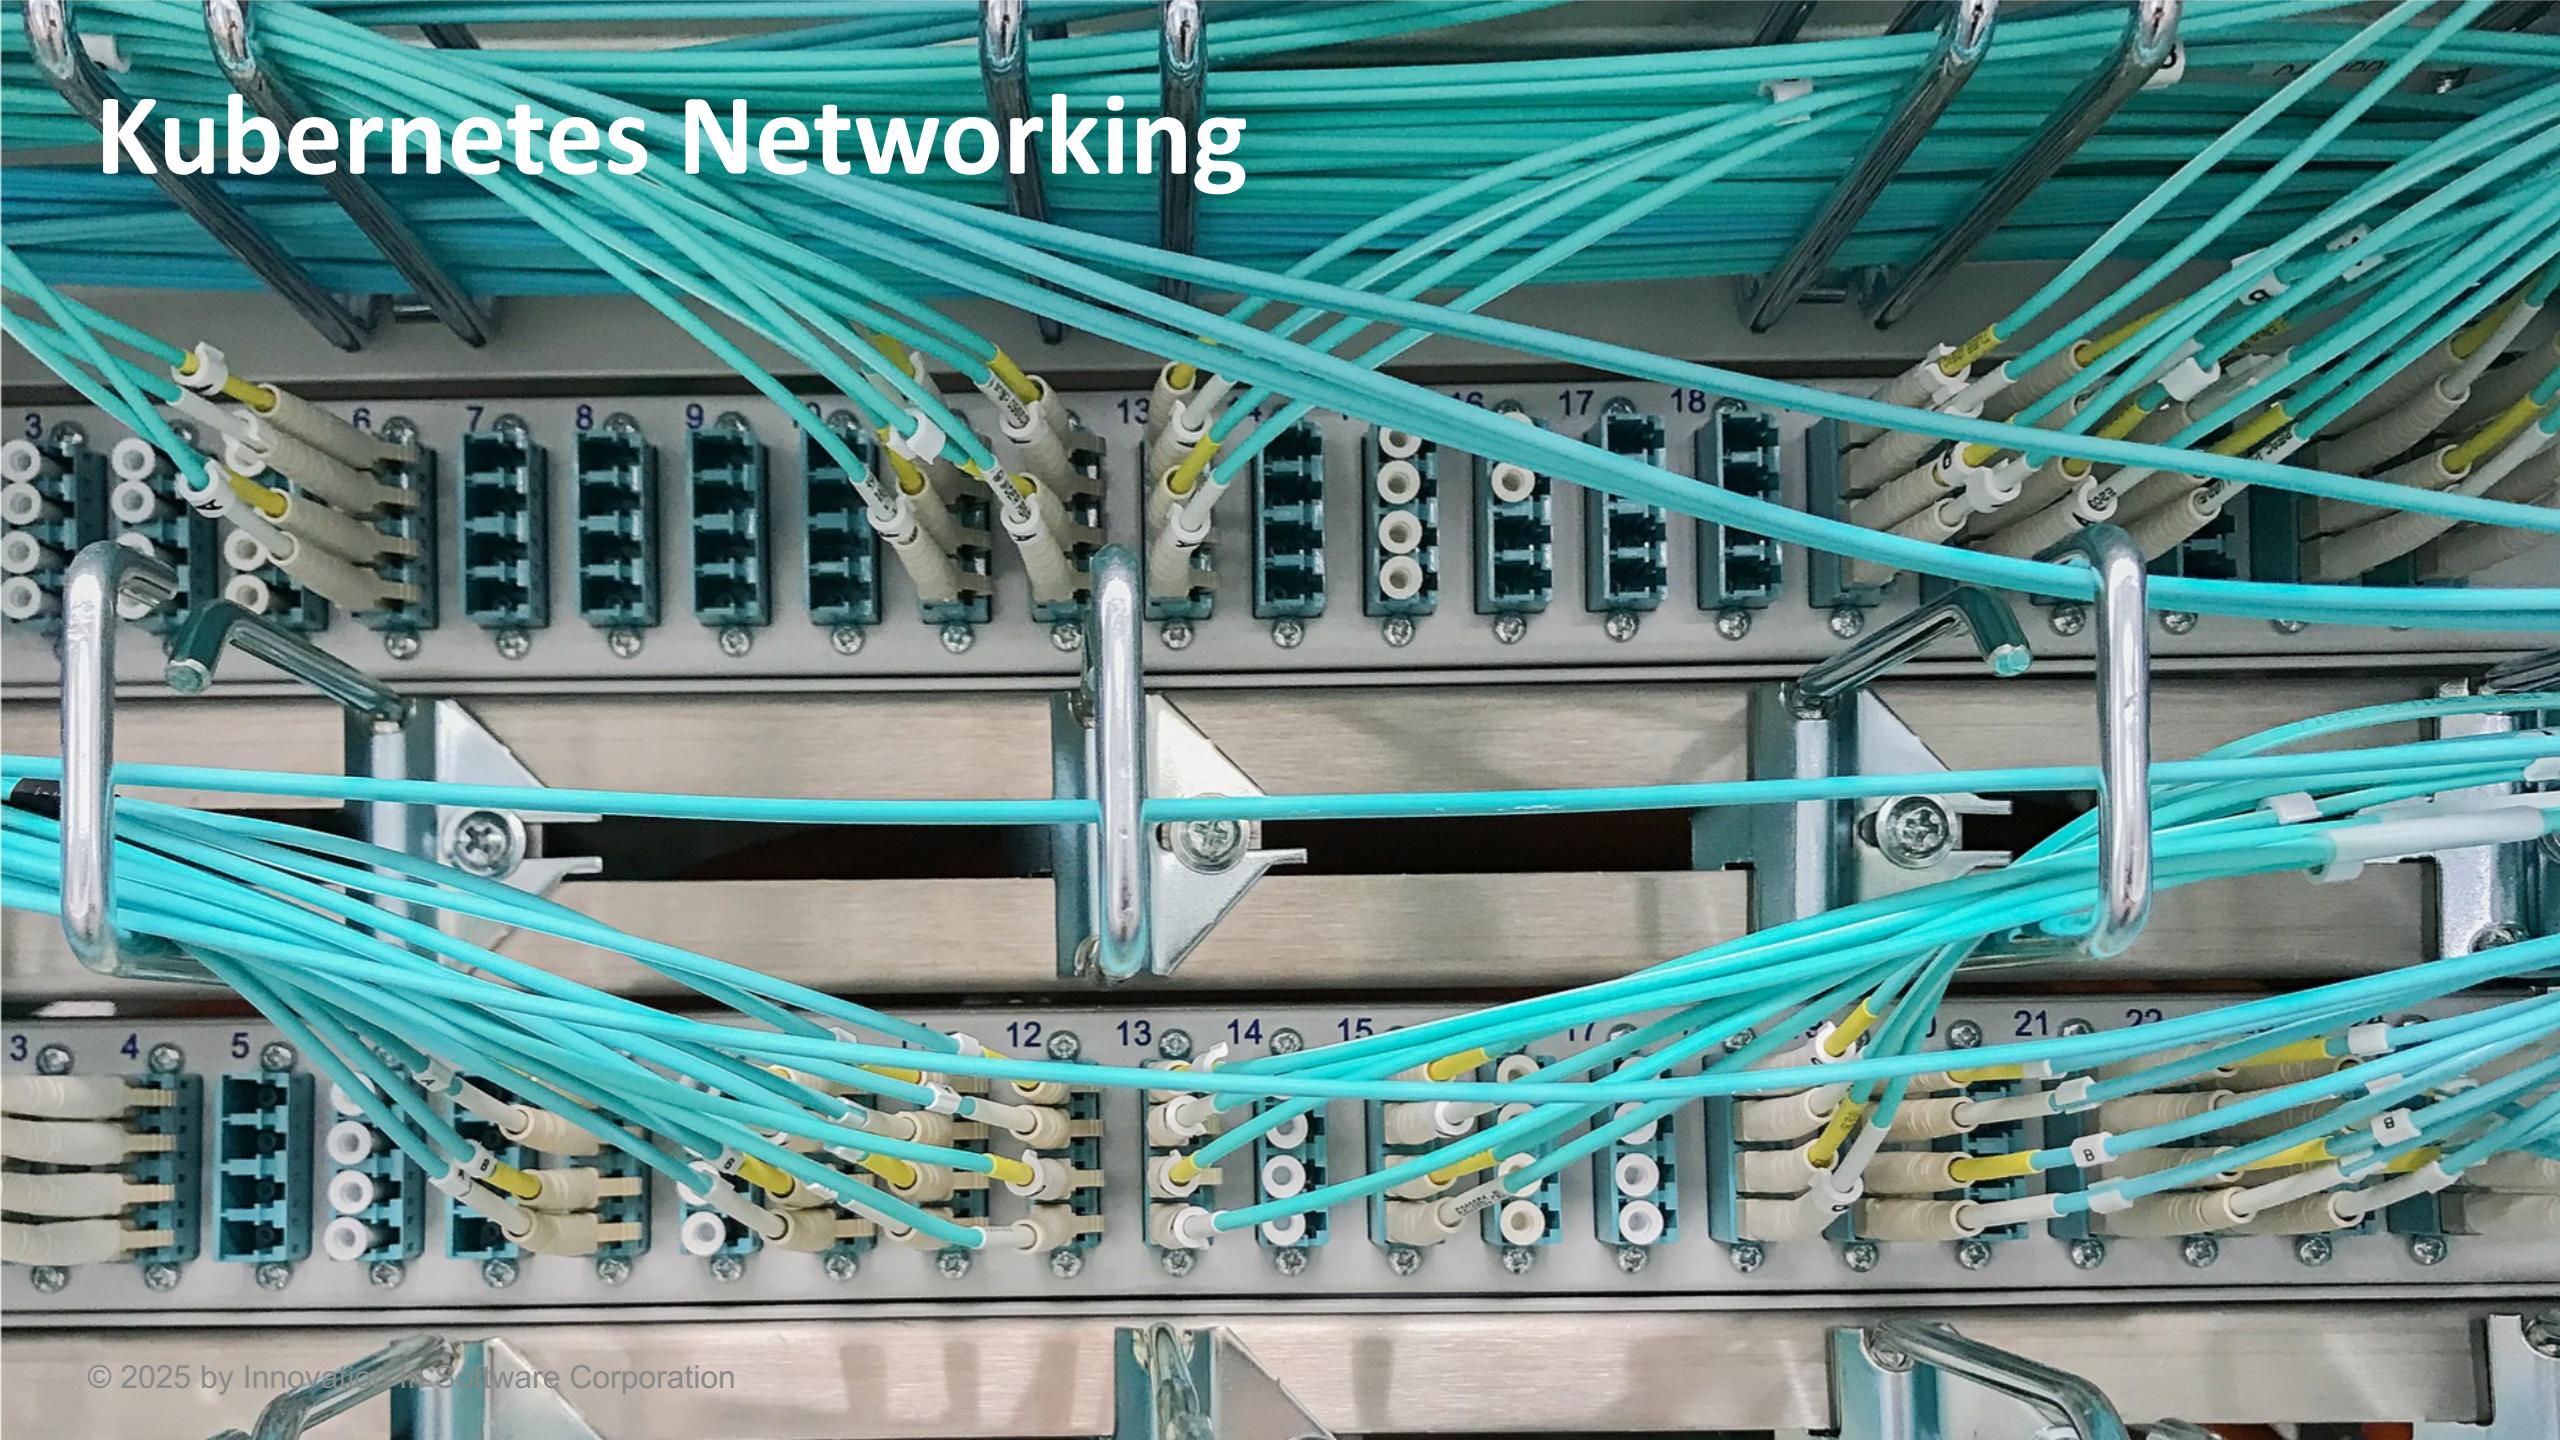
WORKFORCE DEVELOPMENT



Networking



Kubernetes Networking



Networking Problems and Solutions

1. Highly-coupled container-to-container communications => linux
2. Pod-to-Pod communications => various
3. Pod-to-Service communications => service
4. External-to-Service communications => service or ingress

Kubernetes Networking Model Requirements

- All containers can communicate with all other containers without NAT
- All nodes can communicate with all containers (and vice versa) without NAT
- The IP that a container sees itself as is the same IP that others see it as

Container to Container Communication

- Containers live inside of pods
- Group of one or more containers that are always co-located and co-scheduled, and run in a shared context
- Containers within a POD share an IP address and port space, and can find each other via localhost (net namespace).
- Communicate with each other using standard inter-process communications (ipc namespace).

Pod to Pod Requirements

- Pods can communicate with all other pods without NAT
- Regardless of which host they land on
- Every pod gets its own IP address
- No need to explicitly create links between pods

Pod Networking Approaches

Using an overlay network

- An overlay network obscures the underlying network architecture from the pod network through traffic encapsulation (for example vxlan).
- Encapsulation reduces performance, though exactly how much depends on your solution.

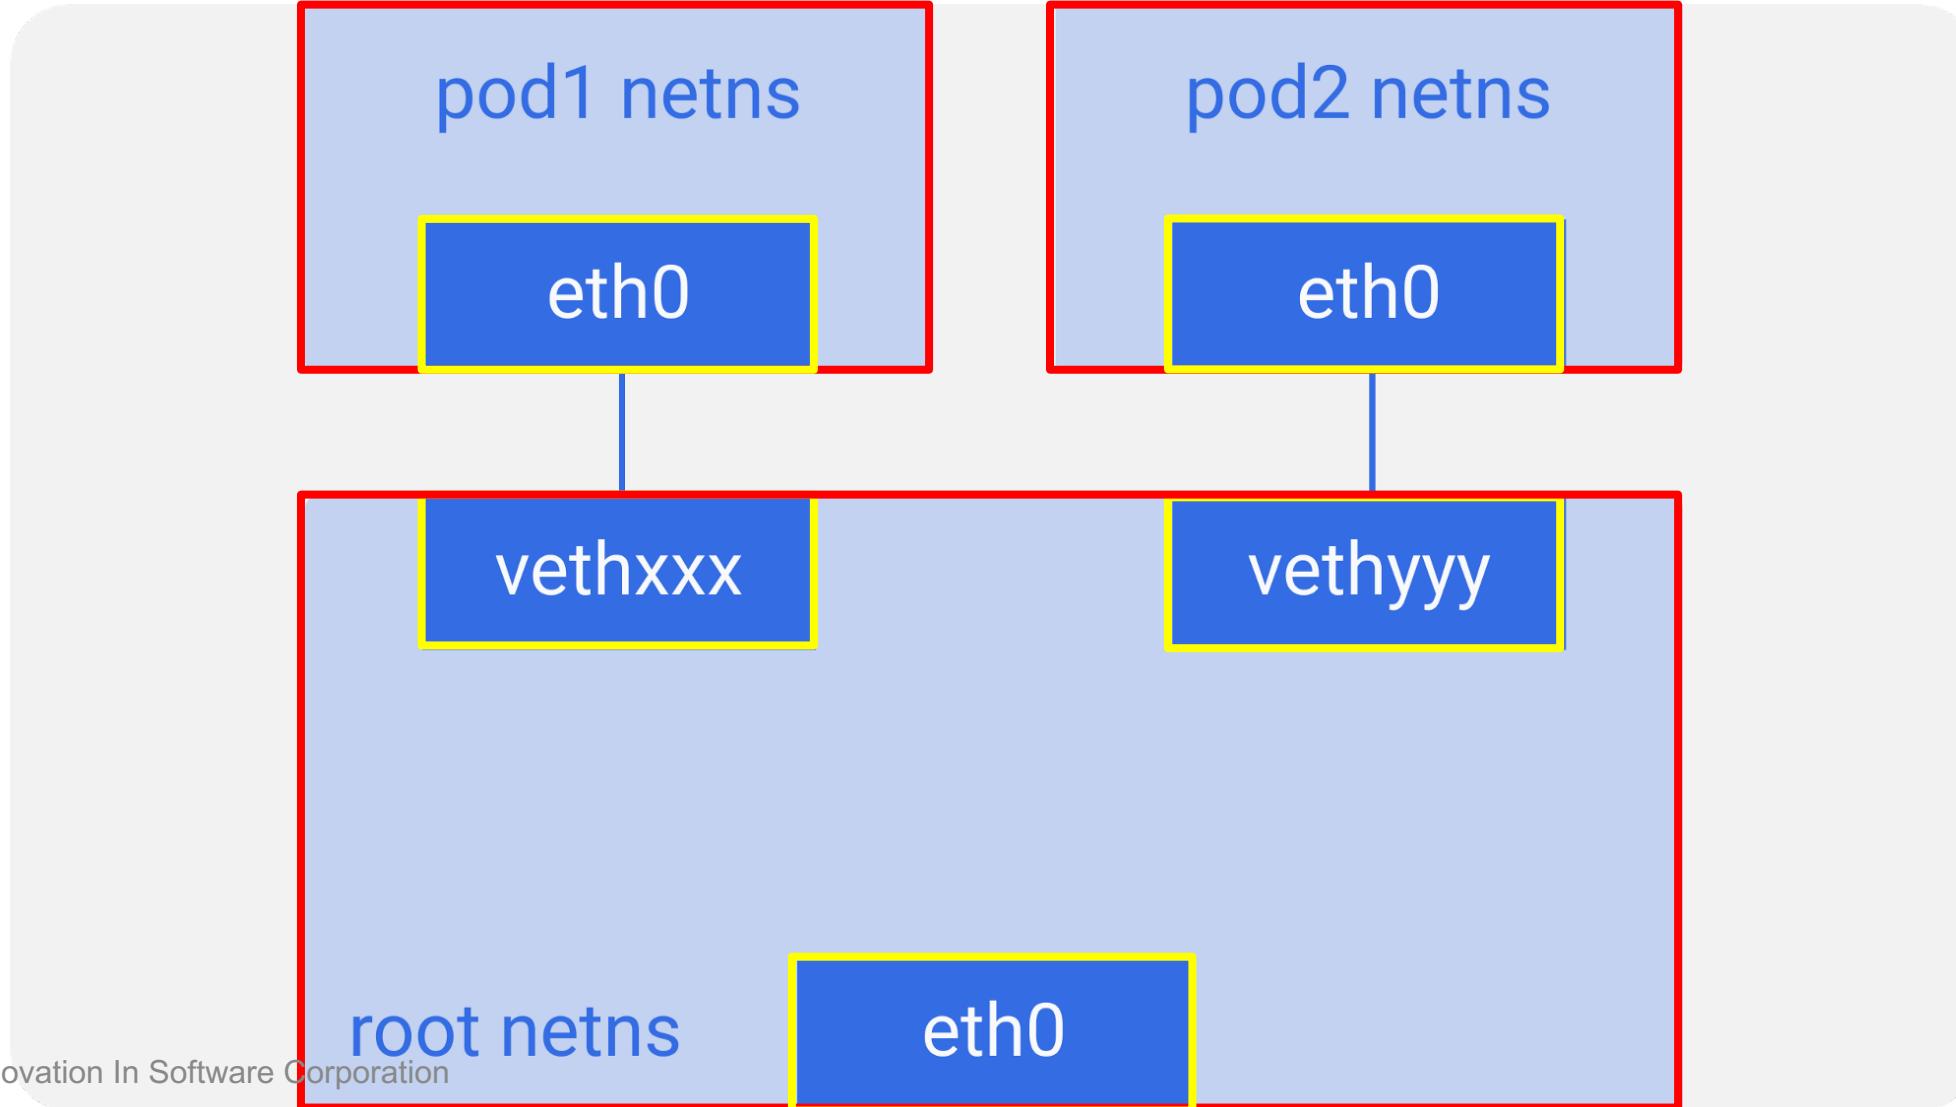
Without an overlay network

- Configure the underlying network fabric (switches, routers, etc.) to be aware of pod IP addresses.
- This does not require the encapsulation provided by an overlay and so can achieve better performance.

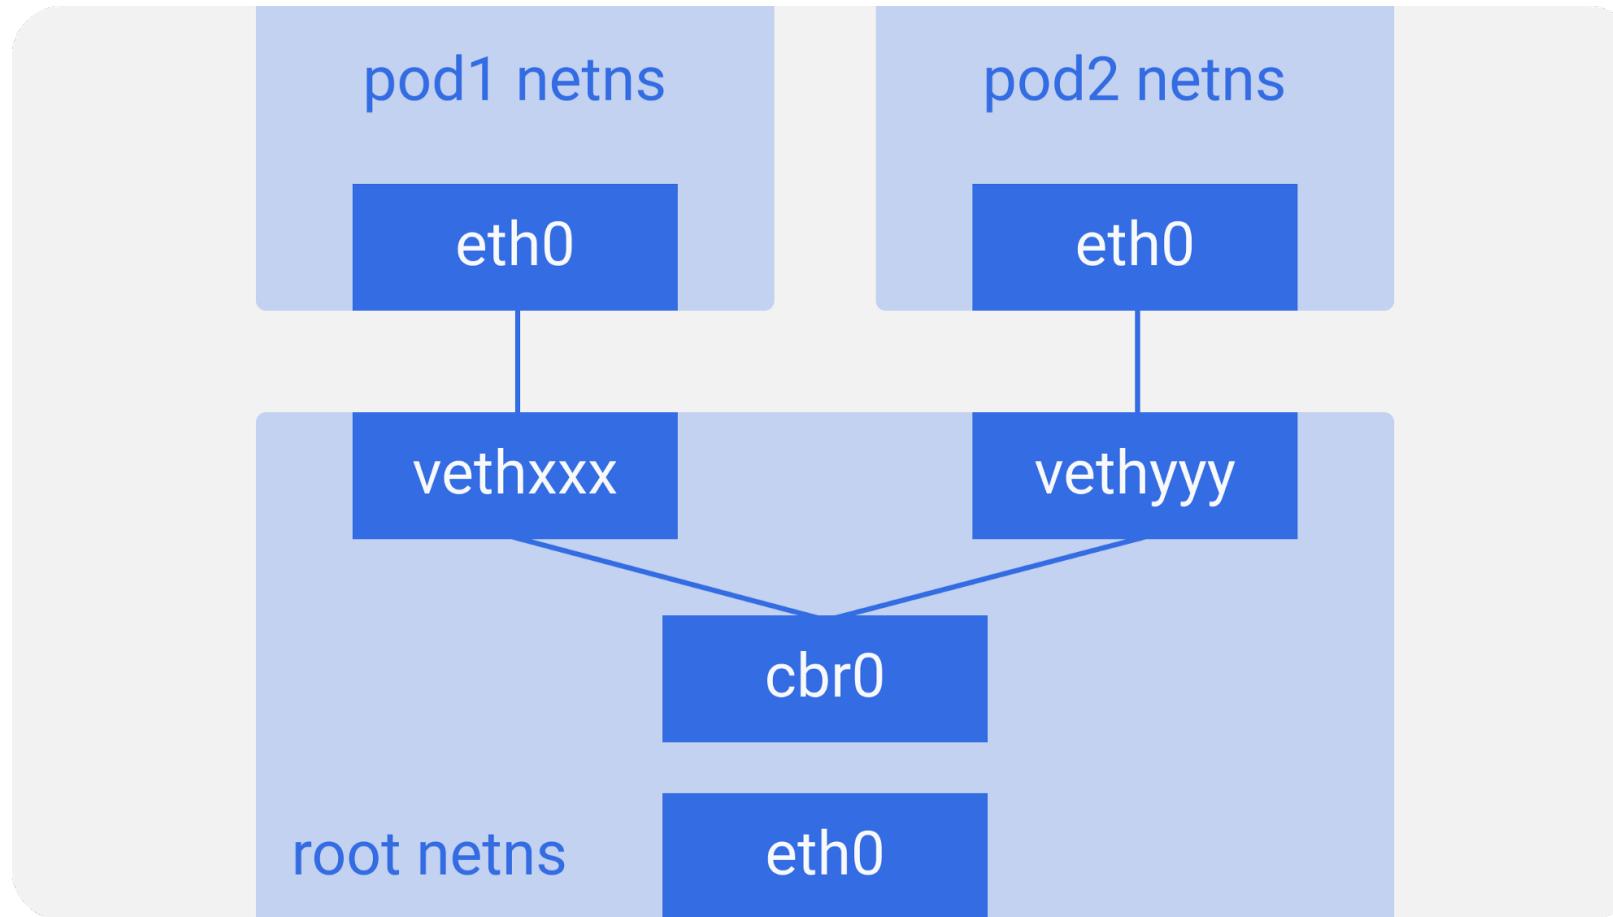
Node Namespace



Node Namespaces

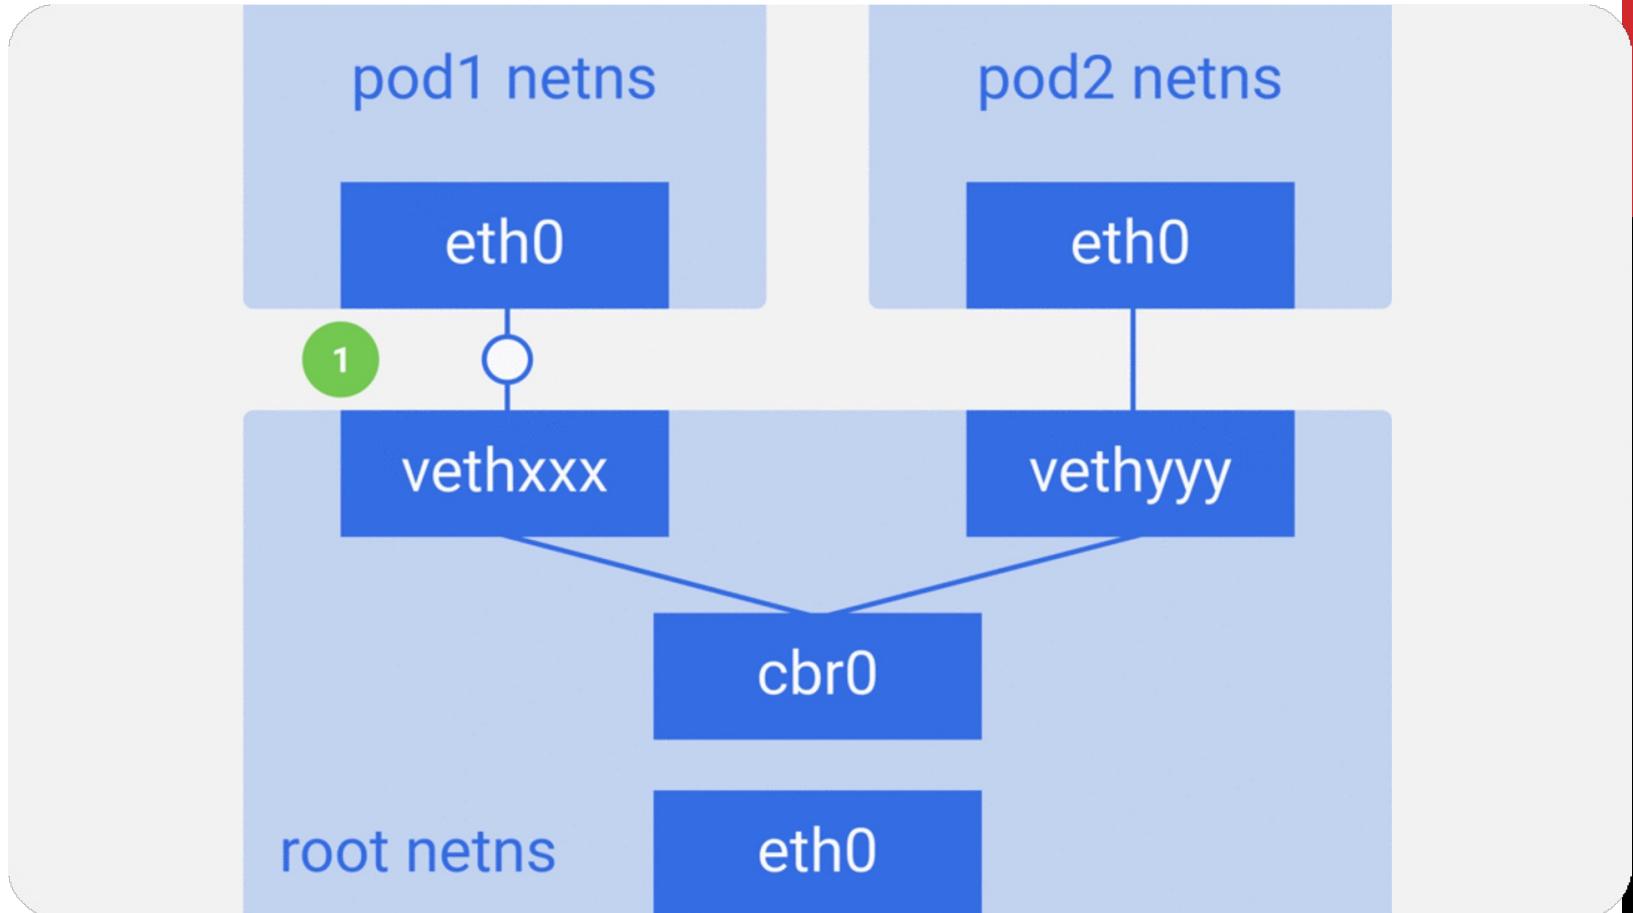


Pod to Host



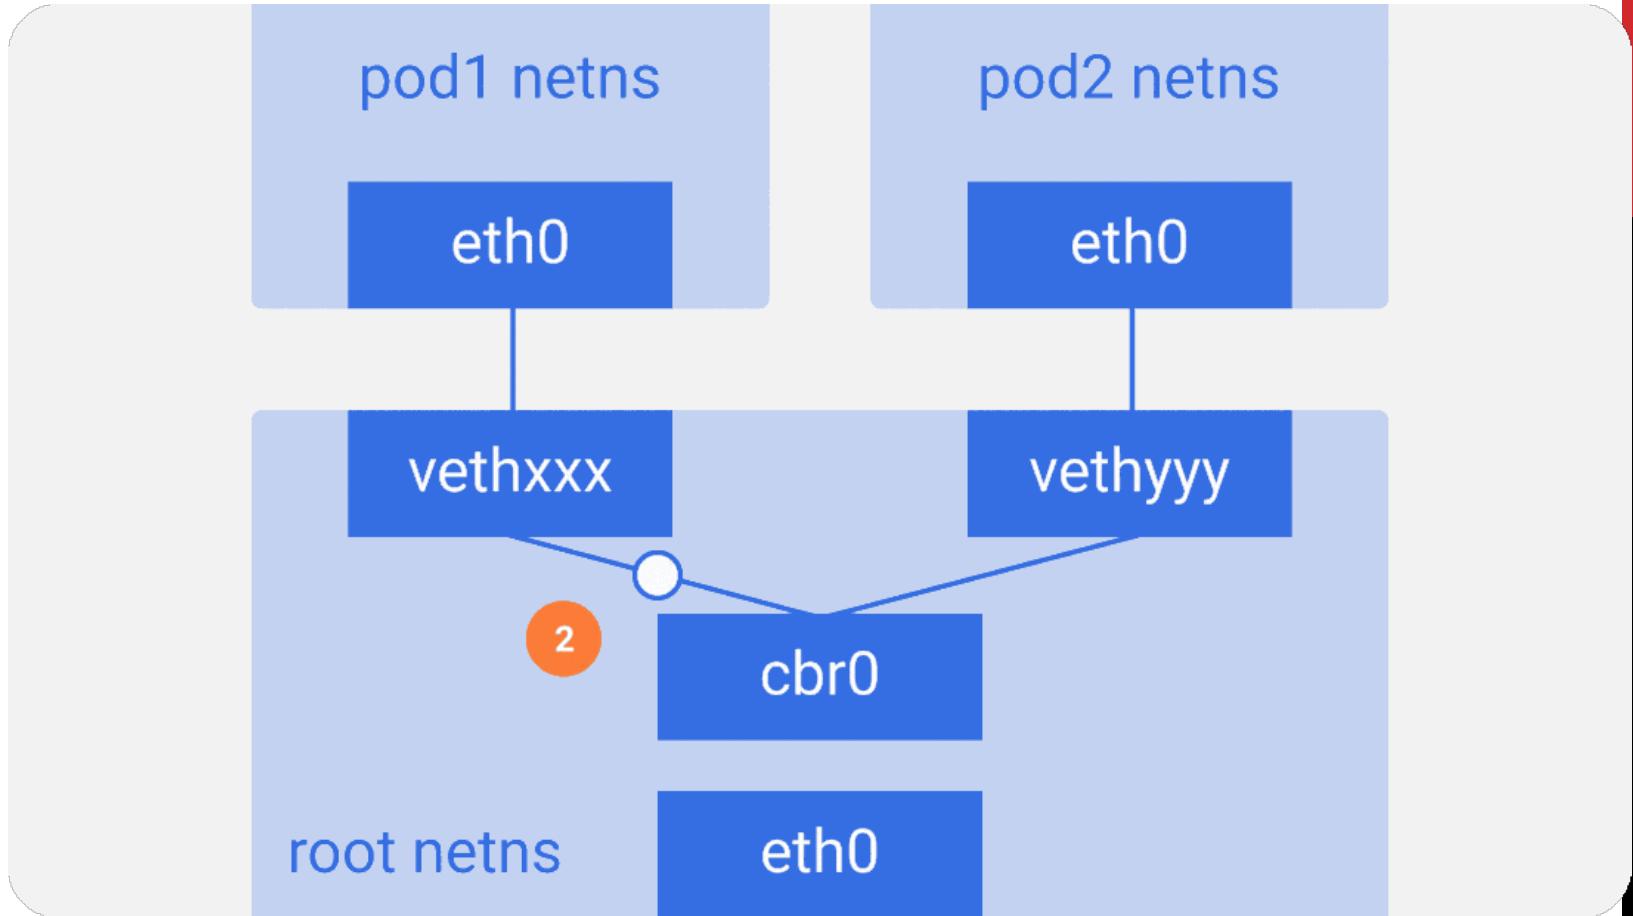
Single Host POD to POD Network

1. Packet leaves pod1's netns at eth0 and enters the root netns at vethxxx.



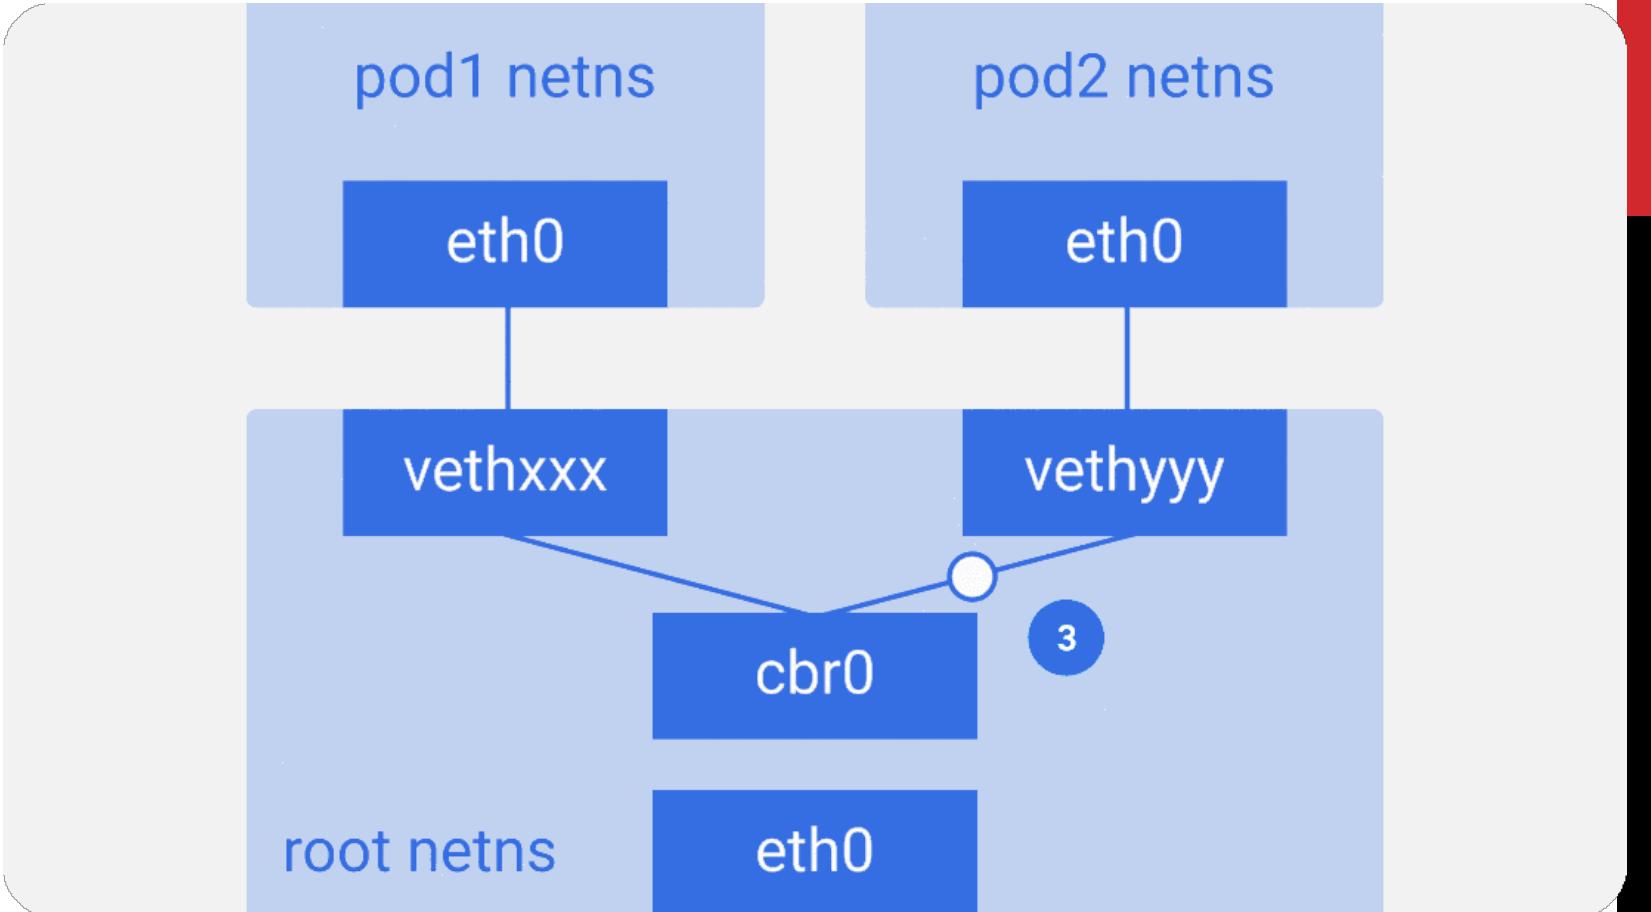
Single Host POD to POD Networking

2. It's passed on to cbr0, which discovers the destination using an ARP request, saying "who has this IP?"



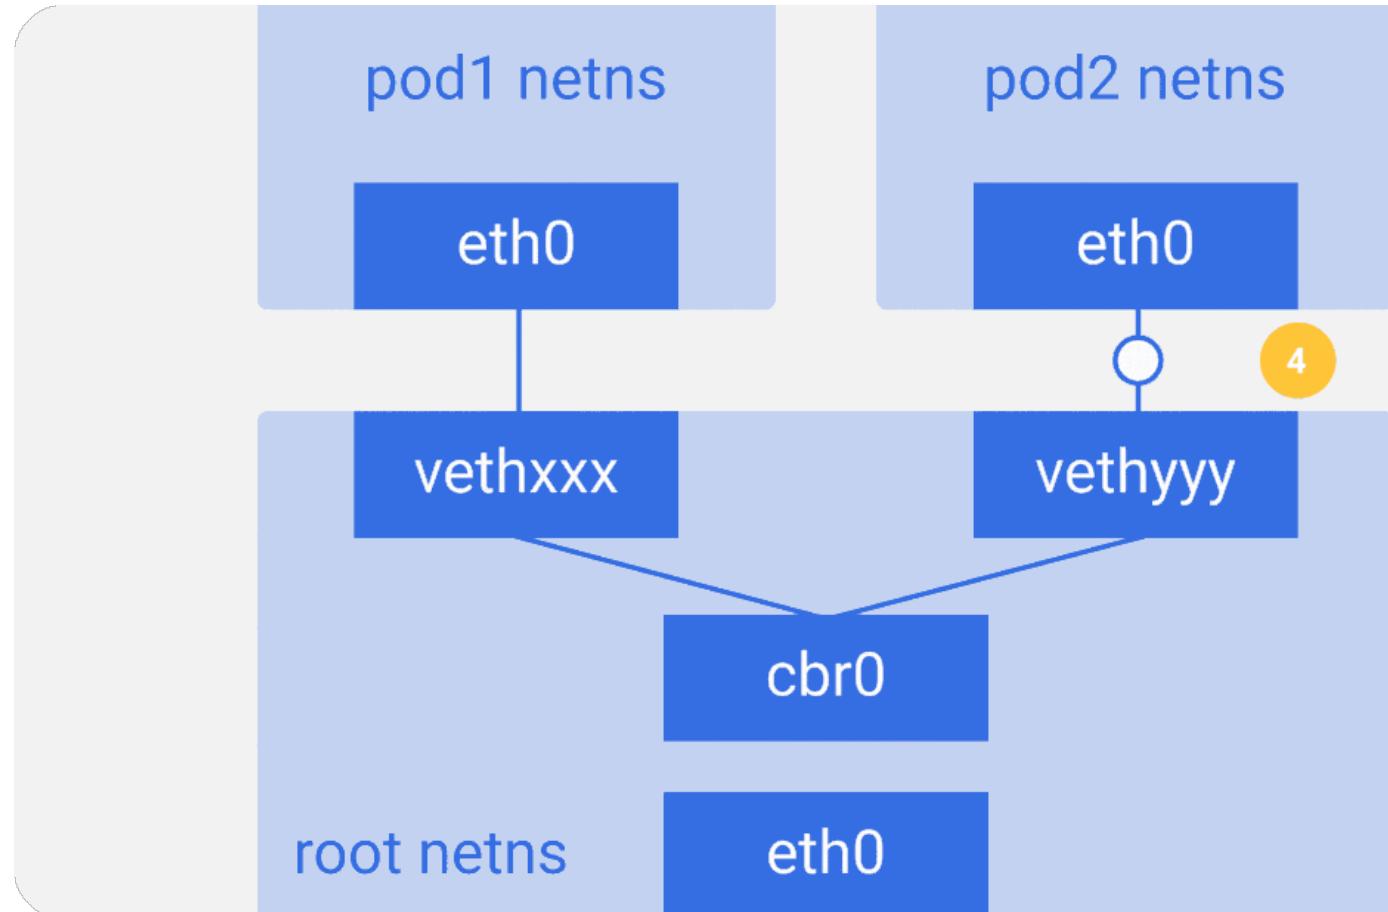
Single Host POD to POD Networking

3. vethyyy says it has that IP, so the bridge knows where to forward the packet.



Single Host POD to POD Networking

4. The packet reaches vethyyy, crosses the pipe-pair and reaches pod2's netns.



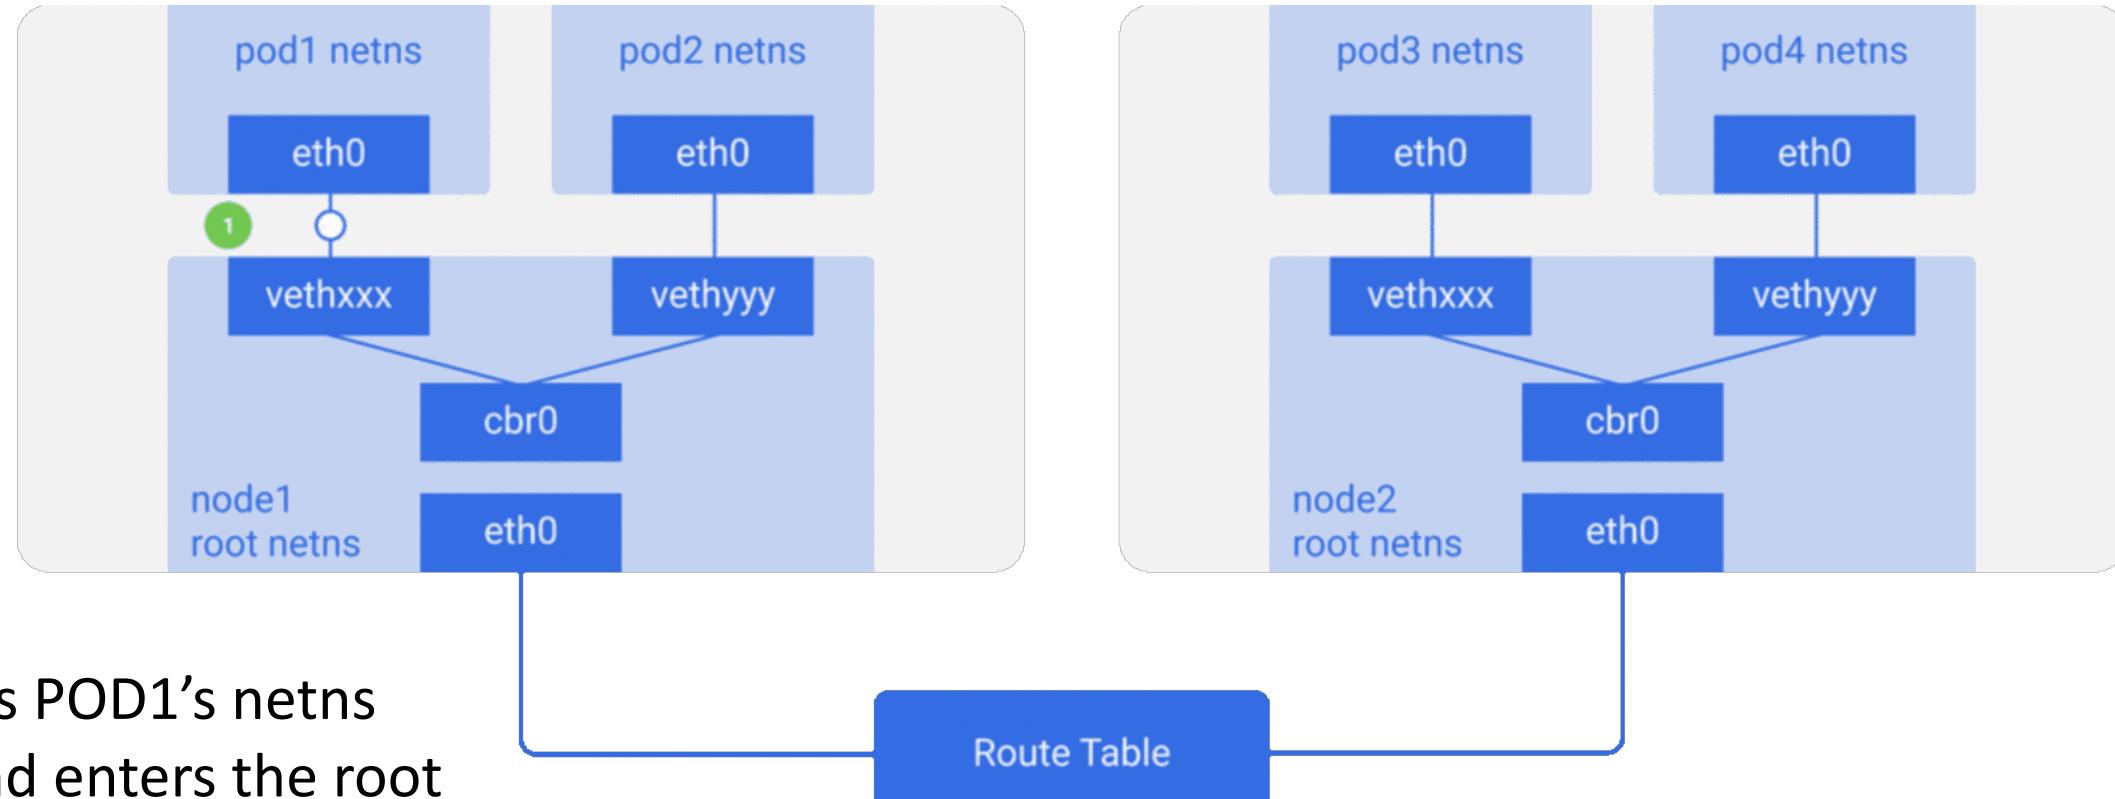
Kubernetes Inter-node Communication



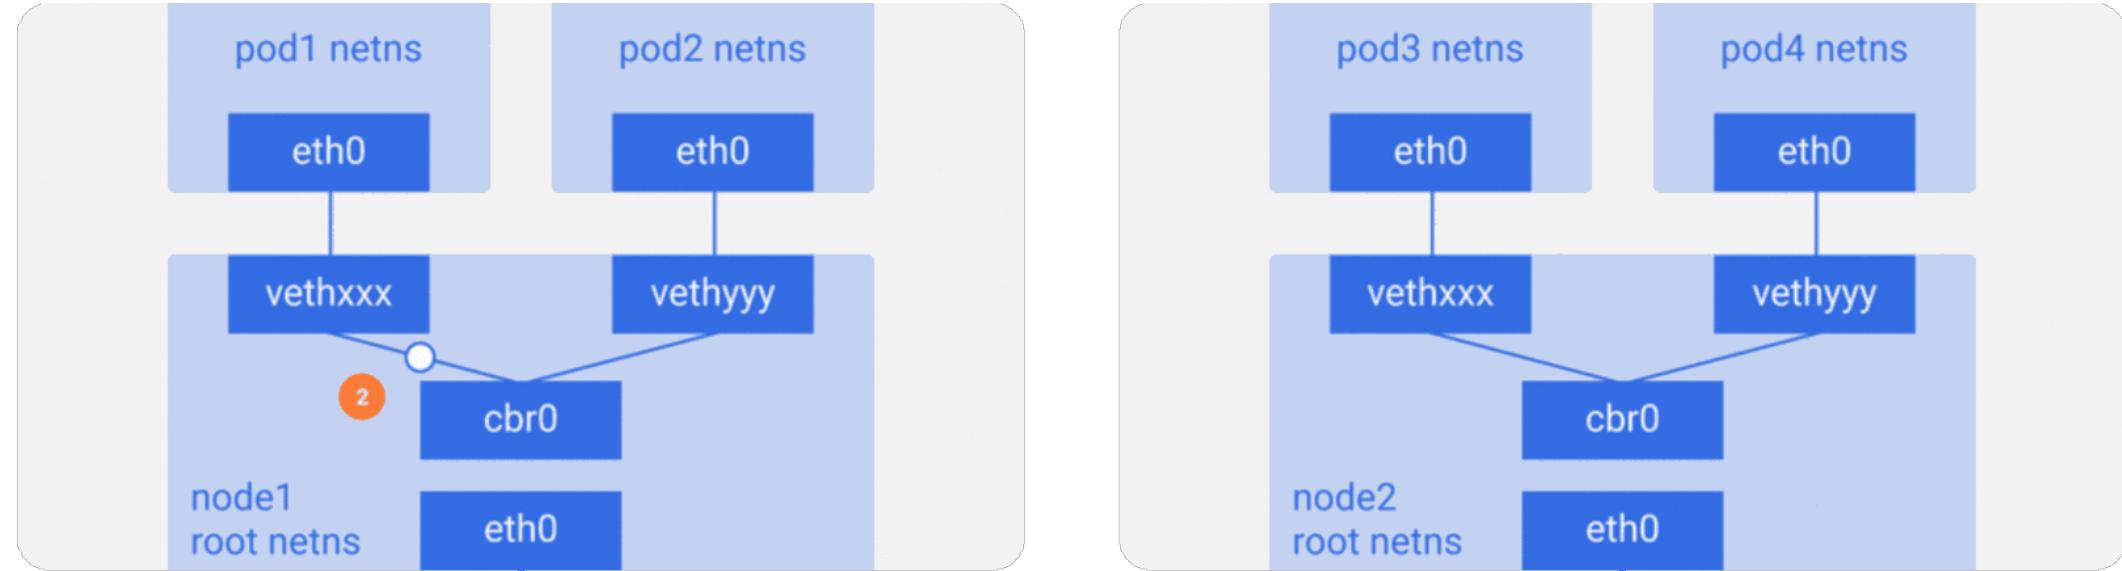
Inter-node Communication

- PODs must be reachable across nodes
- Kubernetes doesn't care how:
 - L2 (ARP across nodes)
 - L3 (IP routing across nodes)
 - Overlay networks

Inter-node Communication

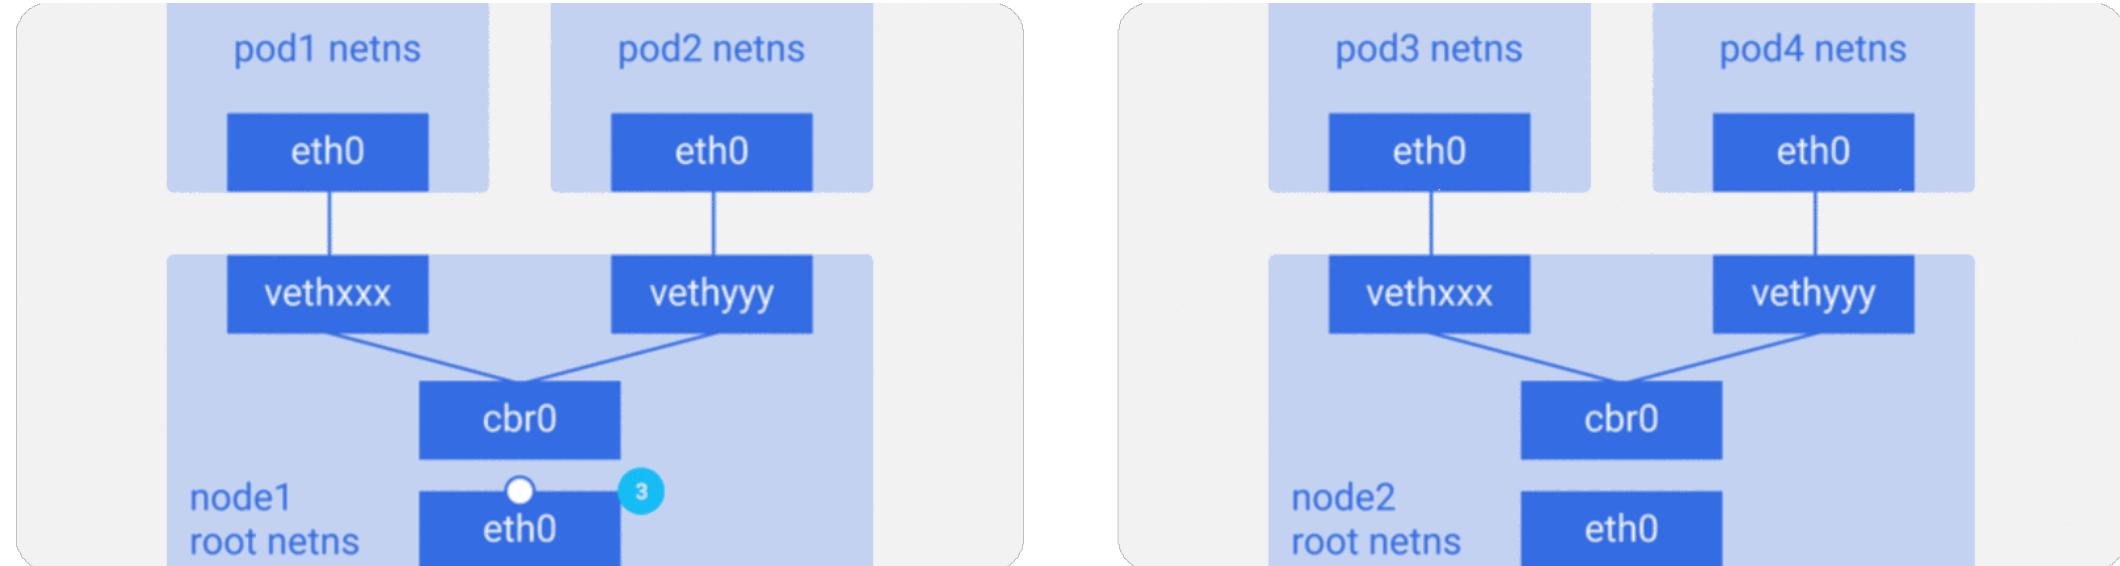


Inter-node Communication



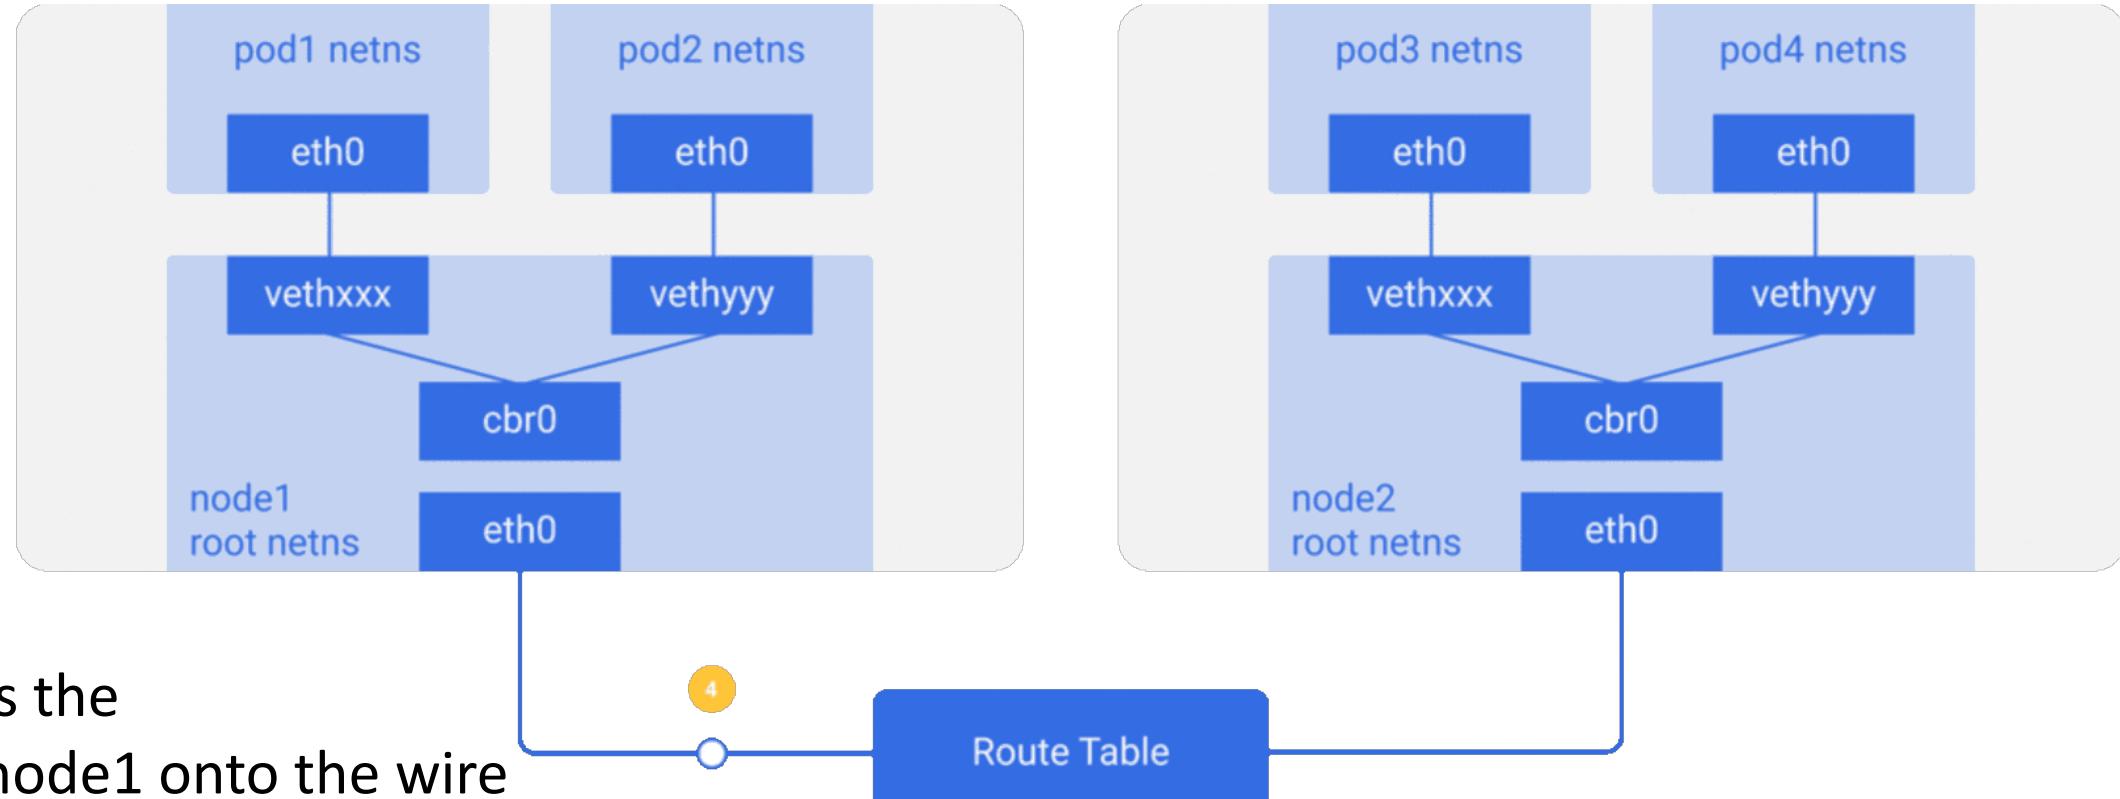
2. It's passed on to `cbr0`,
which makes the ARP request
to find the destination.

Inter-node Communication

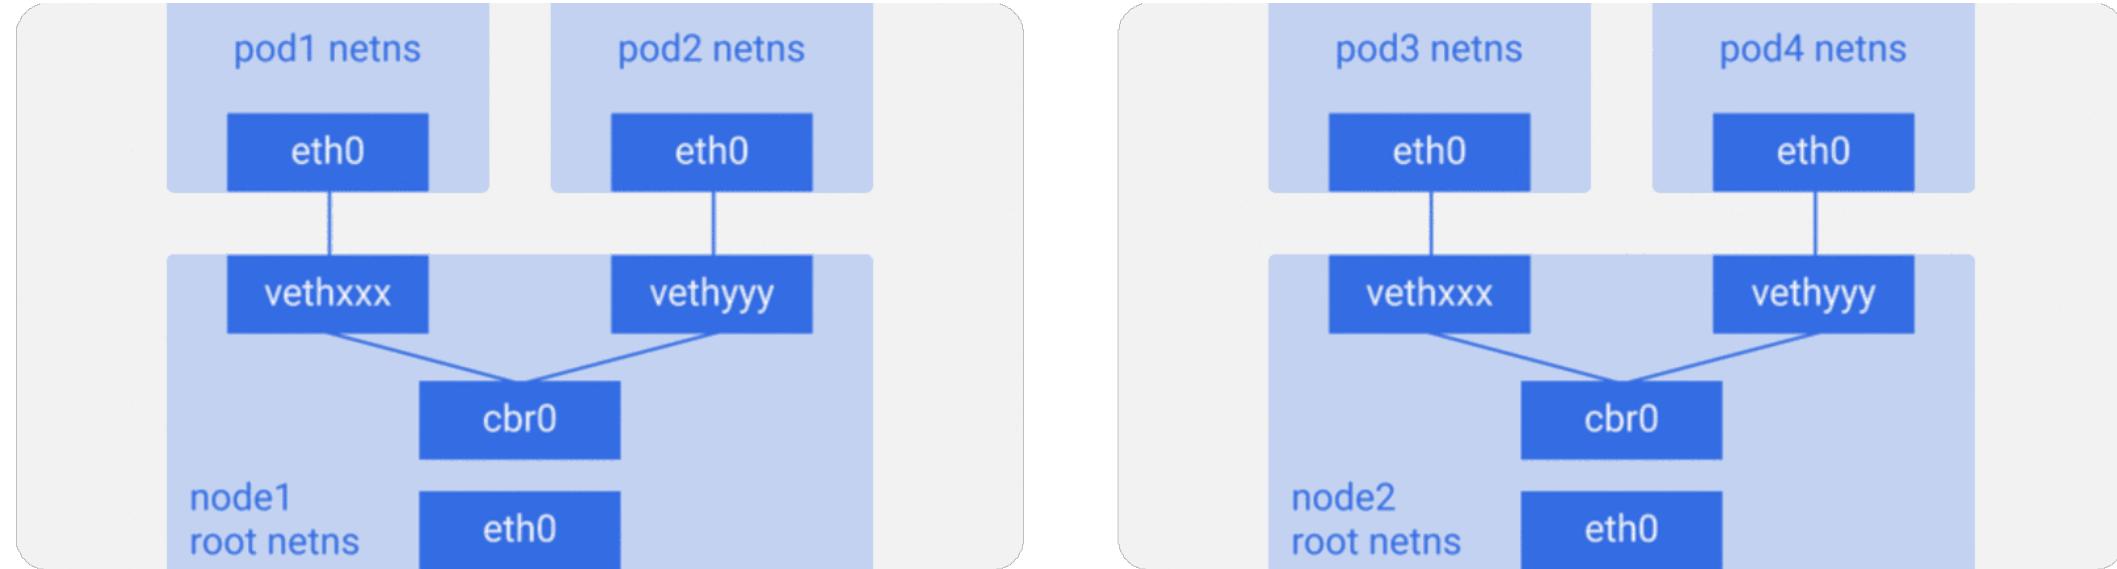


1. It goes to the `cbr0` interface of the node1 root netns
2. It goes to the `cbr0` interface of the node2 root netns
3. It comes out of `cbr0` to the main network interface `eth0` since nobody on this node has the IP address for POD4.

Inter-node Communication

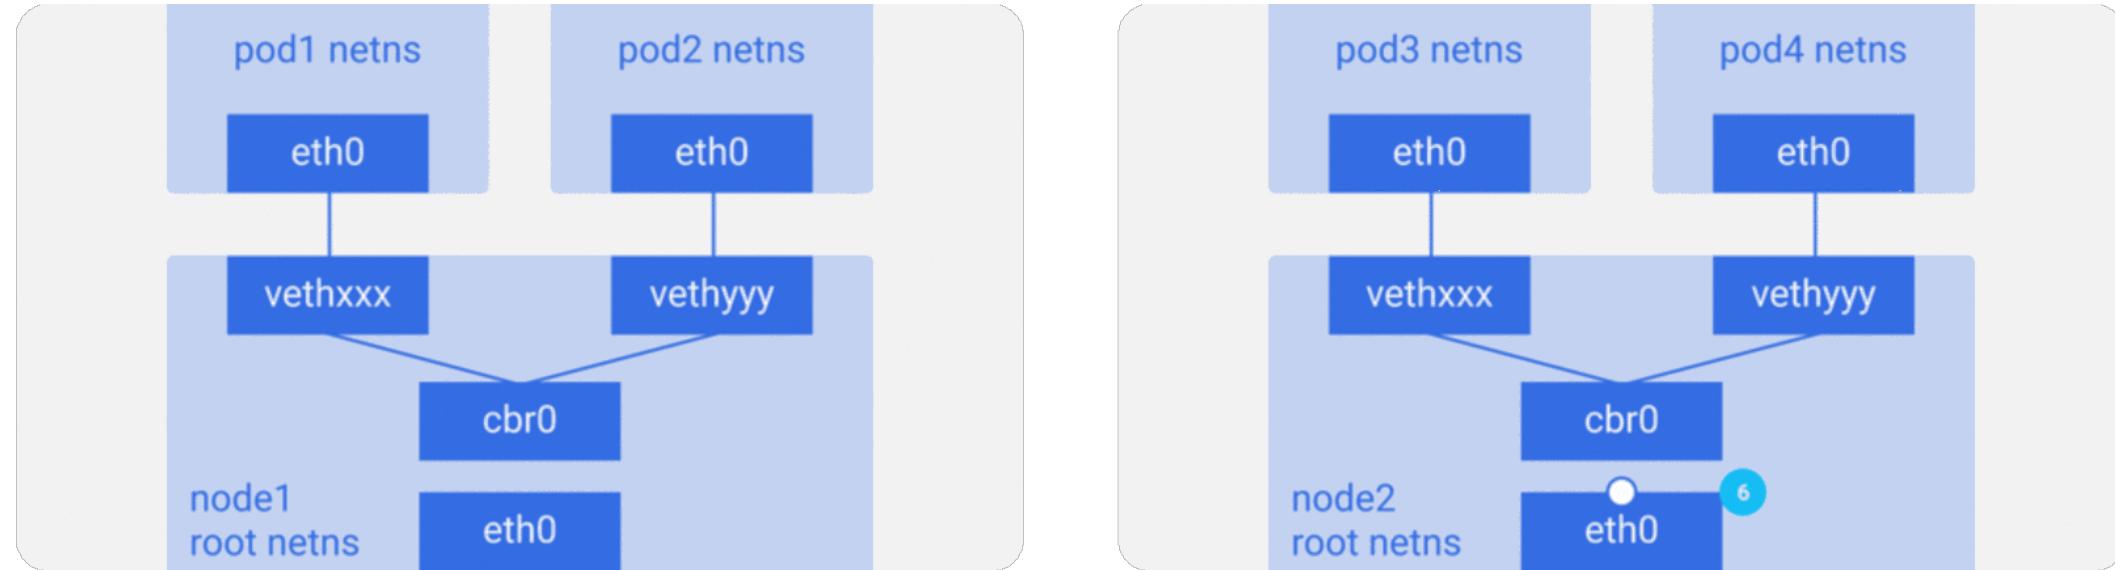


Inter-node Communication



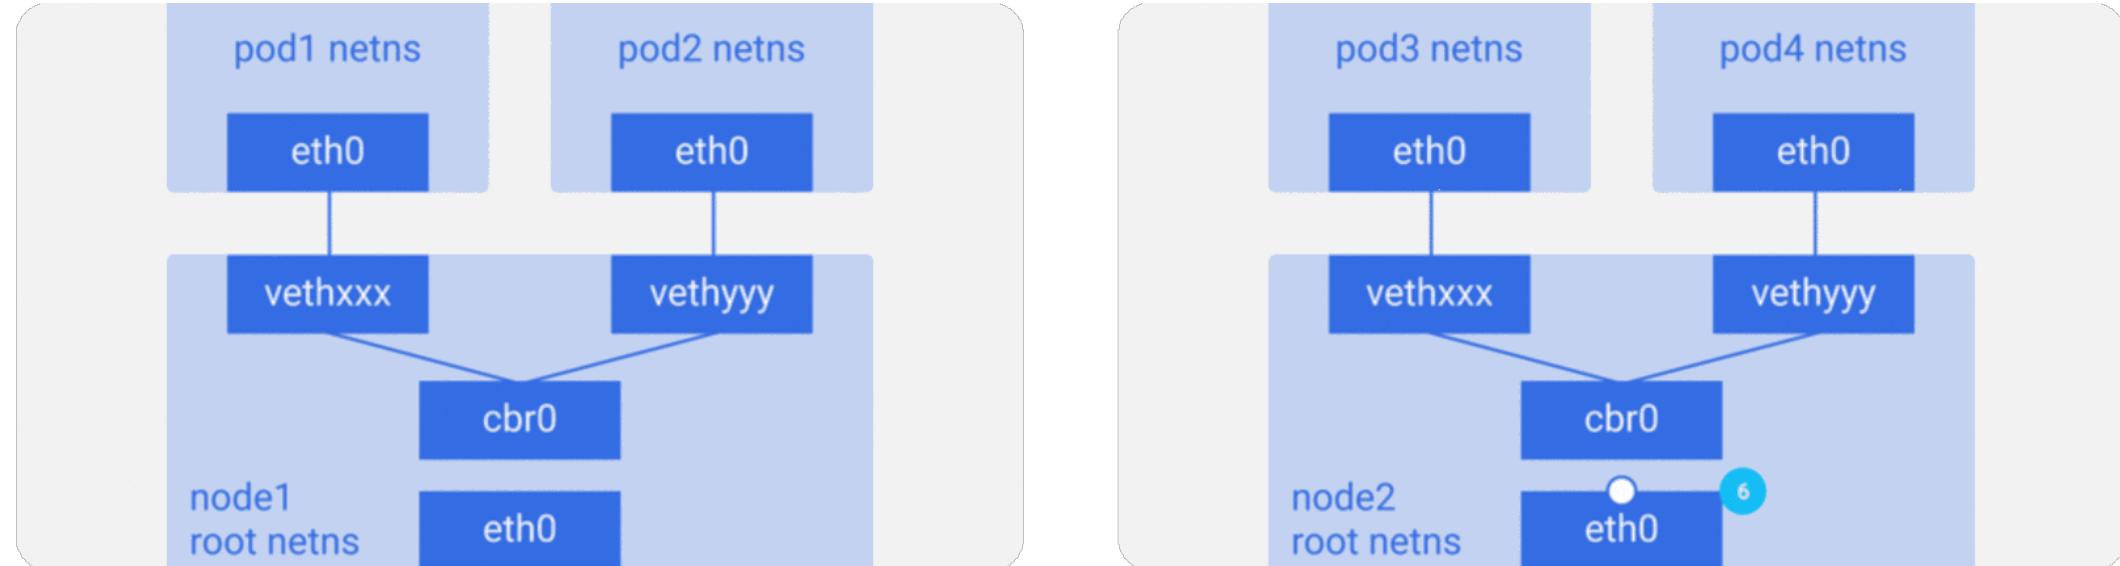
5. The route table has routes setup for each of the node CIDR blocks, and it routes the packet to the node whose CIDR block contains the POD4 IP.

Inter-node Communication



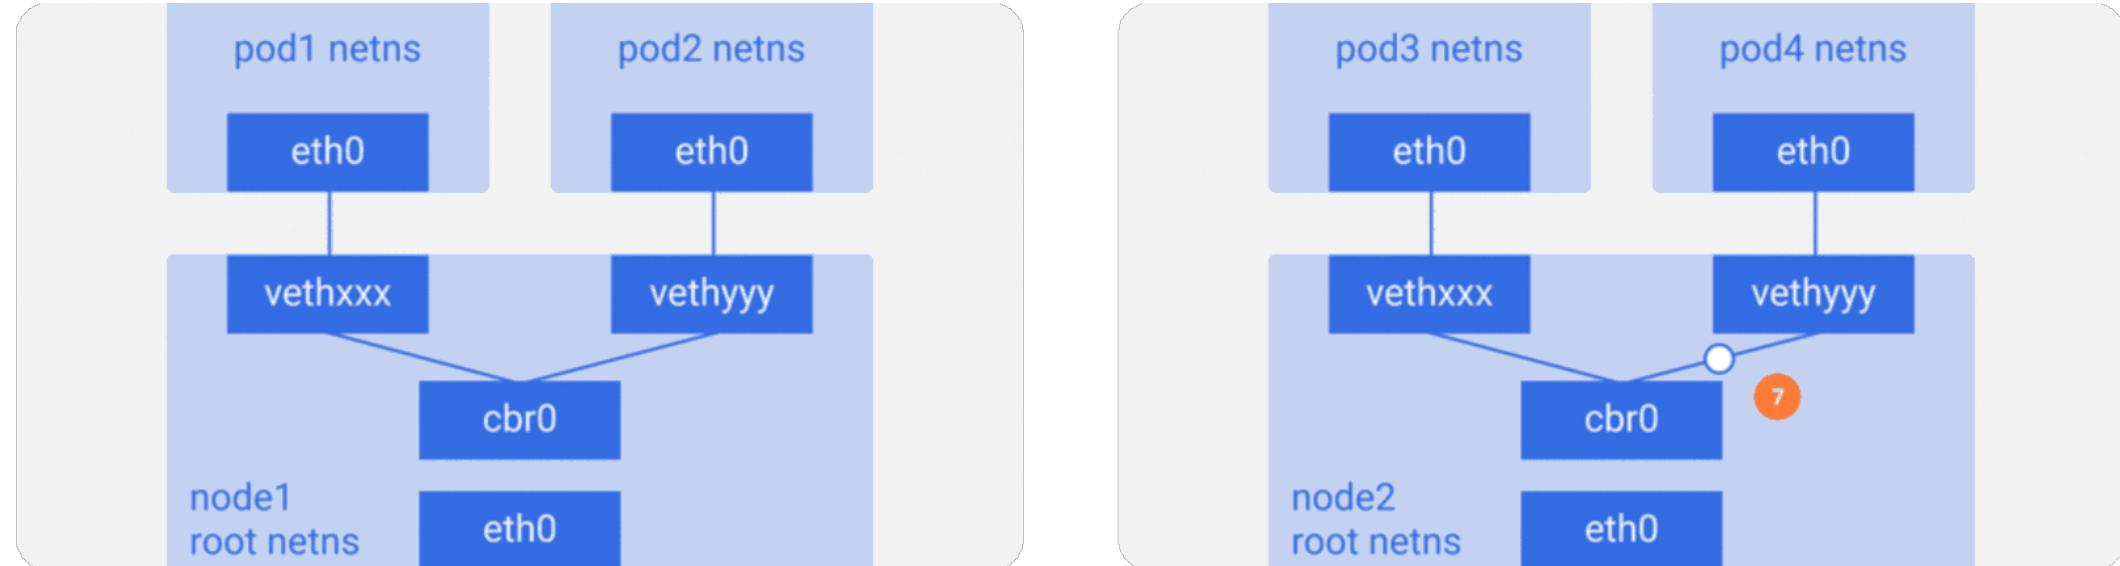
6. Packet arrives at node2 at eth0. POD4 isn't the IP of eth0, the packet is forwarded to cbr0, nodes are configured with IP forwarding enabled.

Inter-node Communication



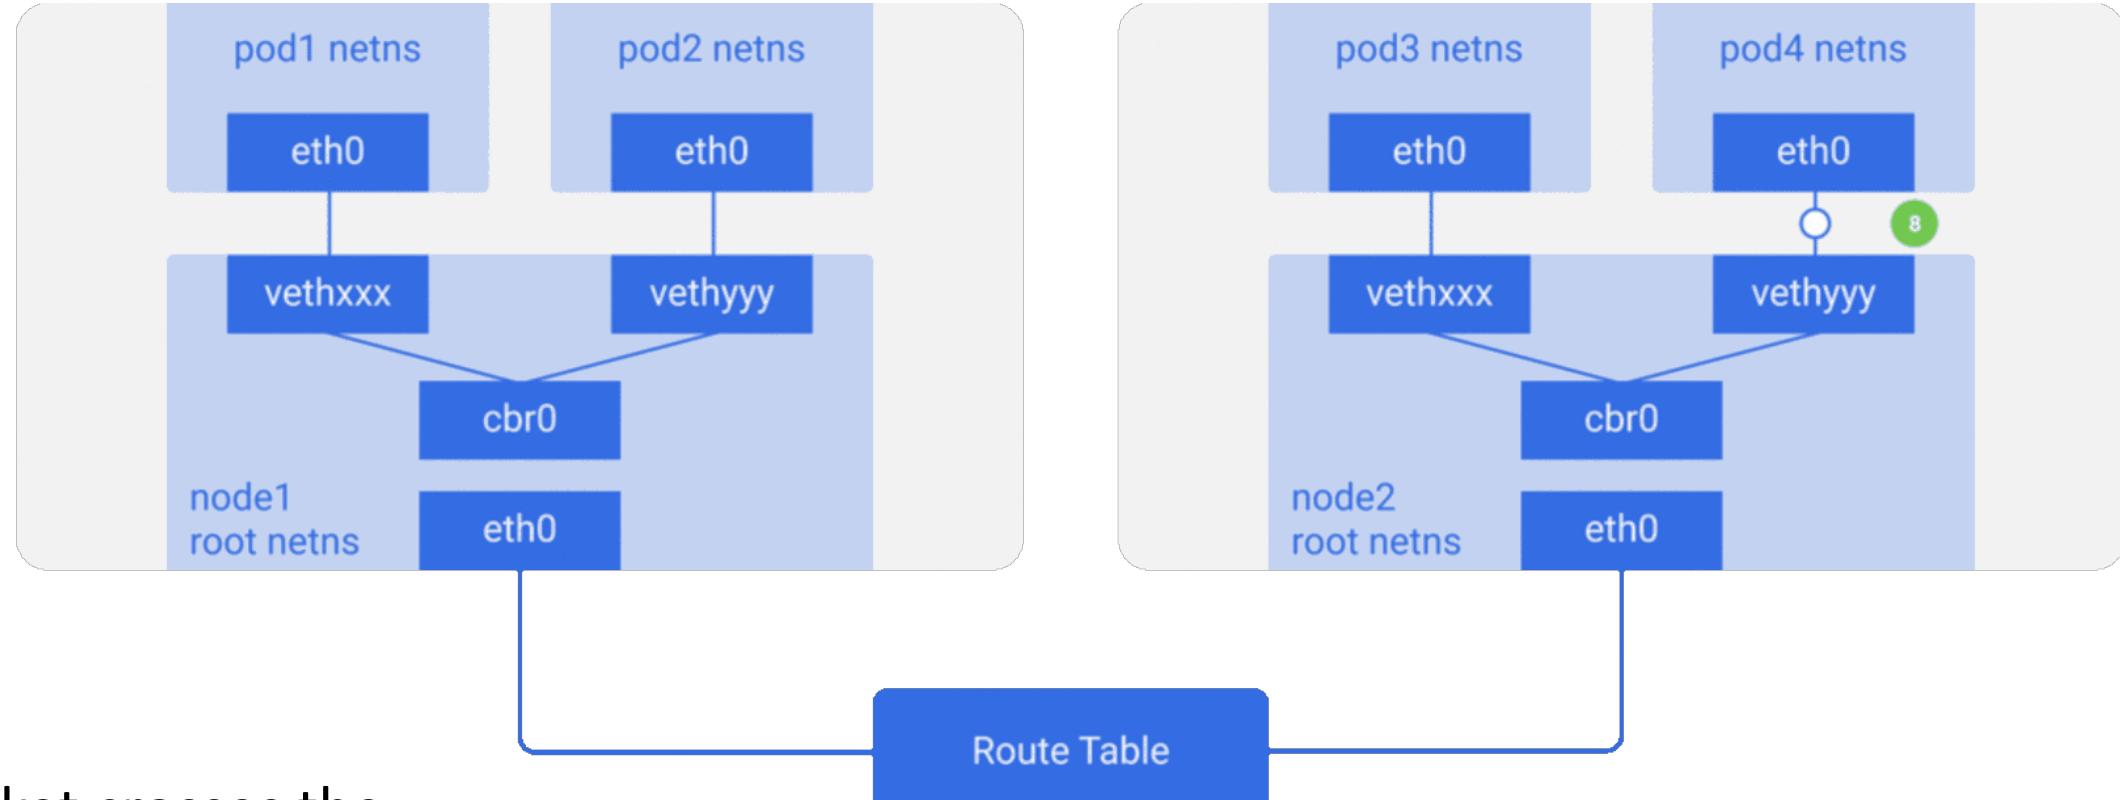
6. The node's routing table is looked up for any routes matching the POD4 IP. It finds cbr0 as the destination for this node's CIDR block.

Inter-node Communication



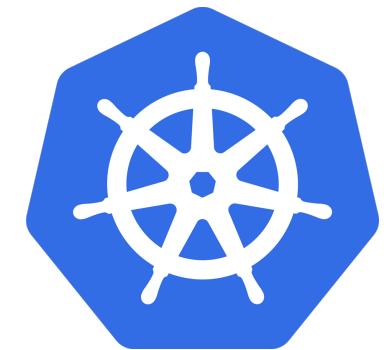
7. The bridge takes the packet, makes an ARP request and finds out that the IP belongs to vethyyy.

Inter-node Communication



8. The packet crosses the pipe-pair and reaches POD4

Kubernetes Overlay Networking



Overlay Networking

- Not required but very common
- Provide virtual IP space
- Avoid route limits (Public Cloud)
- Encapsulating a packet-in-packet, which traverses the native network across nodes.
- Can reduce throughput

Flannel Overlay Networking



1. The packet leaves POD1's netns at eth0 and enters the root netns at vethxxx.

Flannel Overlay Networking

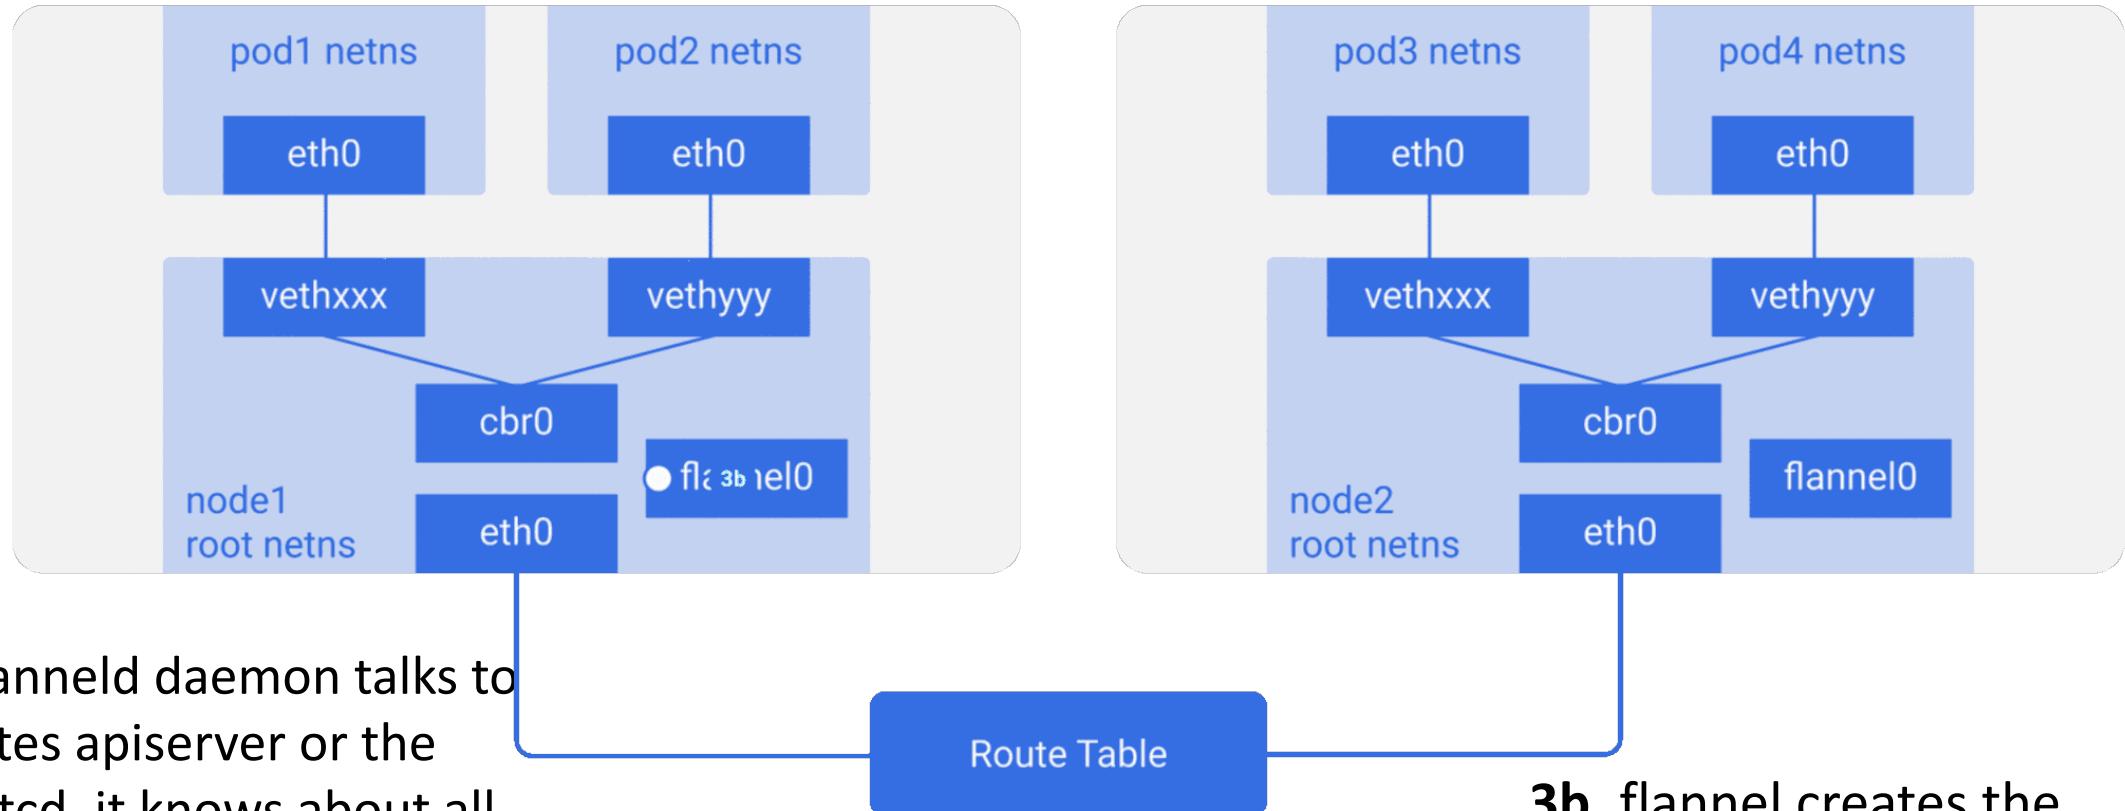


2. It's passed on to `cbr0`, which makes the ARP request to find the destination.

Flannel Overlay Networking



Flannel Overlay Networking



3b. As the flanneld daemon talks to the Kubernetes apiserver or the underlying etcd, it knows about all the pod IPs, what nodes they're on.

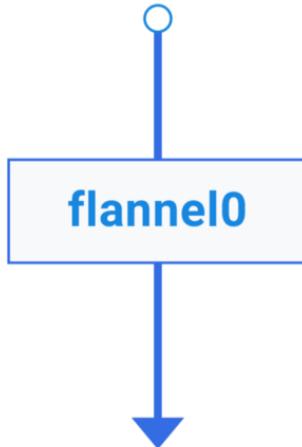
Route Table

3b. flannel creates the mappings (in userspace) for pods IPs to node IPs.

Flannel Overlay Networking



flannel0 takes this packet and wraps it in a UDP packet with extra headers, changing the source and destination IPs to the respective nodes, and sends it to a special VXLAN port (generally 8472).

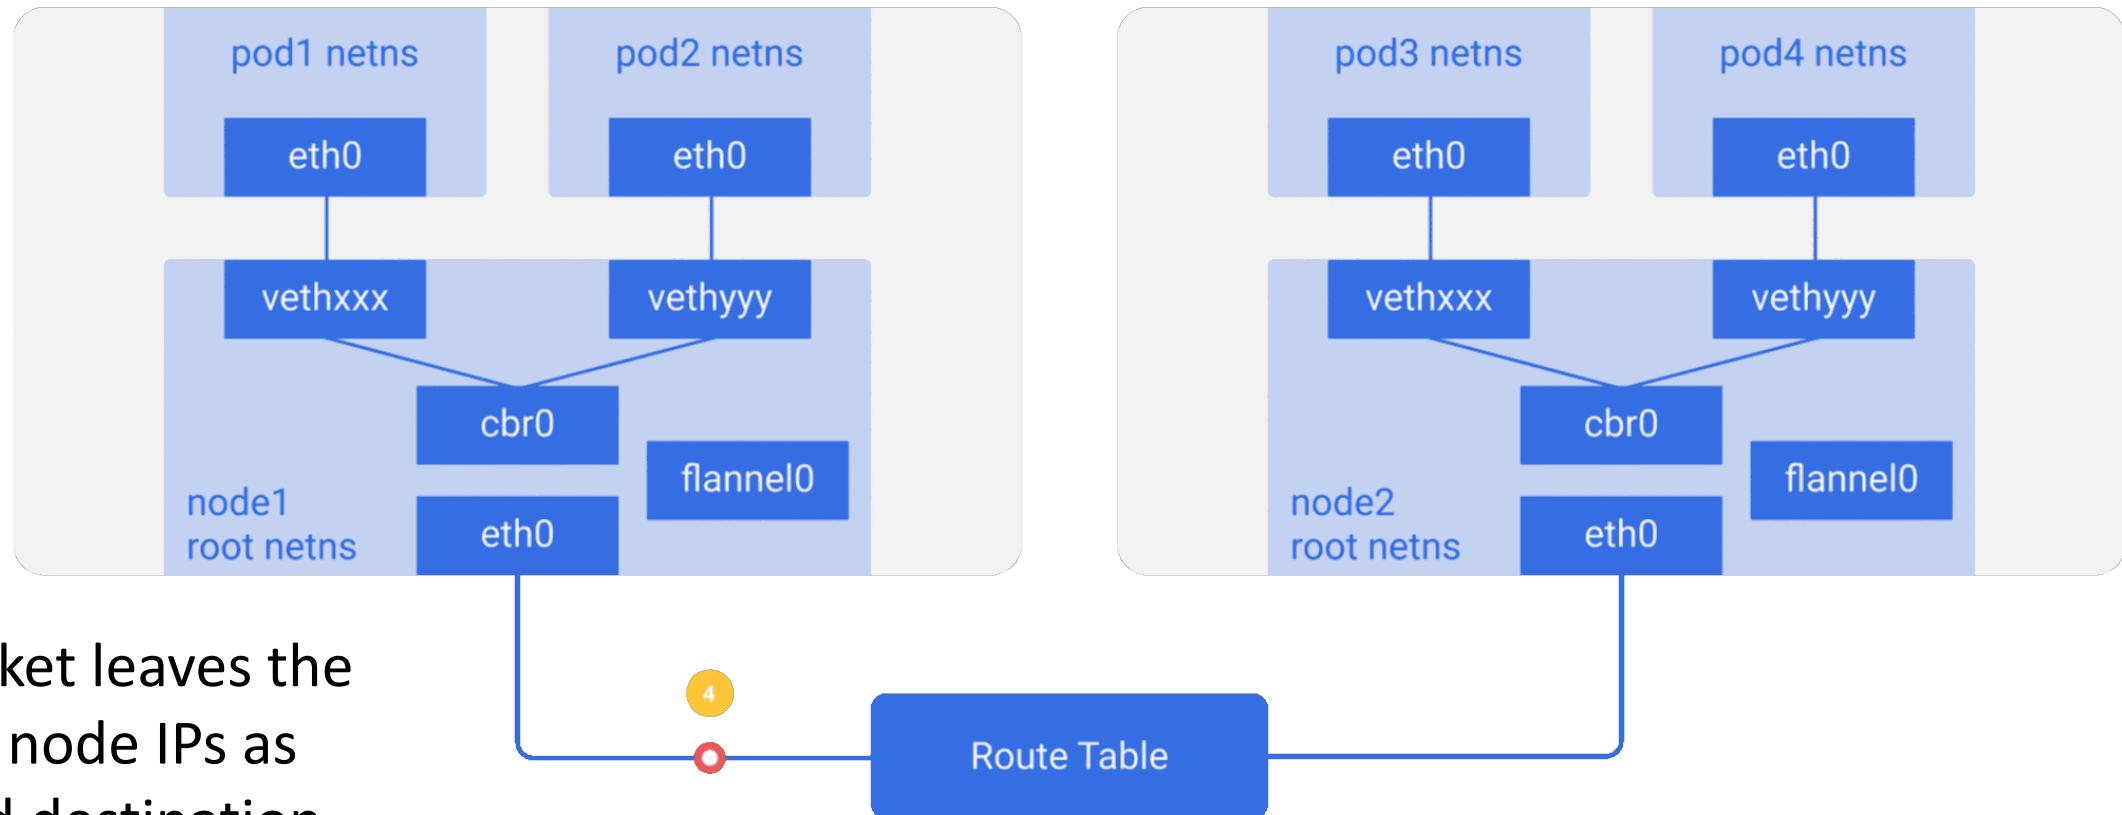


Flannel Overlay Networking

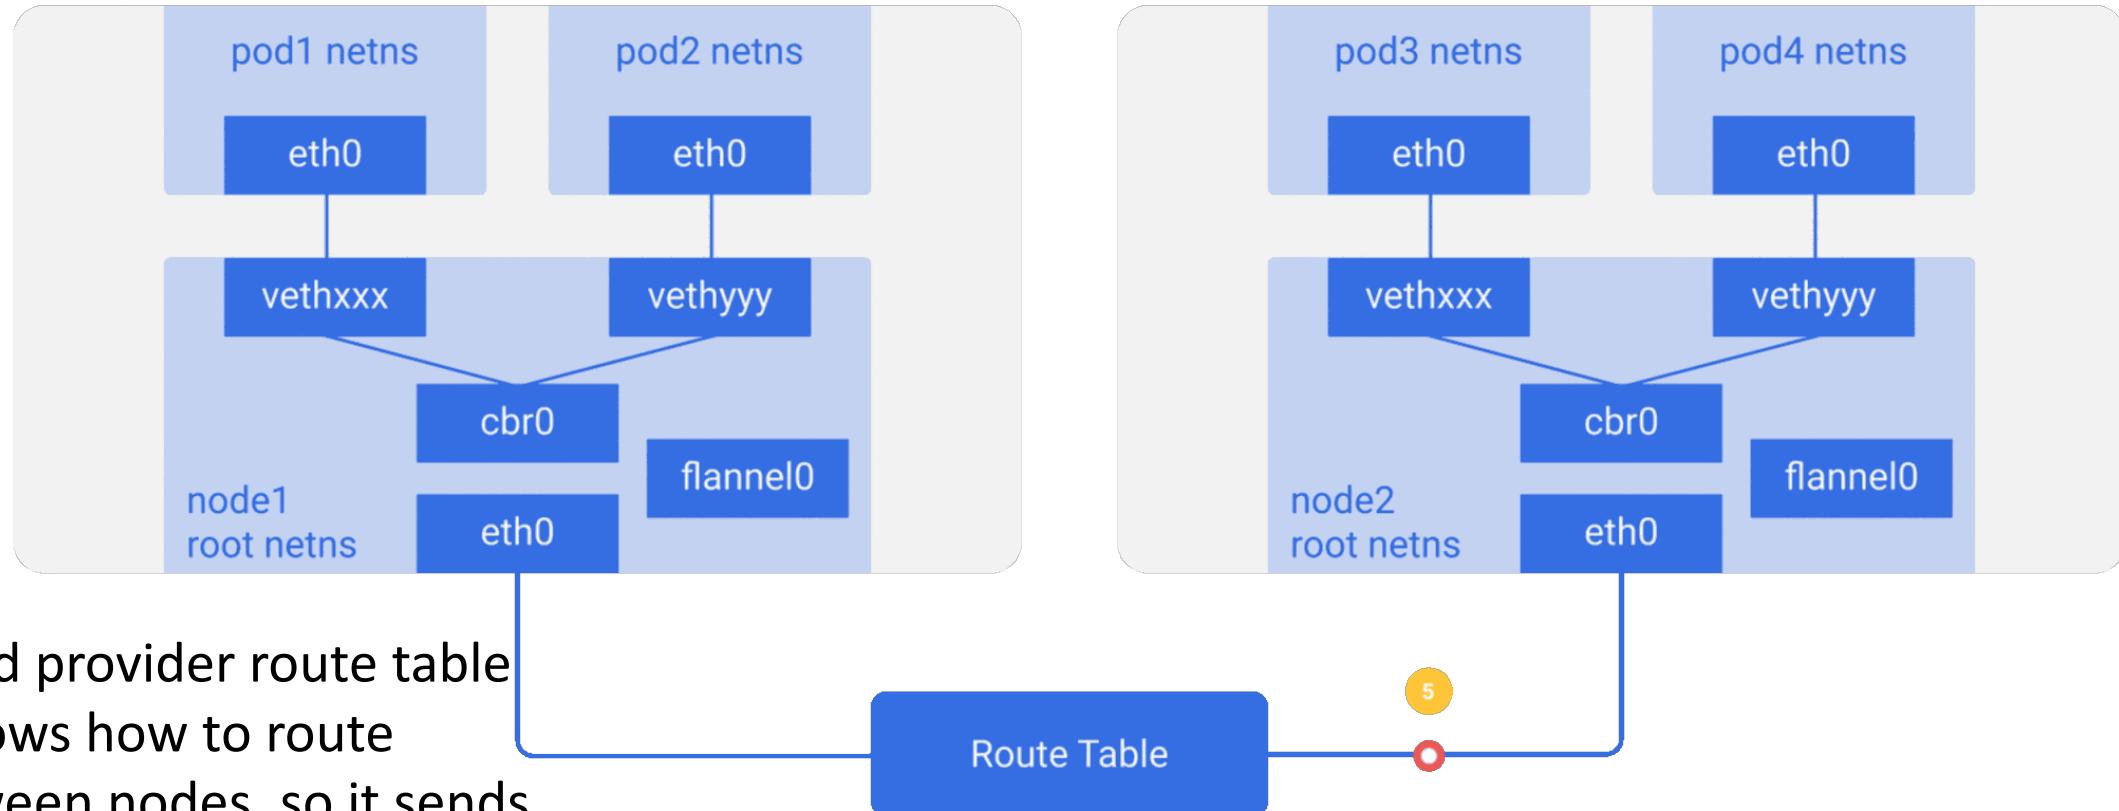


3c. The encapsulated packet is sent out via **eth0** since it is involved in routing the node traffic.

Flannel Overlay Networking



Flannel Overlay Networking

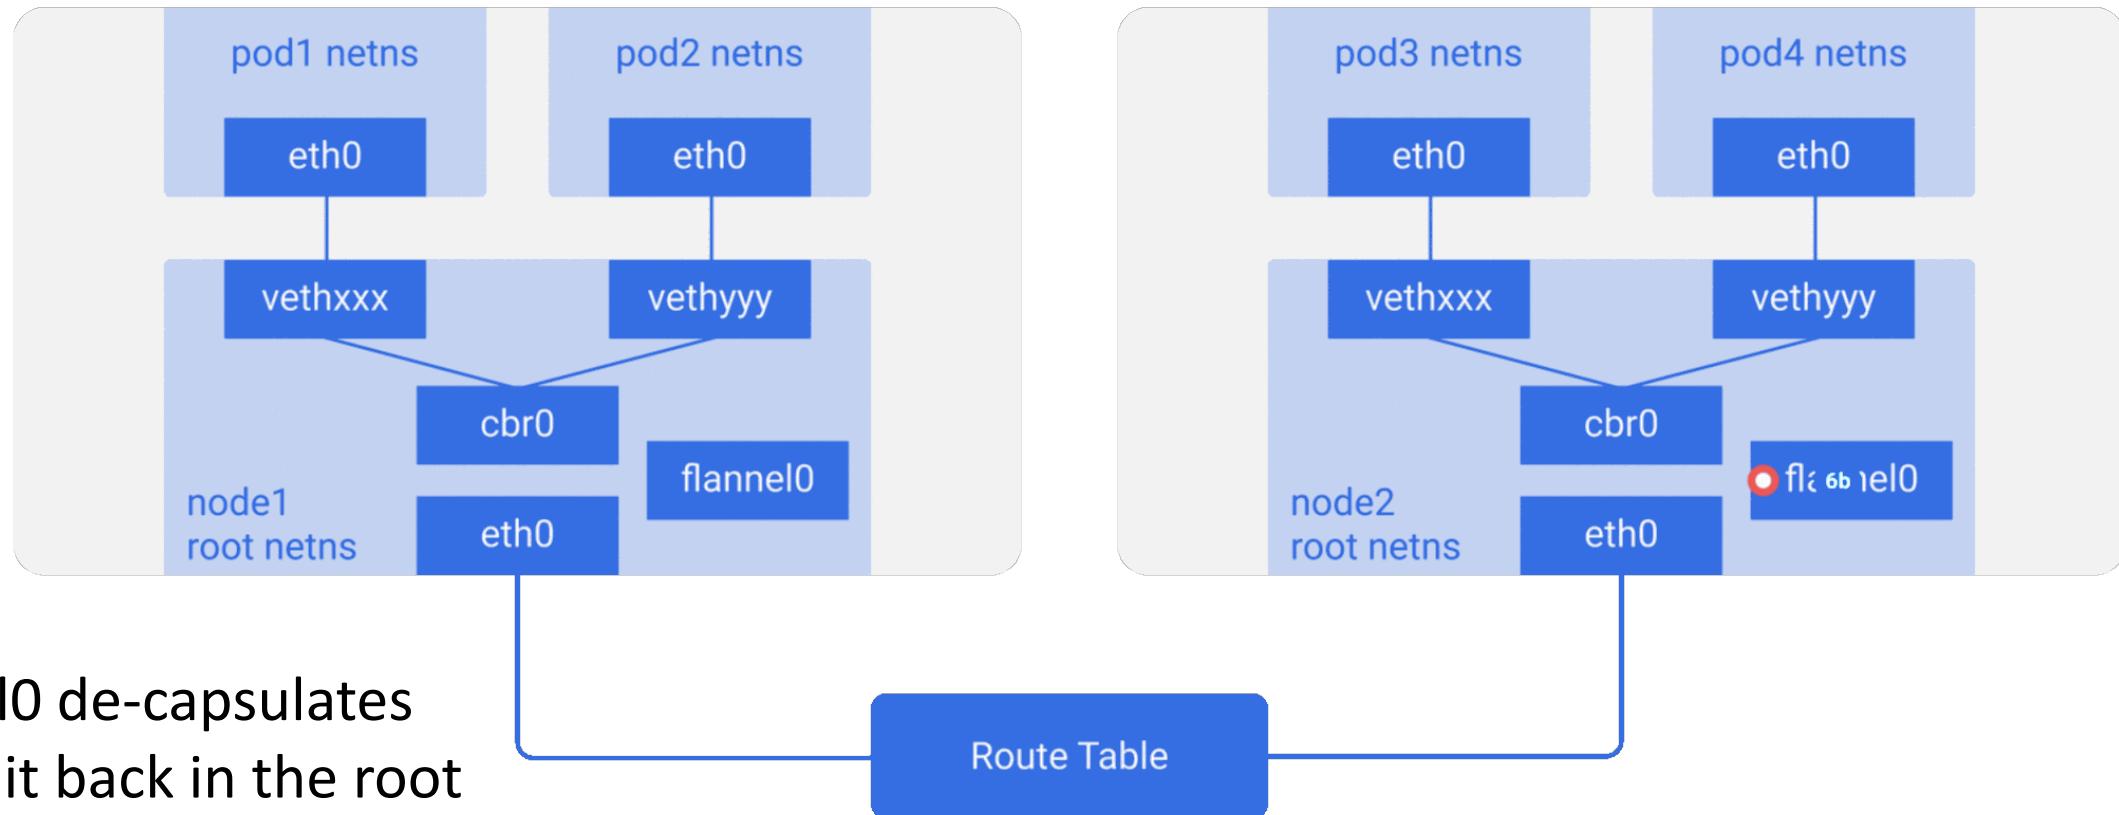


5. The cloud provider route table already knows how to route traffic between nodes, so it sends the packet to destination node2.

Flannel Overlay Networking

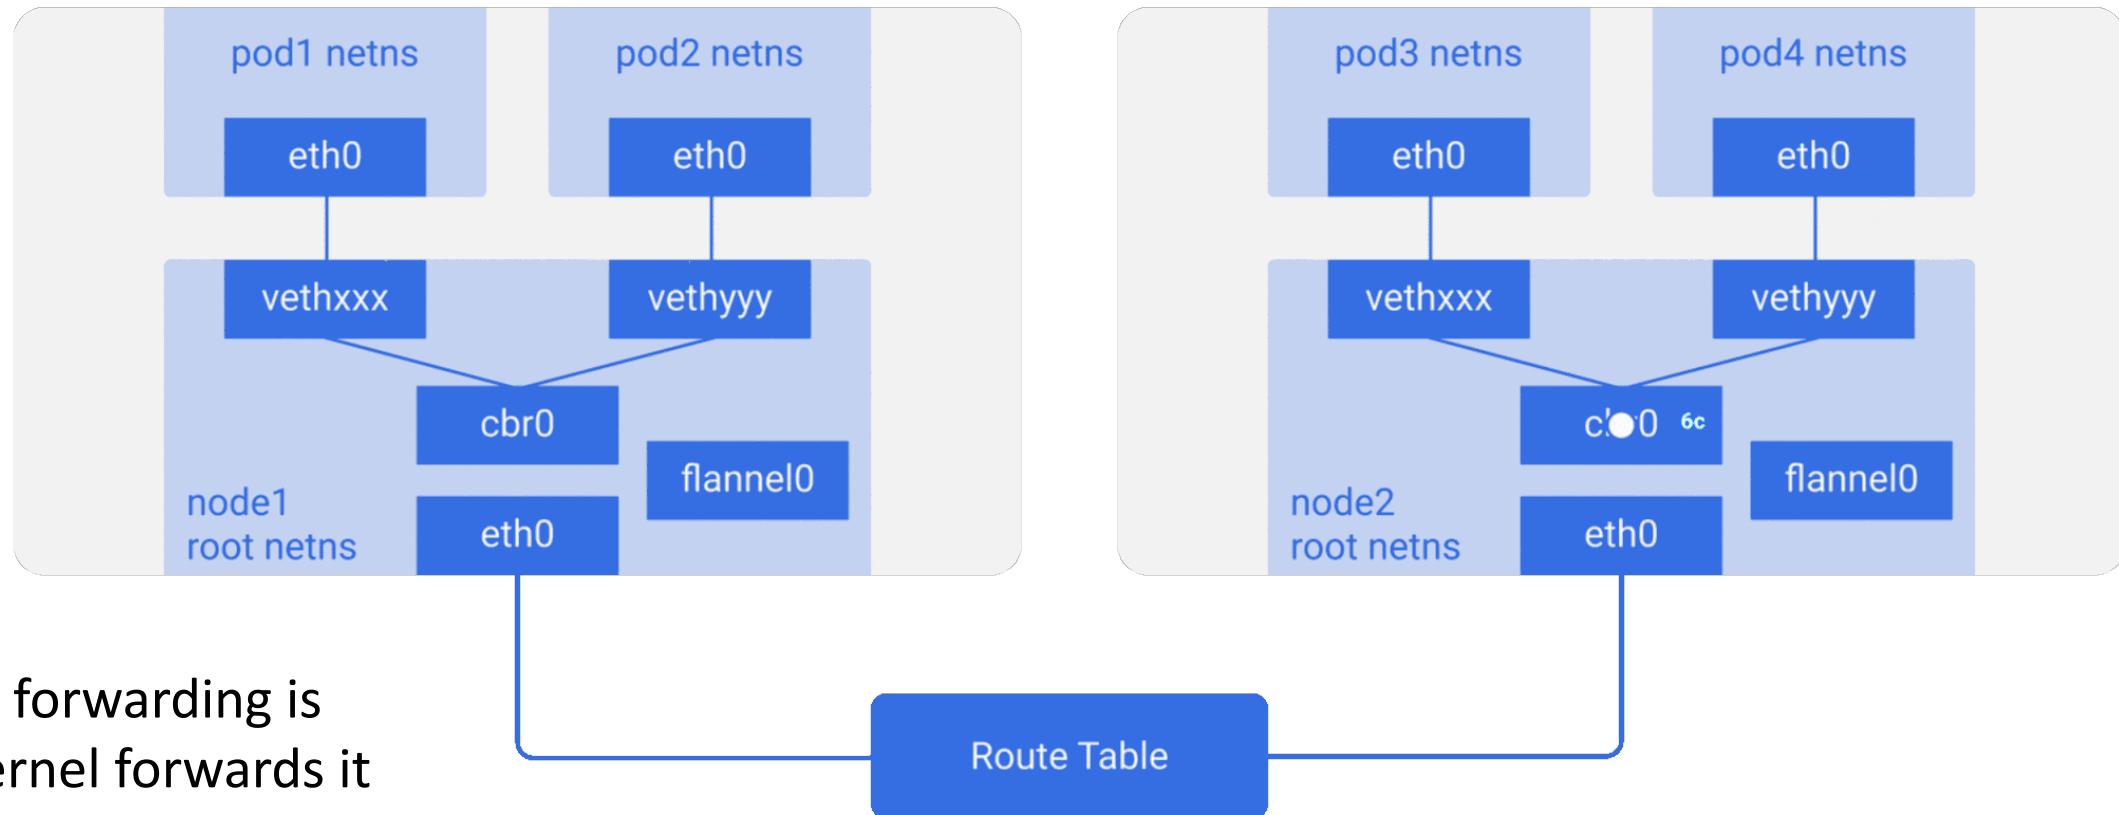


Flannel Overlay Networking



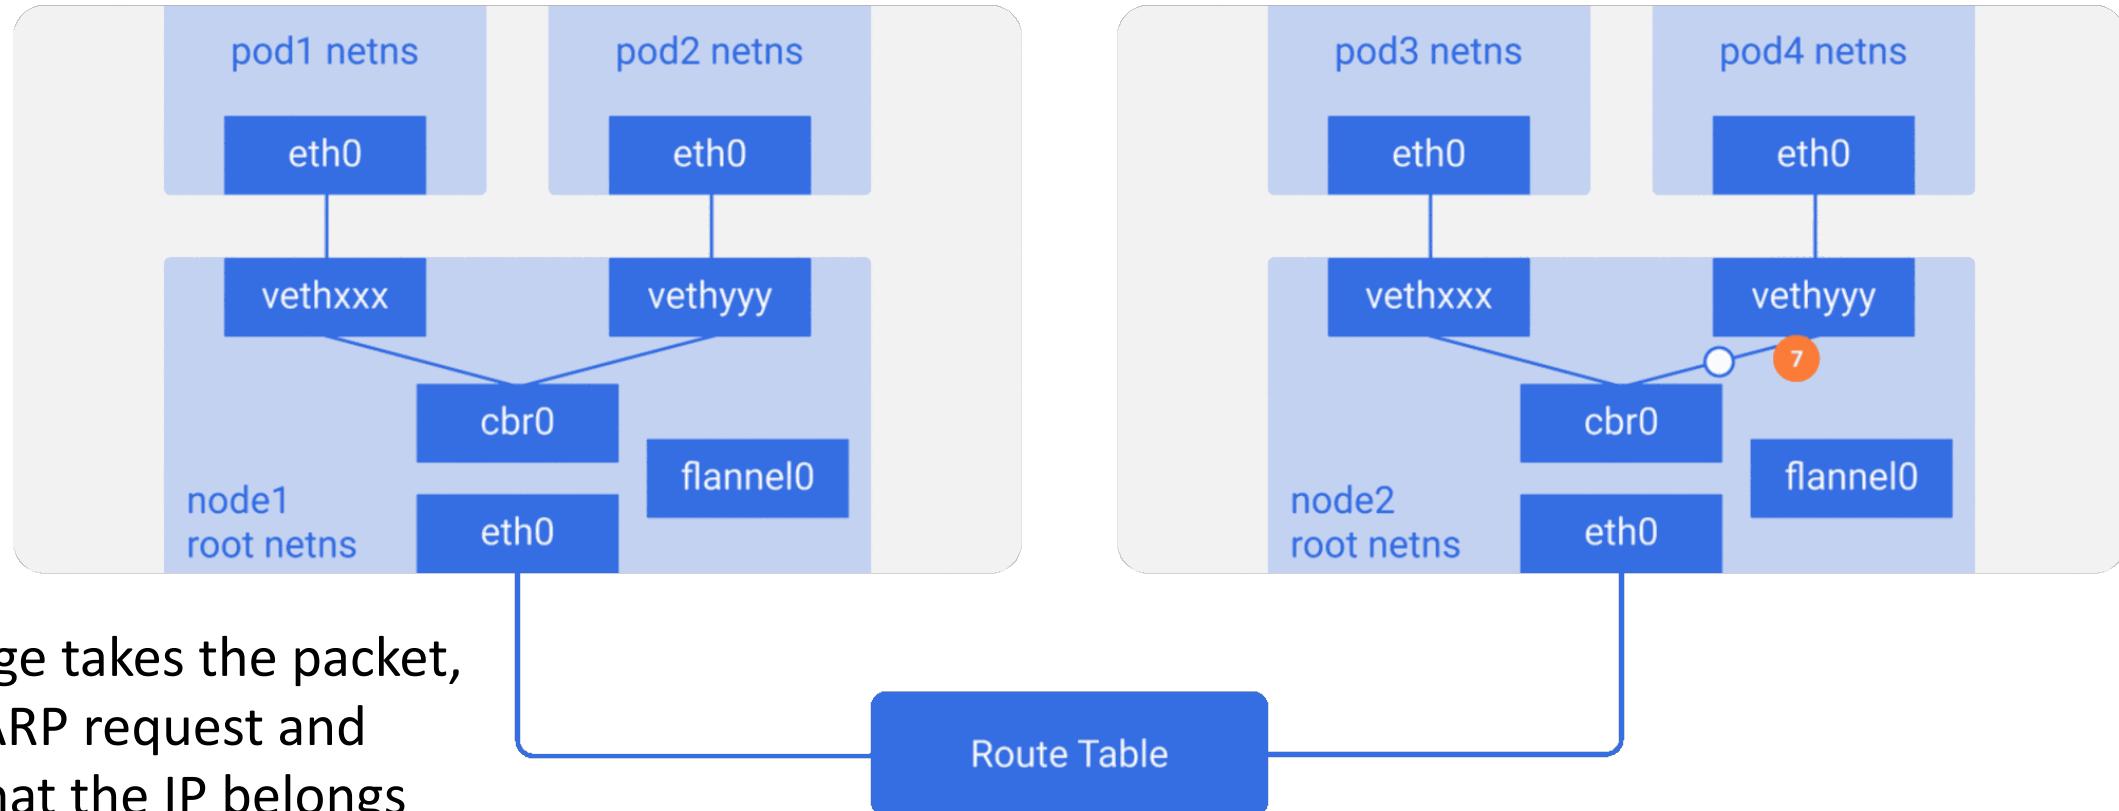
6b. flannel0 de-capsulates
and emits it back in the root
network namespace.

Flannel Overlay Networking

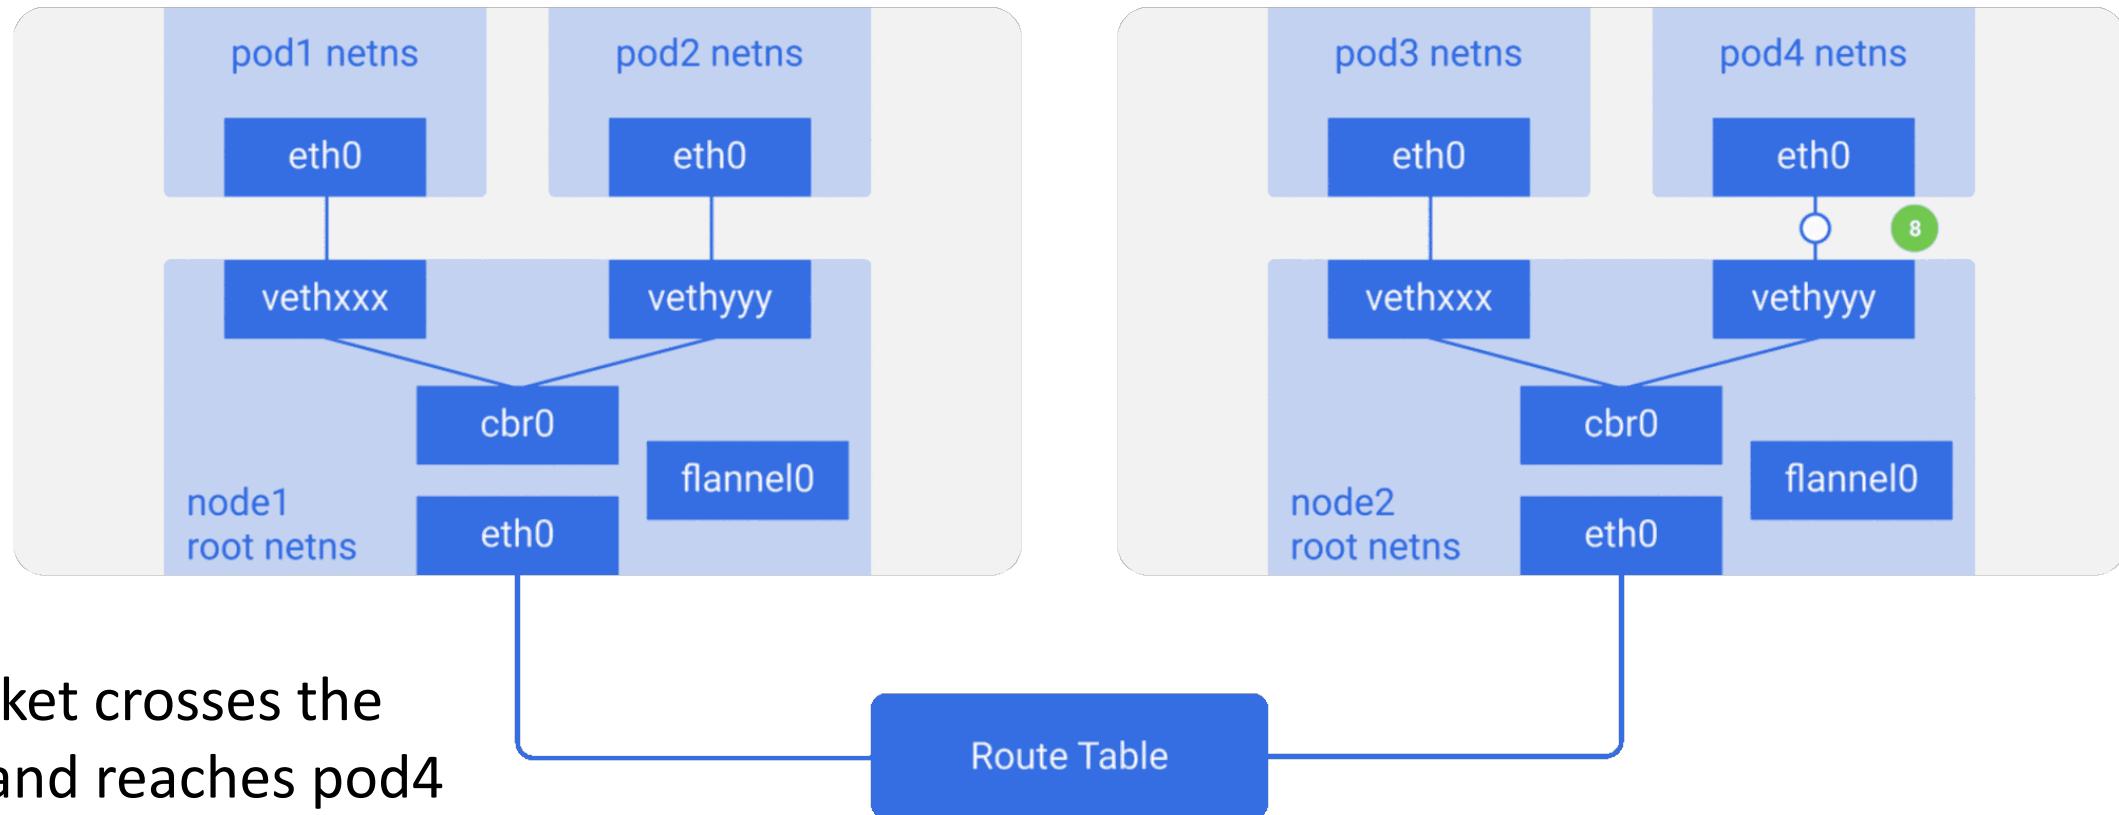


6c. Since IP forwarding is enabled, kernel forwards it to cbr0 as per the route tables.

Flannel Overlay Networking



Flannel Overlay Networking



Kubernetes Services Overview



Kubernetes Service Objects:Microservices

Services provide abstraction between layers of an application

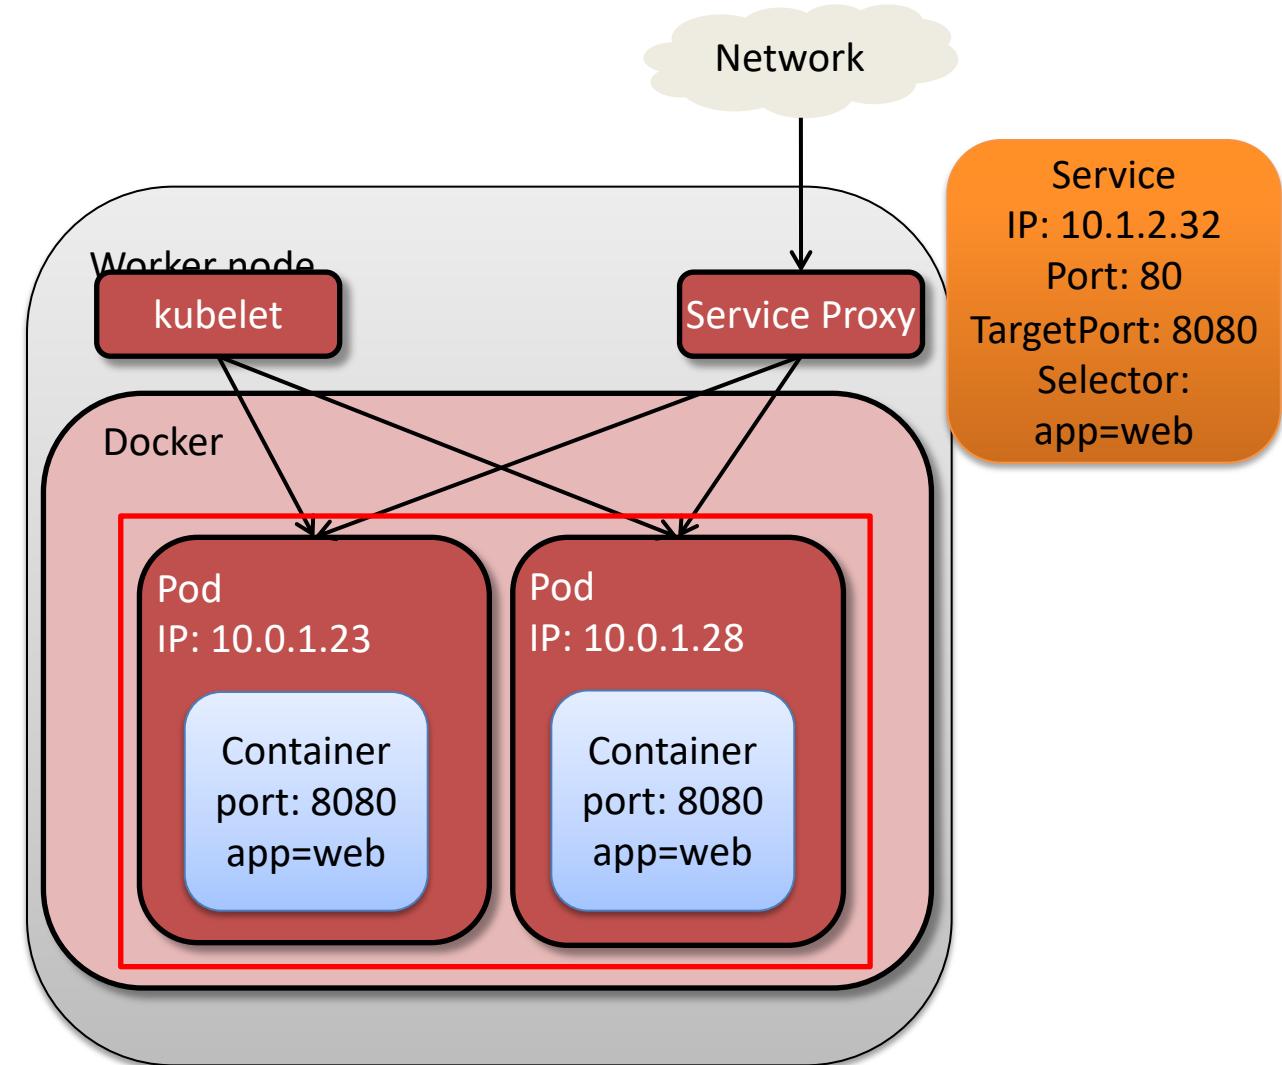
- **Service** object provides a stable IP for a collection of Pods
- Services use a label **selector** to target a specific set of pods as endpoints to receive proxied traffic
- Clients can reliably connect via the service IP and port(s), even as individual endpoint pods are dynamically created & destroyed
- Can model other types of backends using services without selectors

```
apiVersion: v1
kind: Service
metadata:
  name: wordpress
  labels:
    app: wordpress
spec:
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: wordpress
    tier: frontend
  type: LoadBalancer
```

sampleservice.yaml

Services Provide Abstraction Layer for Applications

- Services can be used for communications between application tiers
- Services can also be used to expose applications outside the K8s cluster
- Services distribute requests over the set of Pods matching the service's selector
 - Service functions as TCP and UDP proxy for traffic to Pods
 - Service maps its defined ports to listening ports on Pod endpoints



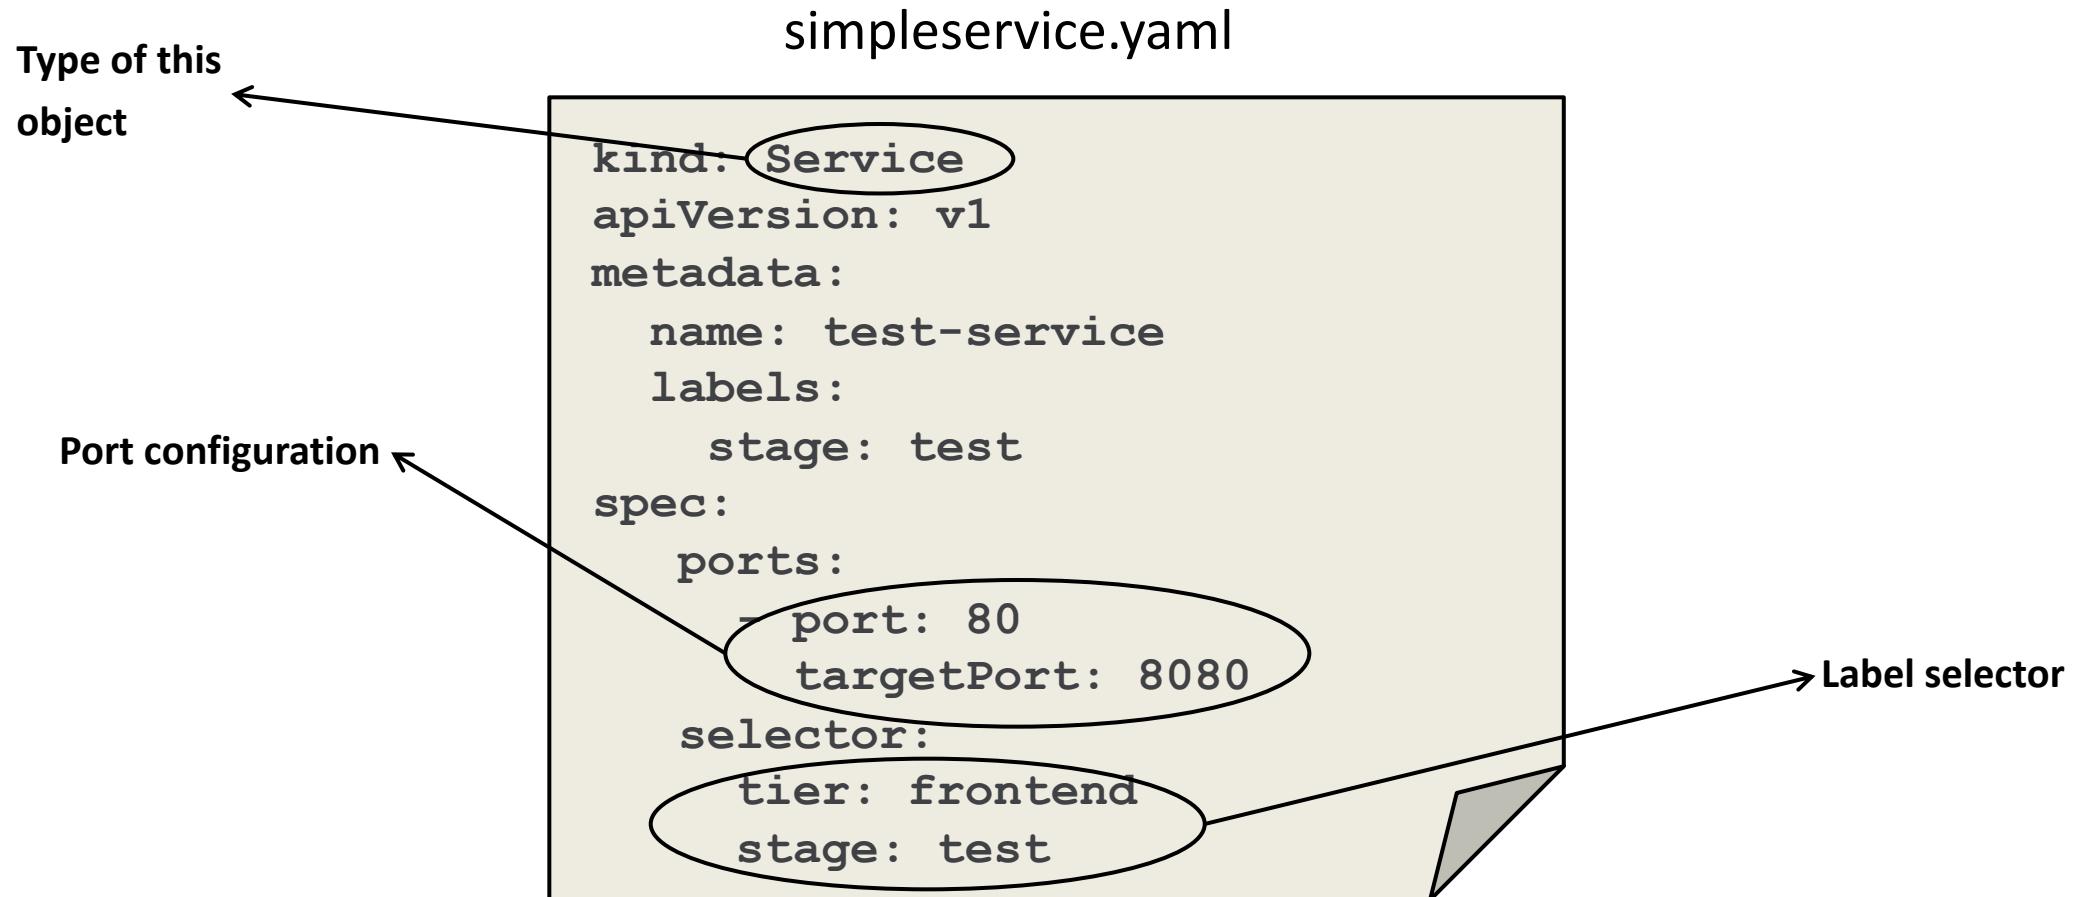
Defining a Service Using a Manifest File

Services can defined in YAML or JSON, like other K8s resources

- **kind** field value is ‘Service’
- **metadata** includes
 - **name** to assign to Service
- **spec** includes the ports associated with the Service
 - **port** is the Service’s port value
 - **targetPort** is connection port on selected pods (default: **port** value)
- **selector** specifies a set of label KV pairs to identify the endpoint pods for the Service

```
kind: Service
apiVersion: v1
metadata:
  name: test-service
  labels:
    stage: test
spec:
  ports:
    - port: 80
      targetPort: 8080
  selector:
    tier: frontend
    stage: test
```

Reviewing a Service Manifest File



ServiceTypes and Exposing Applications

- By default, a Service is assigned a cluster-internal IP – good for back-ends
 - **ServiceType ClusterIP**
- To make a front-end Service accessible outside the cluster, there are other ServiceTypes available
 - **ServiceType NodePort** exposes Service on each Node's IP on a static port
 - **ServiceType LoadBalancer** exposes the Service externally using cloud provider's load balancer
- Can also use a Service to expose an external resource via **ServiceType ExternalName**

```
kind: Service
apiVersion: v1
metadata:
  name: test-service
  labels:
    stage: test
spec:
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30080
  selector:
    tier: frontend
  type: NodePort
```

Exposing Services at L7 through Ingresses

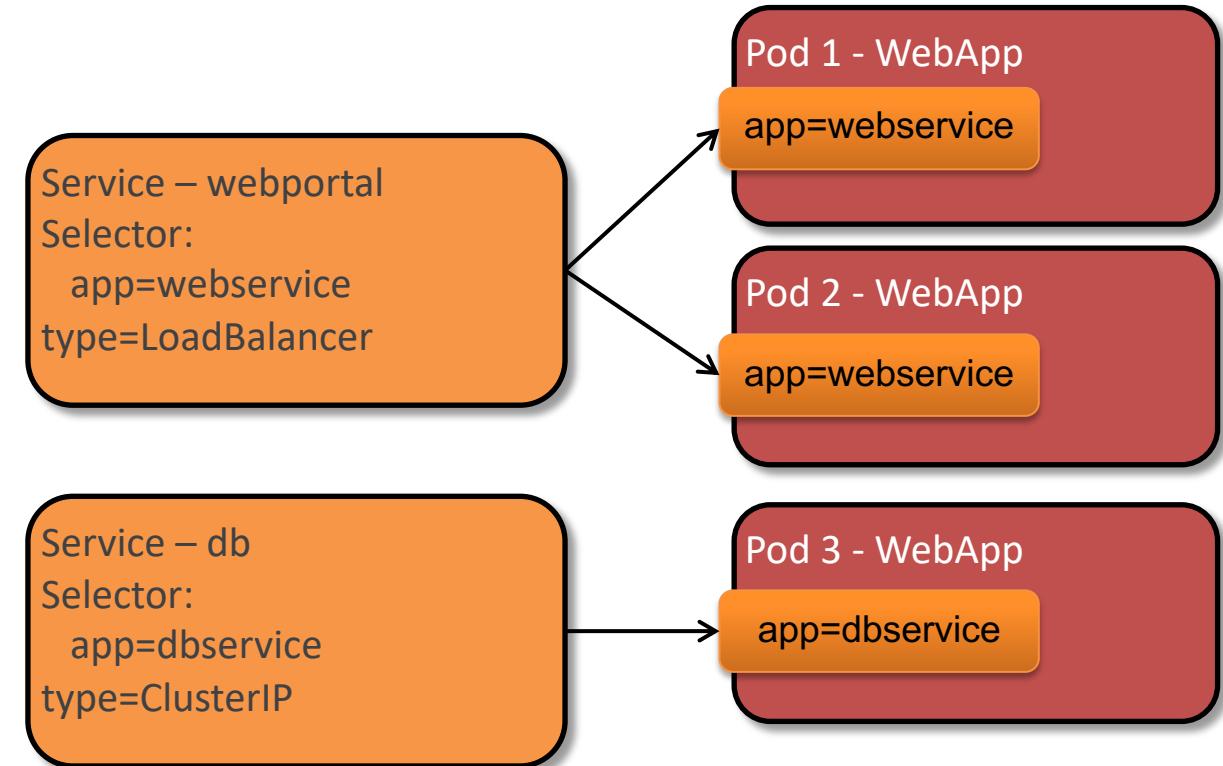
- Kubernetes also provides the facility to define **Ingress** resource to configure an external loadbalancer at L7
- Spec of Ingress resource is a set of rules matching HTTP host/url paths to specific Service backends
- Ingresses require the cluster to be running an appropriately configured Ingress controller to function (e.g. nginx)
- Useful for implementing fanout, Service backends for virtual hosts, etc.

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: test-ingress
spec:
  rules:
  - host: bar.foo.com
    http:
      paths:
      - path: /first
        backend:
          serviceName:
            firstservice
            servicePort: 80
      - path: /second
        backend:
          serviceName:
            secondservice
            servicePort: 80
```

Selecting Pods as Service Endpoints

Service's pod selector based on labels

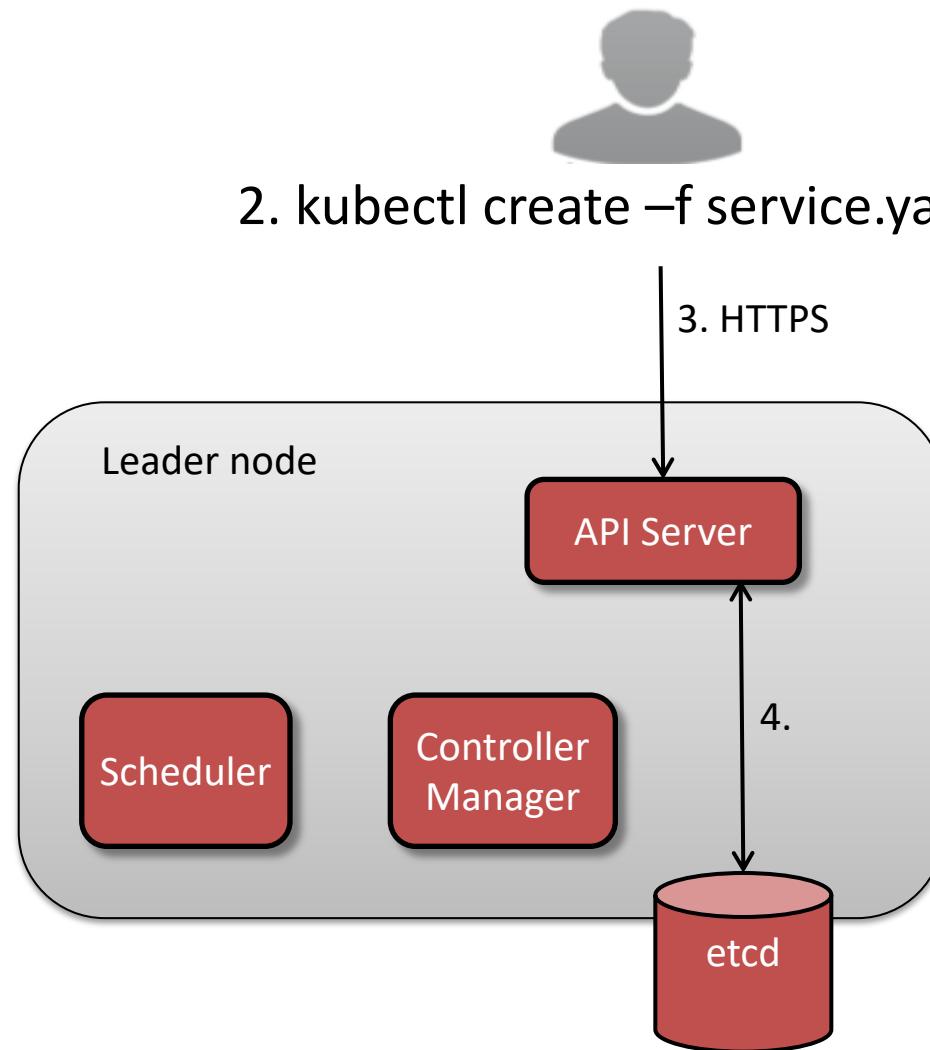
- Multiple pods can have the same label, unlike pod names which are unique in the namespace
- K8s system re-evaluates Service's selector continuously
- K8s maintains **endpoints** object of same name with list of pod IP:port's matching Service's selector



Services Management

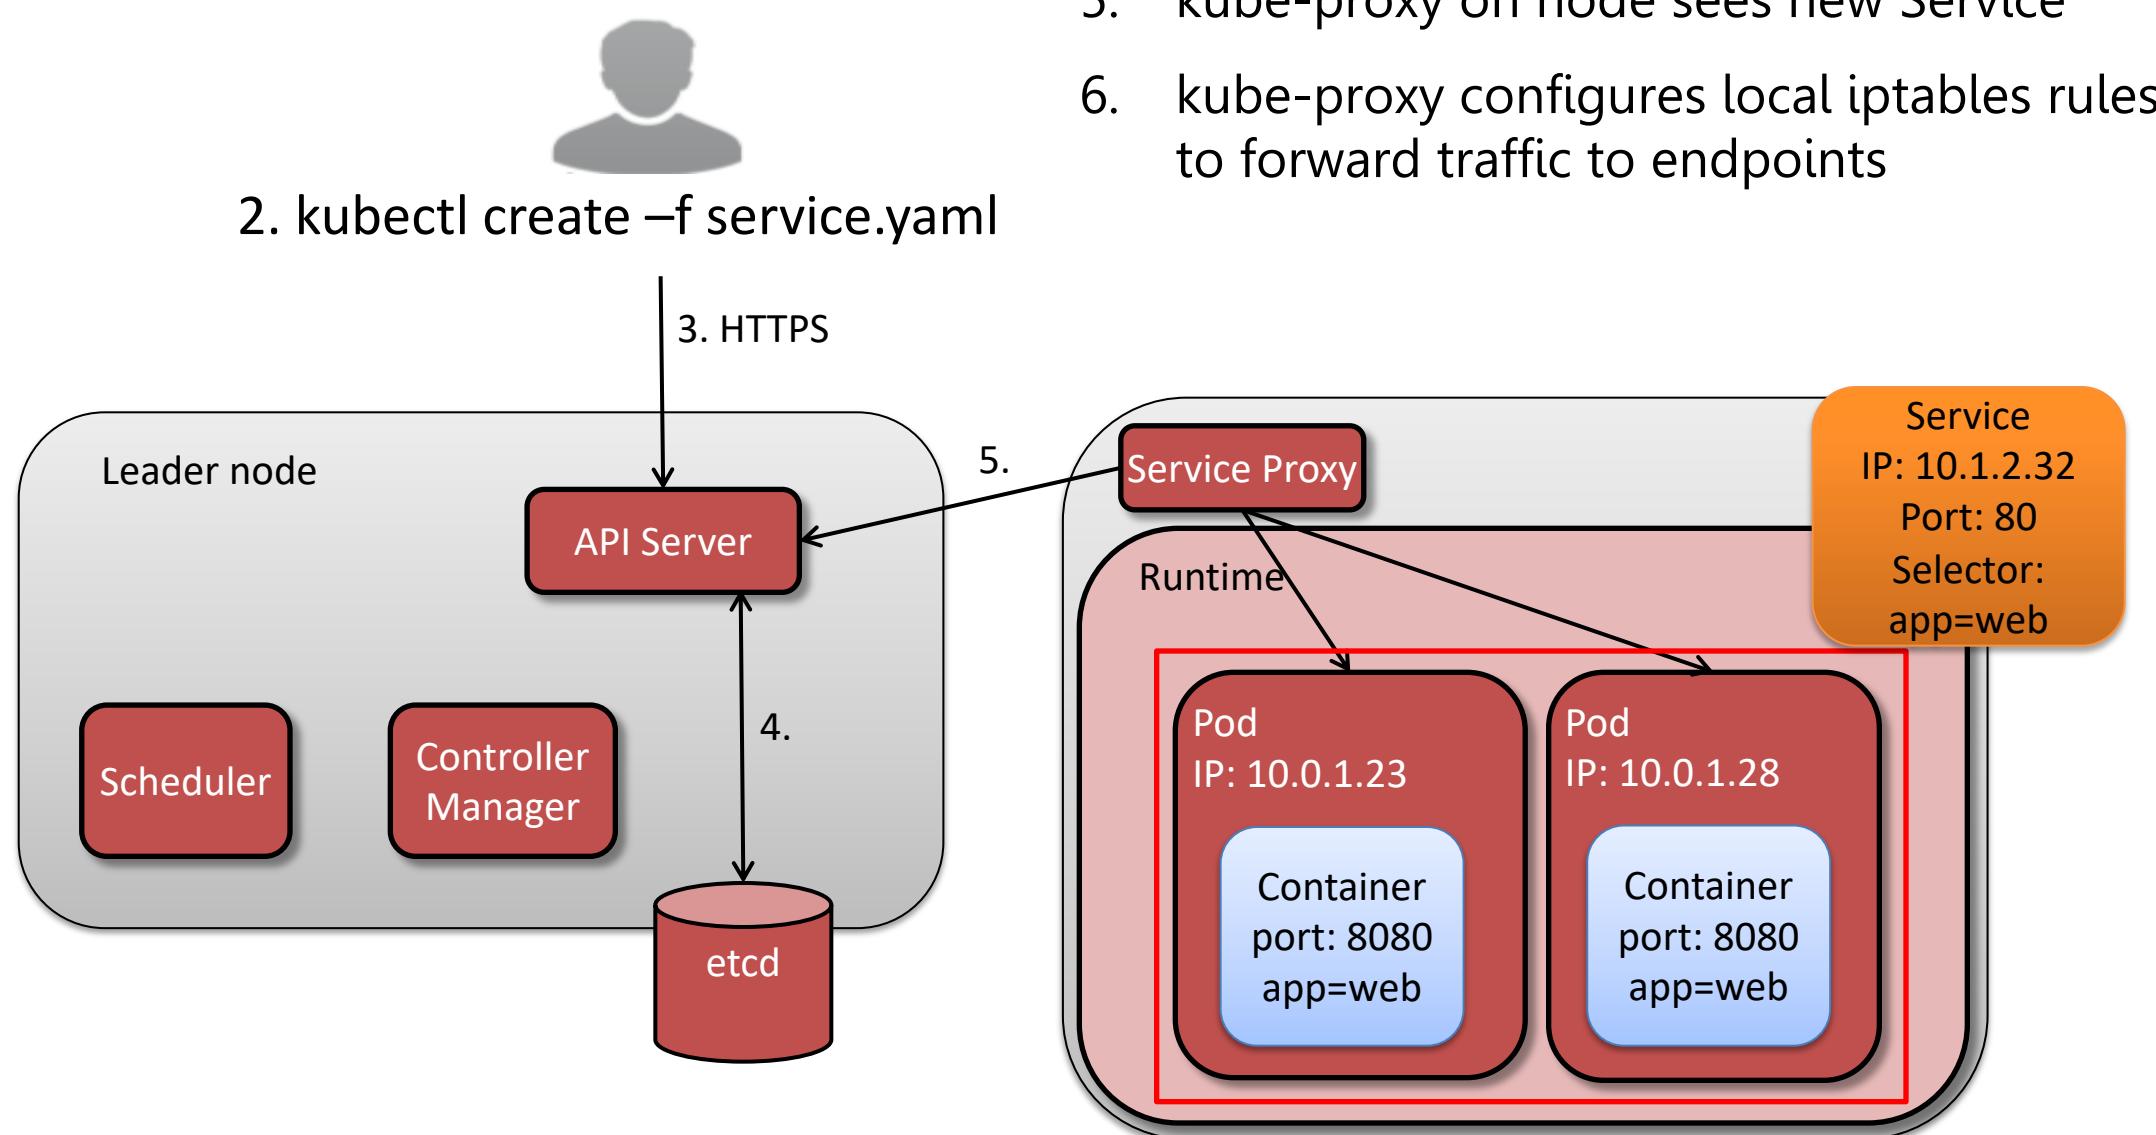


Service Creation Process



1. User writes a Service manifest file
2. User requests creation of Service from manifest via CLI
3. CLI tool marshals parameters into K8s RESTful API request (HTTP POST)
4. kube-apiserver creates new Service object record in etcd

Service Creation Process



Accessing Services Externally

Checking the service external IP

- Configured as type LoadBalancer, a configured cluster will provide an externally accessible IP for your Service

```
$ kubectl get services test-service
NAME      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
test-service  10.3.247.123  104.154.105.198  8080/TCP  4h
```

Deleting Services

Services can be deleted anytime

- Service deletion does not affect pods targeted by the Service's selector
- Can be done referencing the Service name or a manifest for the resource

```
$ kubectl delete -f simpleservice.yaml
```

```
$ kubectl delete service test-service
```

Service Discovery via DNS

Kubernetes advertises services via cluster DNS

- Kubernetes uses a cluster-internal DNS add-on to create and manage records for all Services in the cluster
- Pod's DNS search list includes its own namespace and cluster default domain by default

```
$ kubectl get svc
NAME           CLUSTER-IP      EXTERNAL-IP    PORT(S)        AGE
kubernetes     10.0.0.1        <none>         443/TCP       11d
test-service   10.0.0.102      <nodes>        8080:30464/TCP 2d
```

```
kubectl exec -ti busybox1 -- nslookup test-service.default
Server:  10.0.0.10
Address 1: 10.0.0.10 coredns.kube-system.svc.cluster.local
```

```
Name:      test-service.default
Address 1: 10.0.0.102 test-service.default.svc.cluster.local
```

Deployments



What is a Deployment?

Kubernetes controller optimal for stateless applications

- Deployments allow you to declaratively manage pods, including replication
- Deployments support
 - Creating, rolling out, and rolling back changes to homogeneous set of pods
 - Scaling set of pods out and back declaratively
- Deployments include
 - Implicit Replica Set controller to handle pod replicas
 - Template spec of pods to be created and managed – no need to separately create pods
- Deployments used for web applications, mobile back-ends, API's

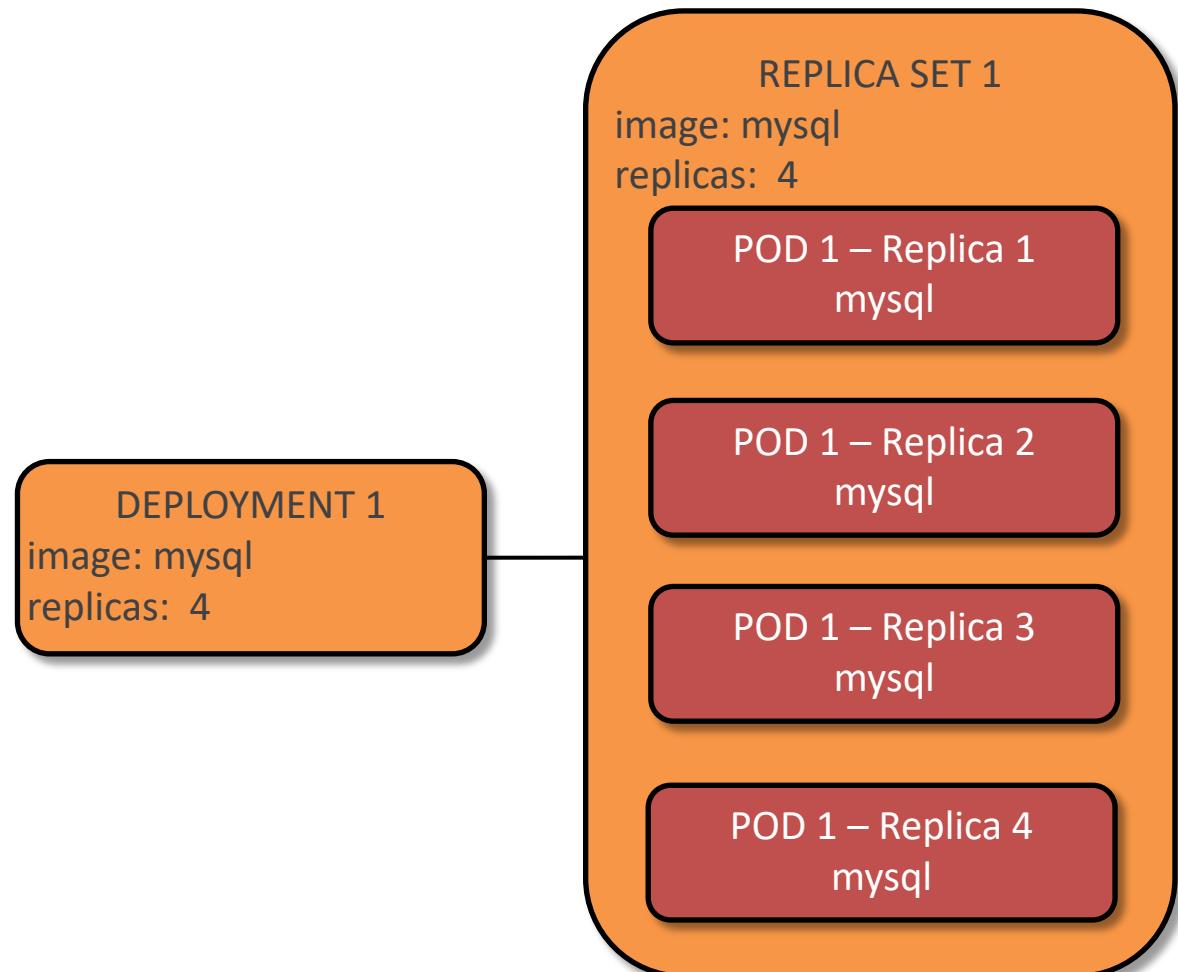
What if my Application isn't Stateless?

- Kubernetes provides other controller objects for applications that need different deployment schemes
- **StatefulSets** (previously PetSets) control deployment of pods for applications that need more stable deployment contexts
 - Pods in StatefulSets have unique ordinal, stable network identity and stable storage using persistent volumes
 - When pods are deployed, they are created in sequence of ordinals 1..N
 - Pod N must be running and ready before Pod N+1 is deployed
 - When pods are destroyed, they are terminated in reverse sequence N..1
- **DaemonSets** ensure that a replica of a specified pod is running on every node (or every selected node) in the cluster
- **Jobs** manage sets of pods where N must run to successful completion

Deployments Control ReplicaSet Controllers

Definition of how many replicated Pods should exist

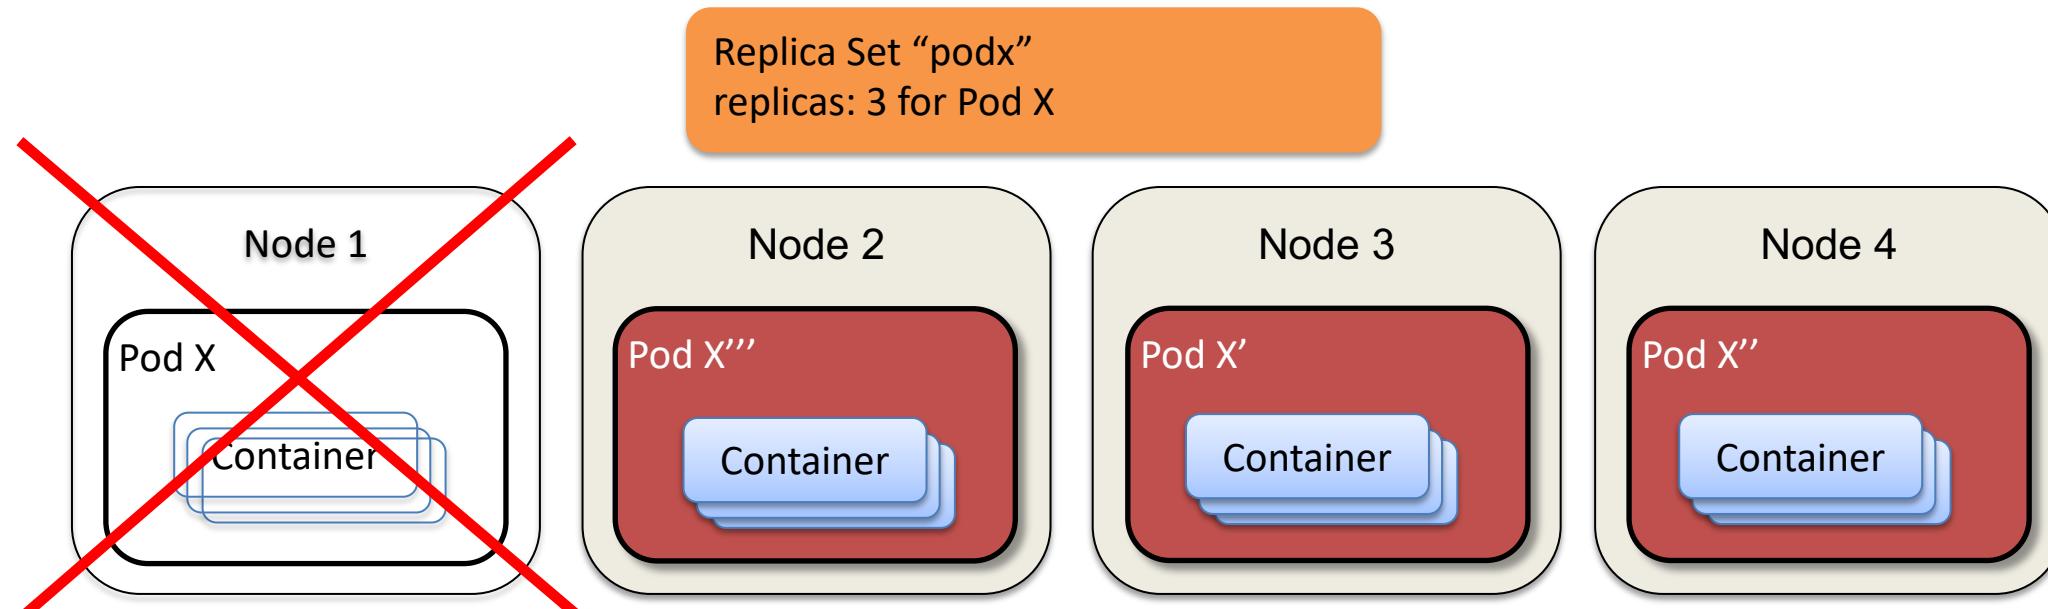
- Deployment creates and manages a Replica Set that manages a set of pods
- Replica count can be adjusted as needed to scale the Replica Set out and back
- Replica Set successor to the ReplicationController object



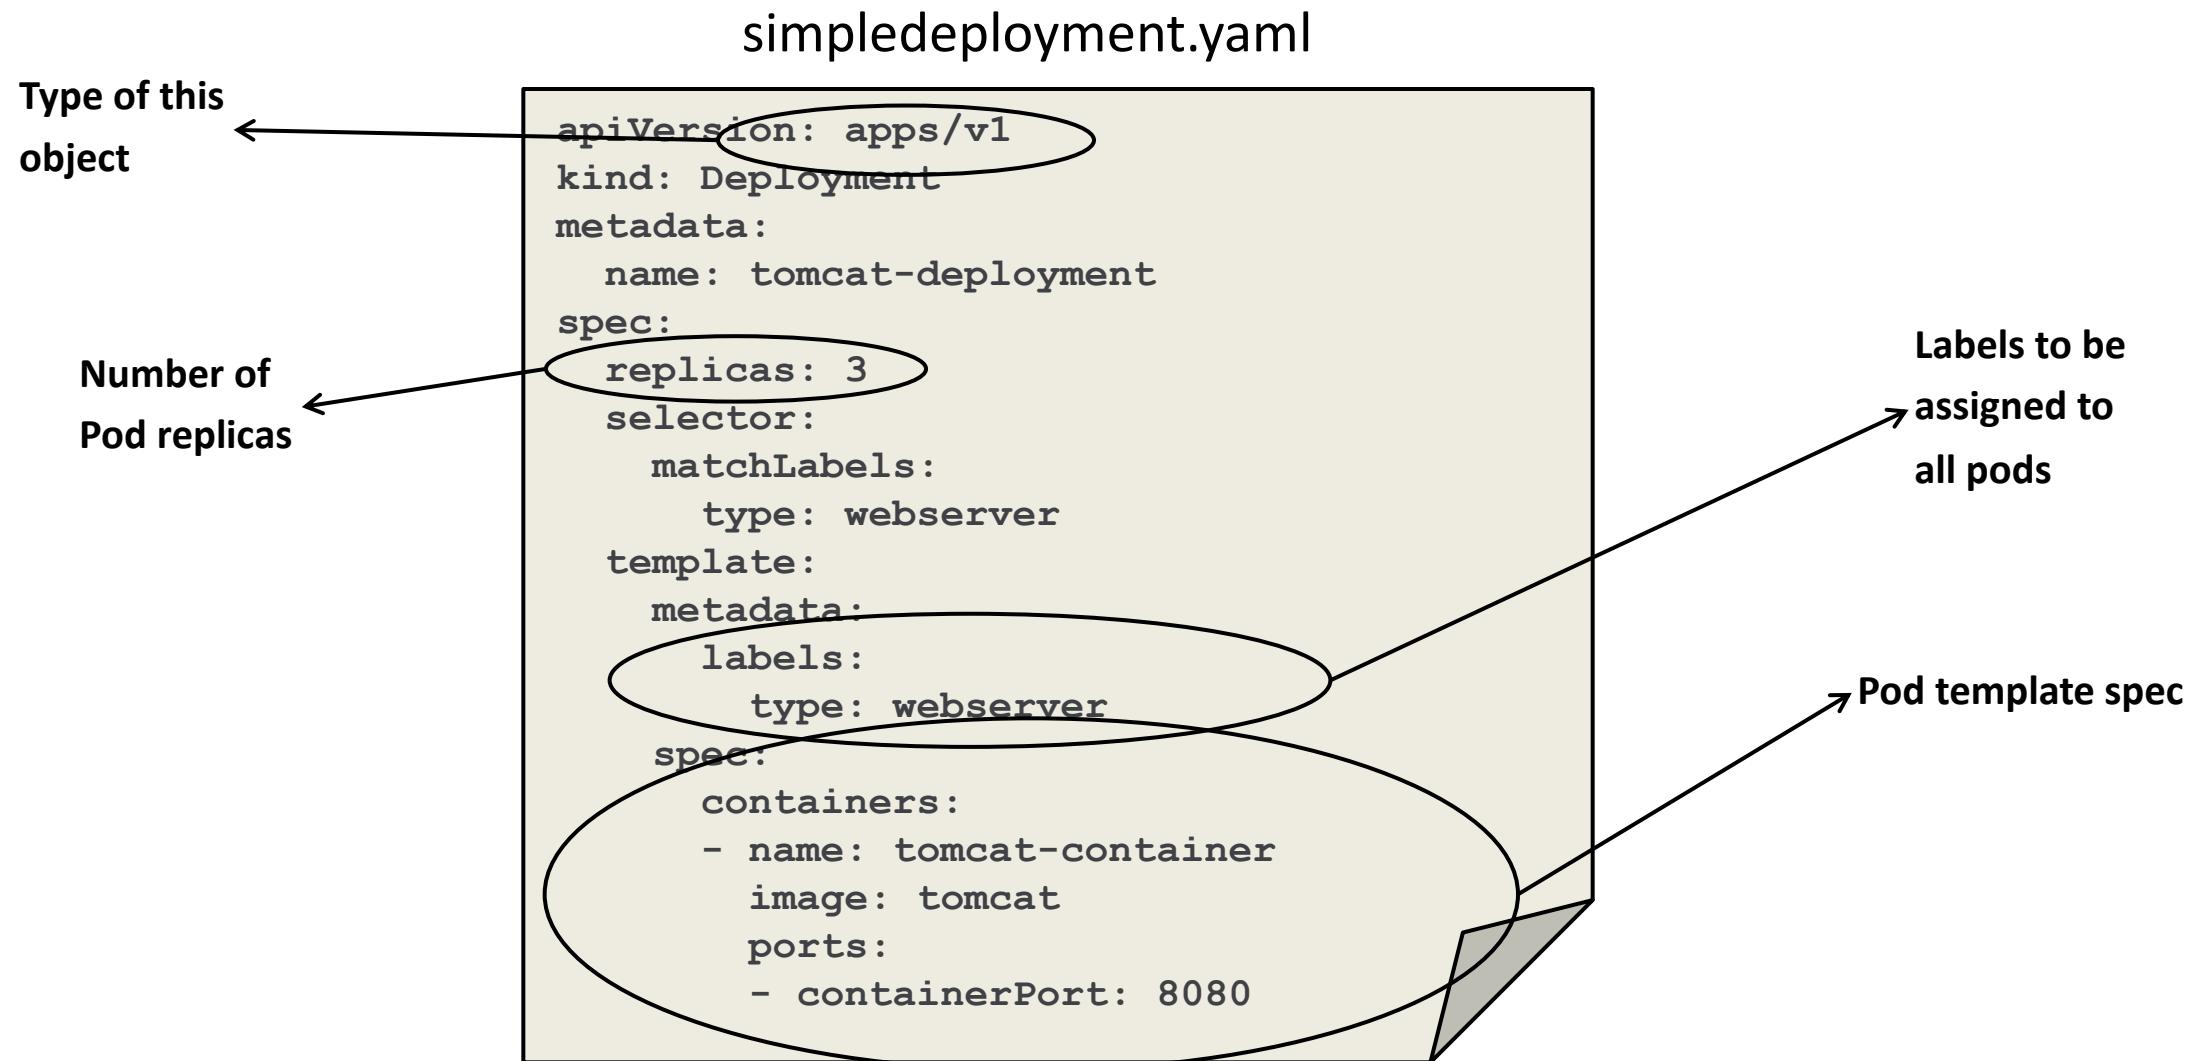
What is a Replica Set?

Provides scaling and high availability

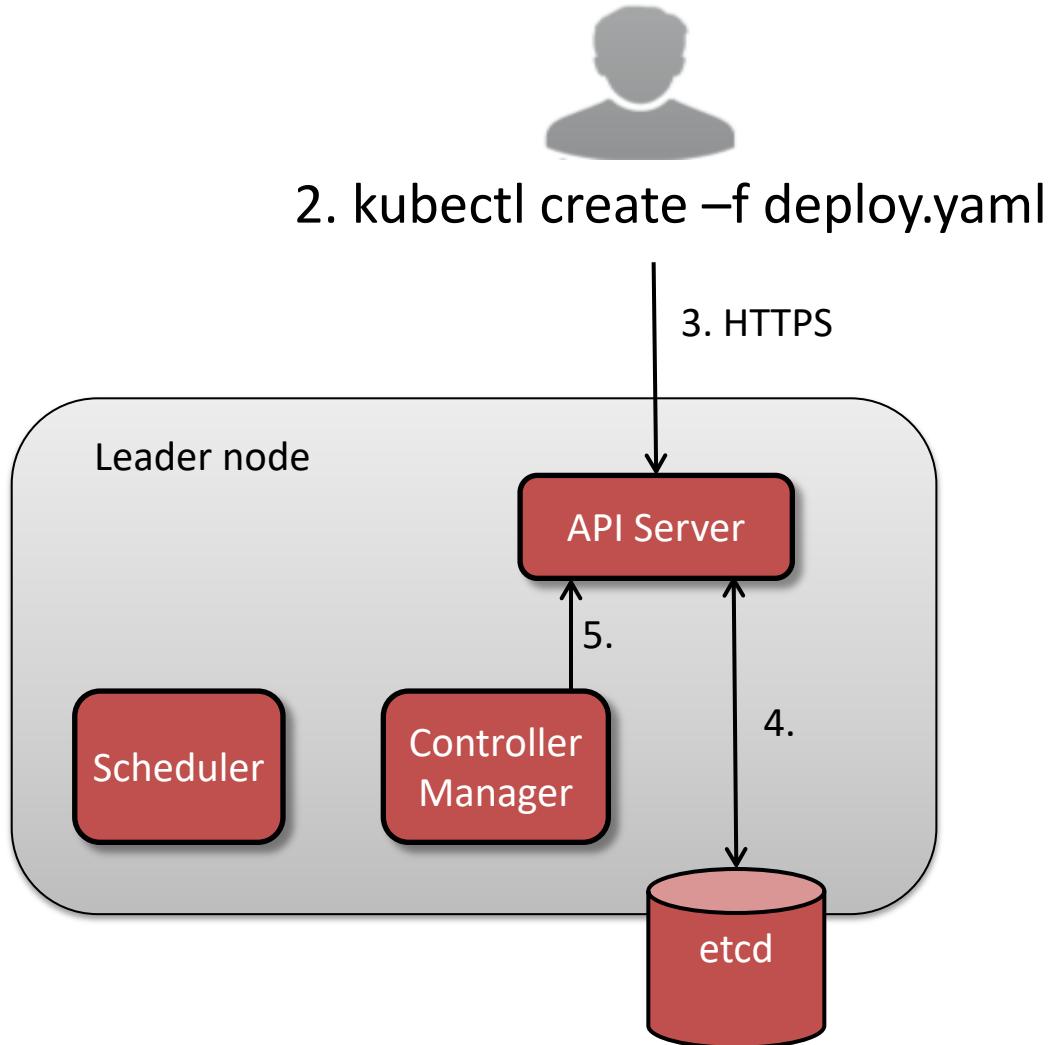
- Replica count can be changed to provide scaling on demand as needed
- If the node hosting a pod fails, the Kubernetes cluster will recreate the pod elsewhere to achieve the target number of replicas



Examining a Deployment Manifest File

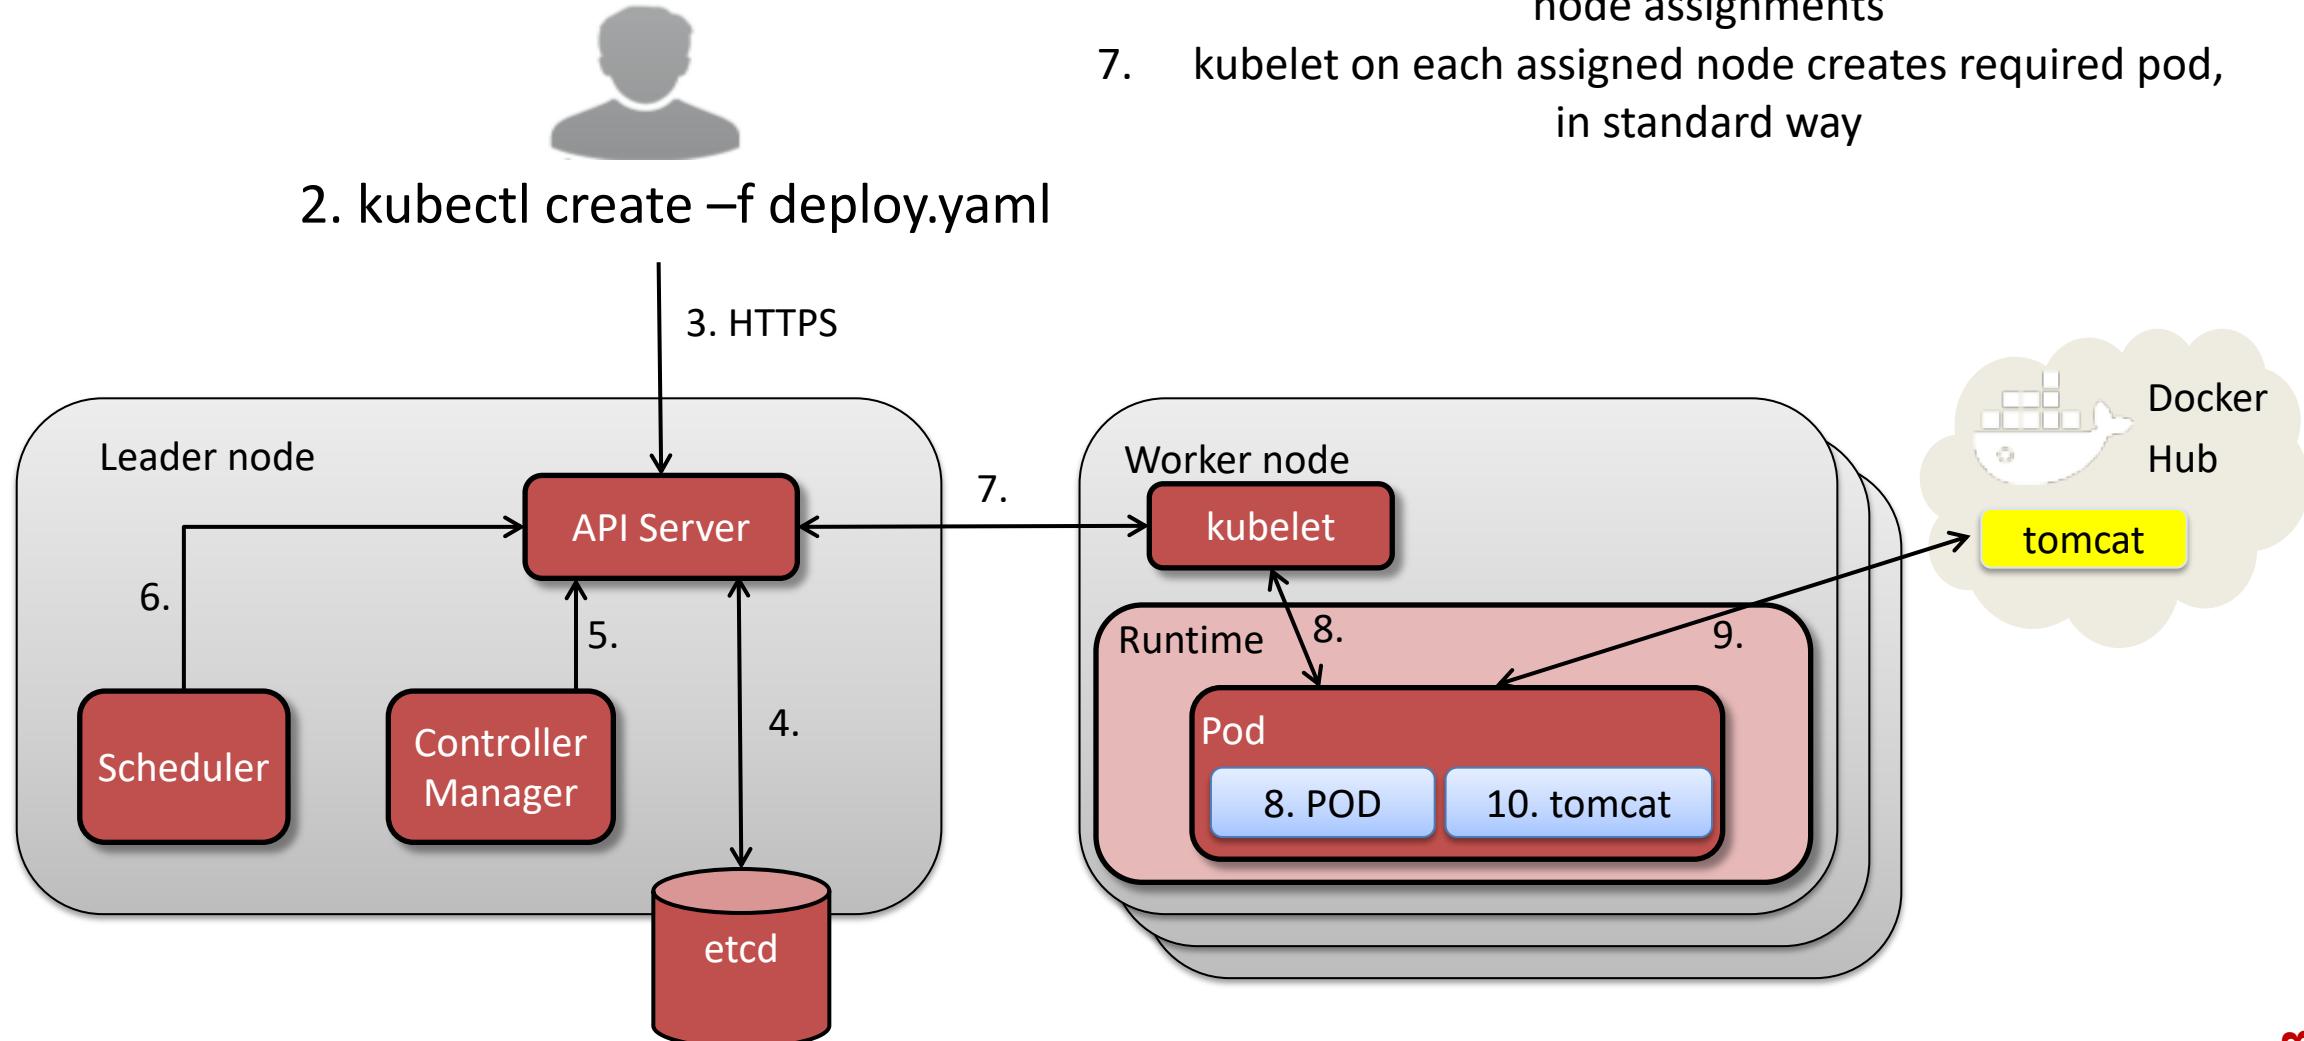


Deployment Creation Process



1. User writes a deployment manifest file
2. User requests creation of deployment from manifest via CLI
3. CLI tool marshals parameters into K8s RESTful API request (HTTP POST)
4. kube-apiserver creates new deployment object record in etcd, and new Replica Set object
5. kube-controller-manager sees new Replica Set and
 - a. Evaluates state of existing vs. required replicas
 - b. Submits pod creation requests to API to create required number of replicas

Deployment Creation Process



Lab: Deployments



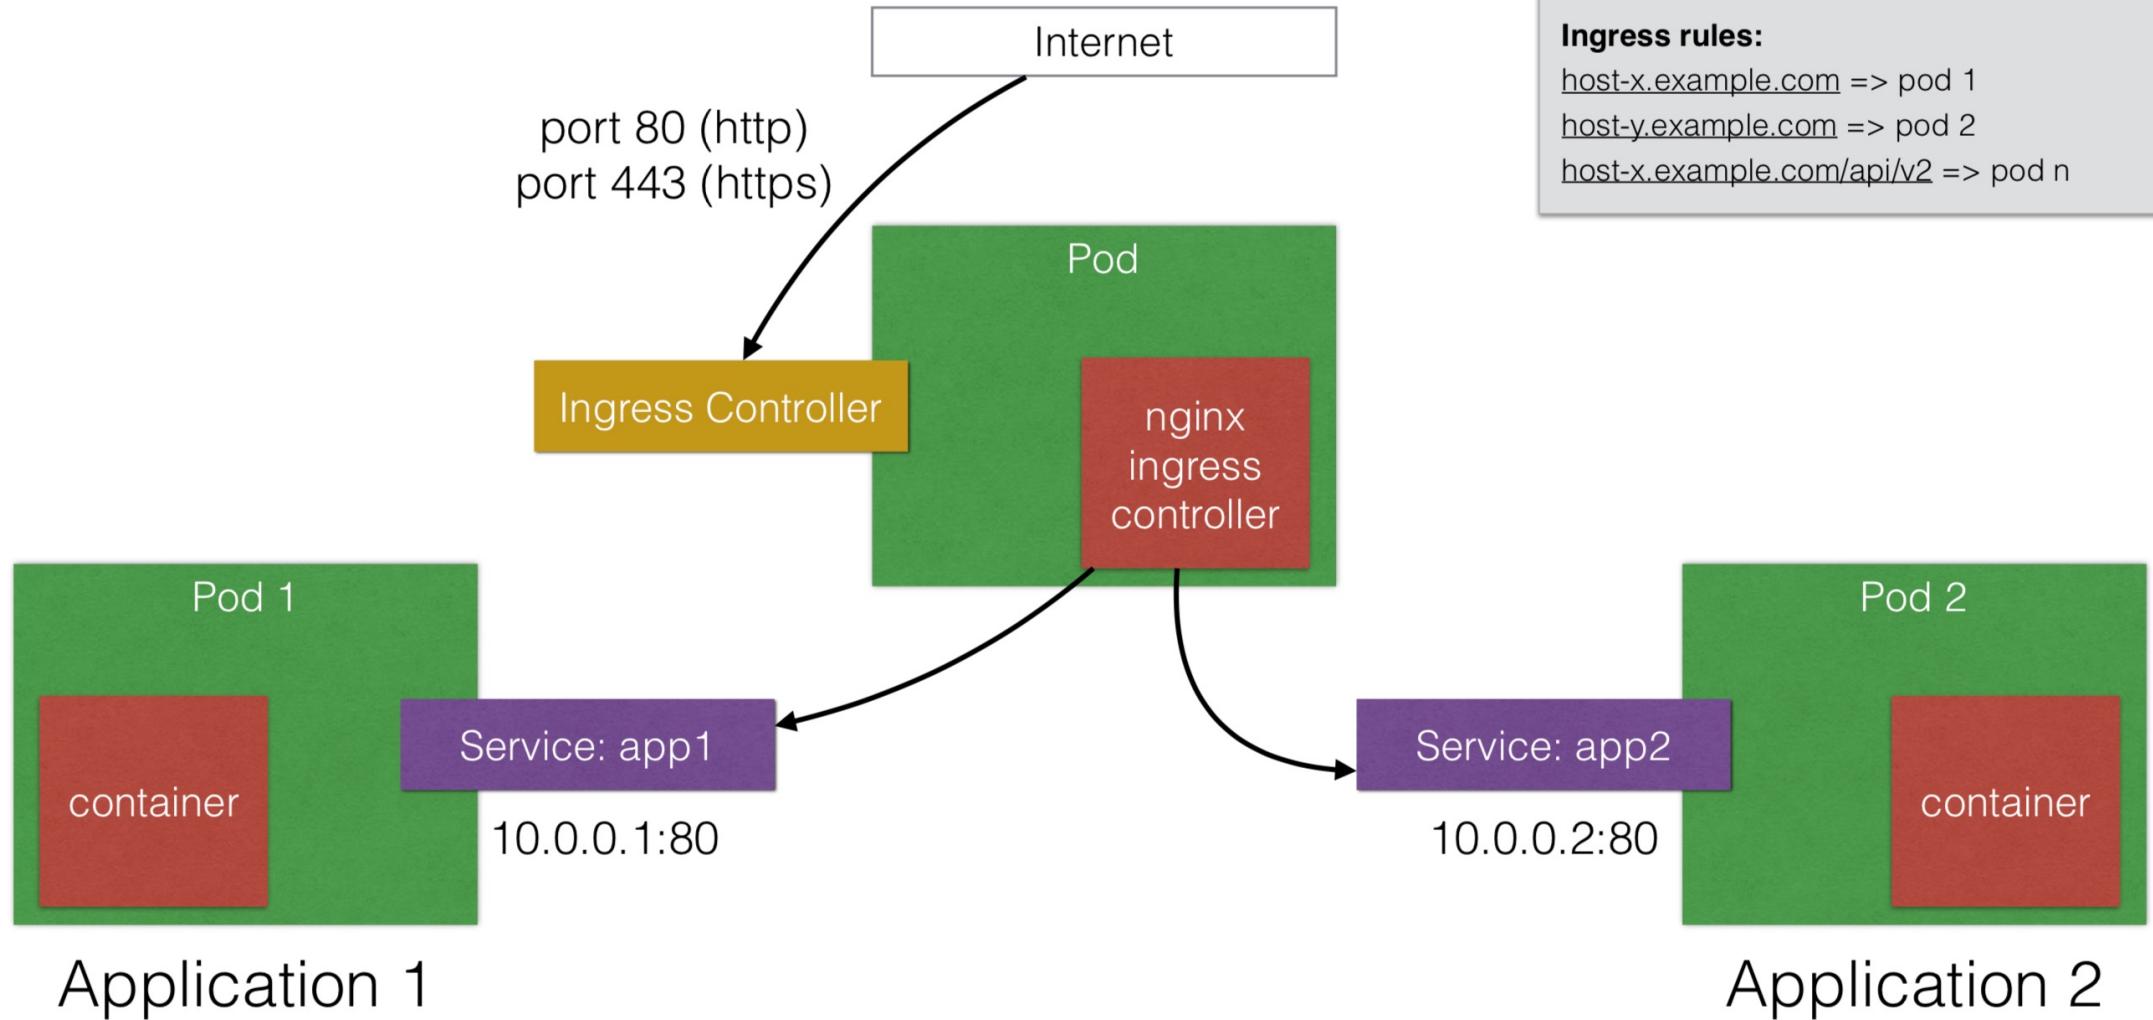
Kubernetes Ingress Overview



Ingress

- L7 (instead of L4)
- HTTP (instead of tcp/udp)
- Ingress can route paths differently:
 - /api/app1 -> app1 pods
 - /api/app2 -> app2 pods
- Requires ingress provider (might be cloud, might be in-cluster)

Ingress



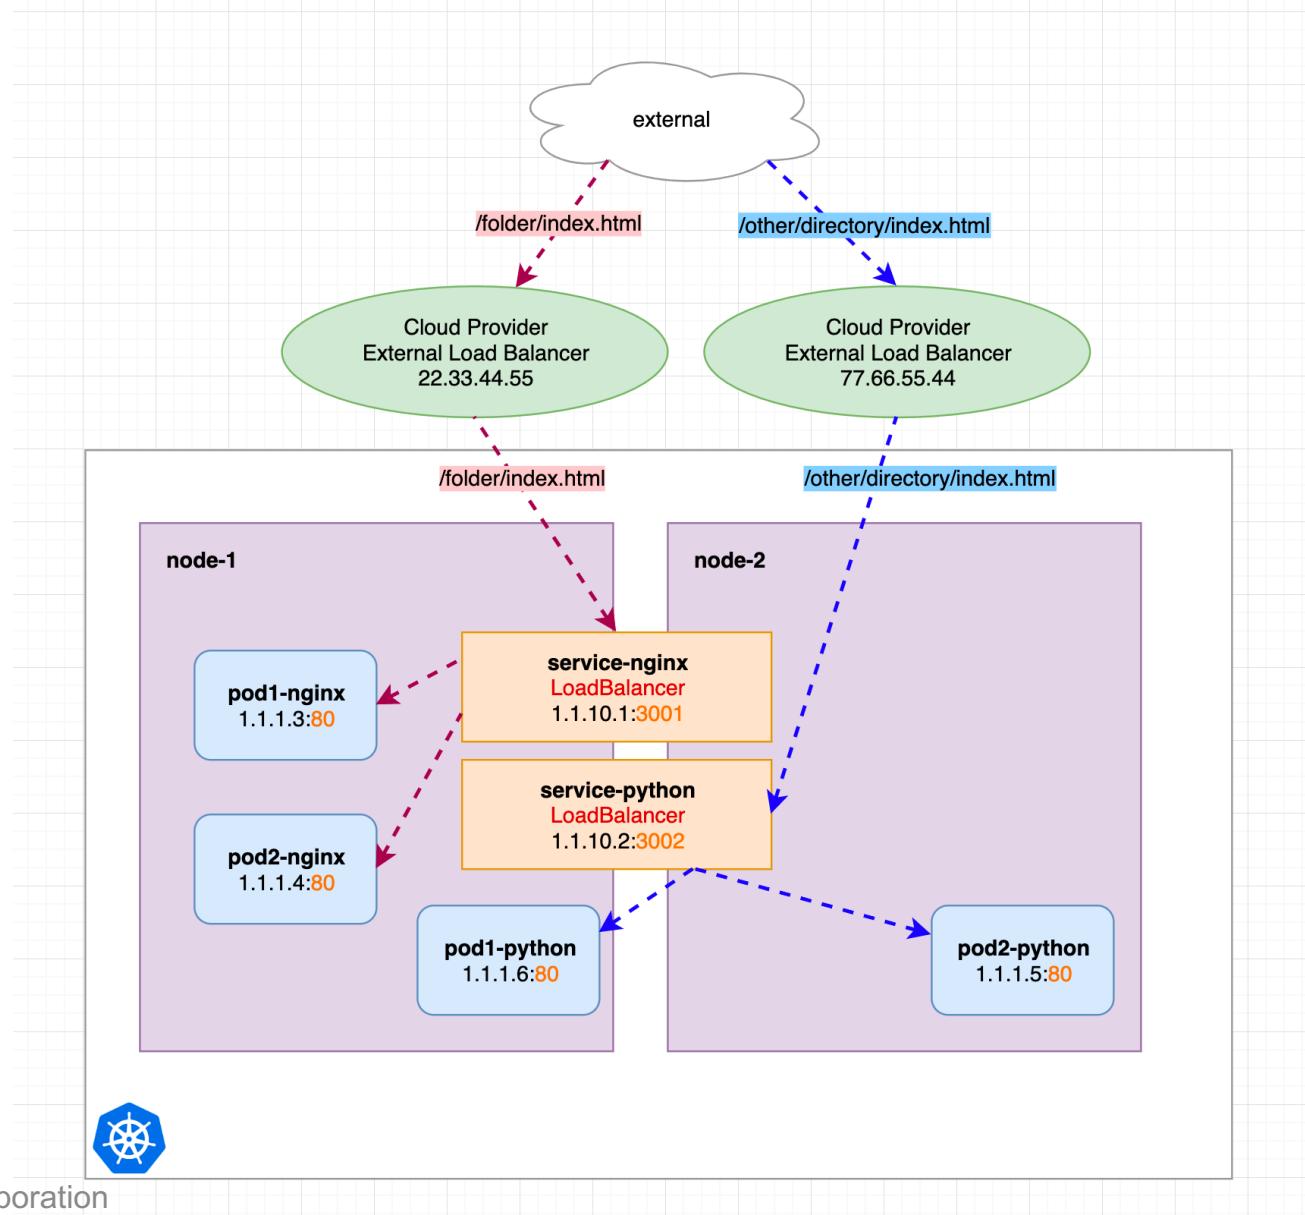
Ingress Definition

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: helloworld-rules
spec:
  rules:
    - host: helloworld-v1.example.com
      http:
        paths:
          - path: /
            backend:
              serviceName: helloworld-v1
              servicePort: 80
    - host: helloworld-v2.example.com
      http:
        paths:
          - path: /
            backend:
              serviceName: helloworld-v2
              servicePort: 80
```

Ingress

- Create Ingress in same namespace as service
- Cannot route traffic to a service in a different namespace where you don't have the ingress object.
- An Ingress controller is required
- Support HTTPS and can terminate at the Ingress controller
- Each cloud has their own implementation of Ingress controller.

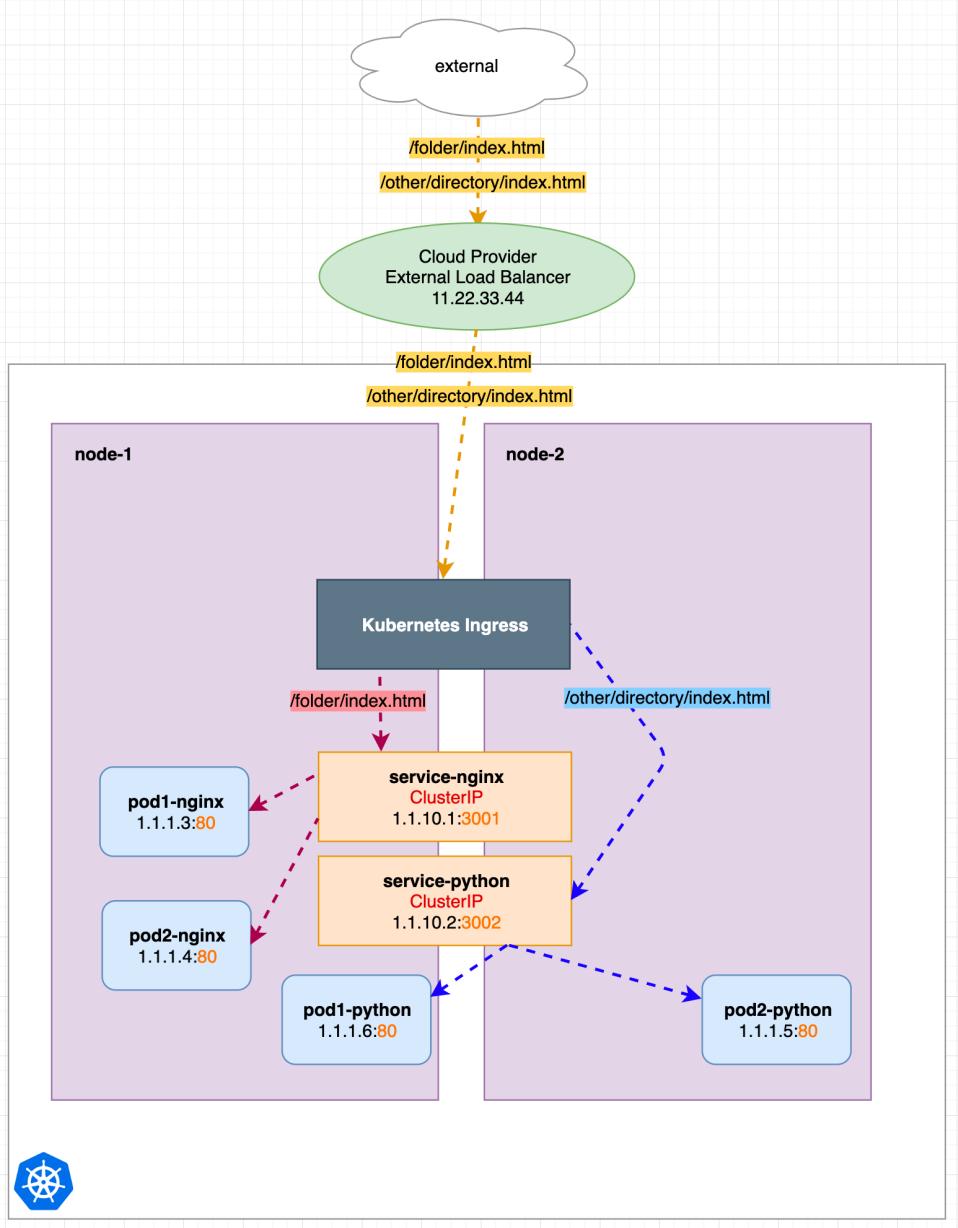
Service Type = Load Balancer



Load Balancer challenges

- Two Load Balancers requires 2 public IPs
- Load Balancers cost money
- The more services we have the more LBs required, which leads to increased bill.

Ingress controller



Ingress controller options



Nginx



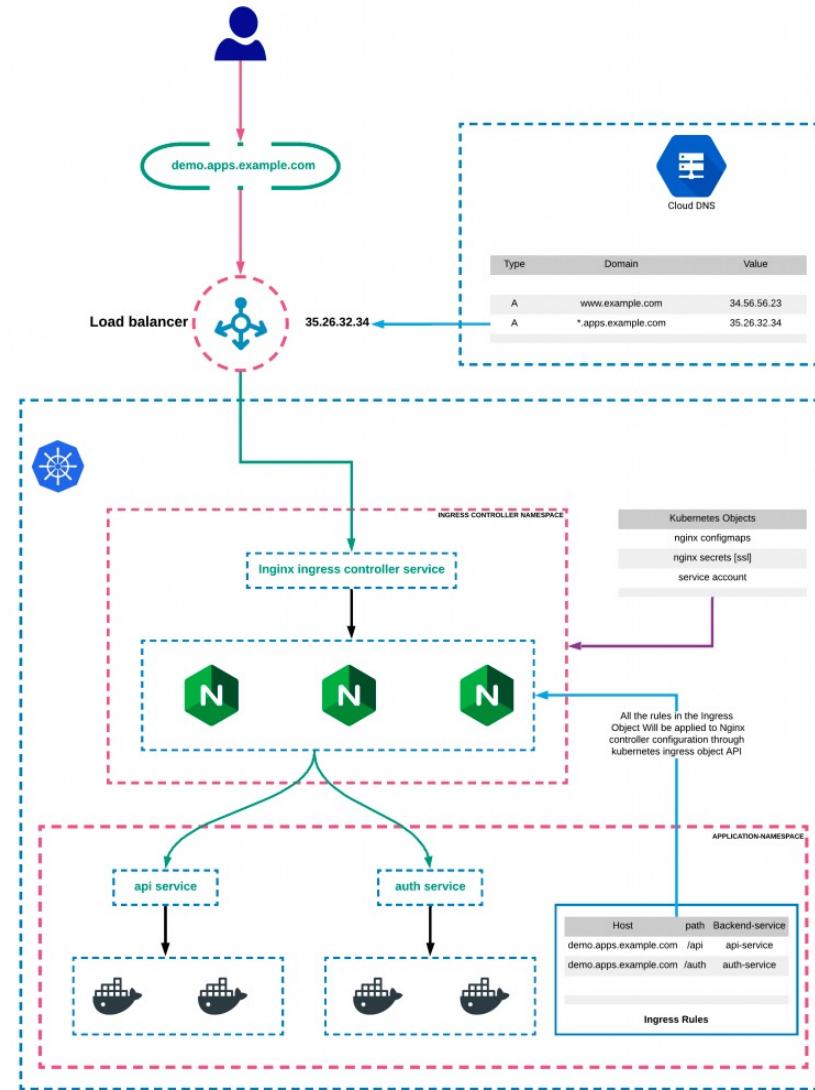
Traefik



Ingress process (nginx)

- nginx.conf is a Go template which can talk to Kubernetes Ingress API and pull latest values for traffic routing in real time.
- Nginx controller talks to Kubernetes Ingress API to check if there is a routing rule created
- If it finds any matching Ingress rules they are applied to Nginx controller configuration nginx.conf inside the pod using the Go template
- PROTIP: nginx.conf available in pod /etc/nginx/nginx.conf with all rules.

Ingress diagram



ConfigMap



KEY VALUES

ConfigMap

ConfigMap

- Many applications require configuration via:
 - Config Files
 - Command-Line Arguments
 - Environment Variables
- These need to be decoupled from images to keep portable
- ConfigMap API provides mechanisms to inject containers with configuration data
- Store individual properties or entire config files/JSON blobs
- Key-Value Pairs

ConfigMap

- Not meant for sensitive information
- PODs or controllers can use ConfigMaps

1. Populate the value of environment variables
2. Set command-line arguments in a container
3. Populate config files in a volume

```
kind: ConfigMap
apiVersion: v1
metadata:
  creationTimestamp: 2016-02-18T19:14:38Z
  name: example-config
  namespace: default
data:
  example.property.1: hello
  example.property.2: world
  example.property.file: |-  
    property.1=value-1  
    property.2=value-2  
    property.3=value-3
```

ConfigMap from directory

- 2 files in docs/user-guide/configmap/kubectl
 - game.properties
 - ui.properties

game.properties

```
enemies=aliens
lives=3
enemies.cheat=true
enemies.cheat.level=noGoodRotten
secret.code.passphrase=UUDDLRLRBABAS
secret.code.allowed=true
secret.code.lives=30
```

ui.properties

```
color.good=purple
color.bad=yellow
allow.textmode=true
how.nice.to.look=fairlyNice
```

ConfigMap from directory

```
kubectl create configmap game-config --from-file=docs/user-guide/configmap/kubectl
```

```
kubectl describe configmaps game-config
```

```
Name: game-config
```

```
Namespace: default
```

```
Labels: <none>
```

```
Annotations: <none>
```

```
Data
```

```
====
```

```
game.properties: 121 bytes
```

```
ui.properties: 83 bytes
```

ConfigMap from directory

```
kubectl get configmaps game-config -o yaml
```

```
apiVersion: v1
data:
    game.properties: |-  
        enemies=aliens  
        lives=3  
        ...  
    ui.properties: |-  
        color.good=purple  
        ...  
kind: ConfigMap
metadata:
    creationTimestamp: 2016-02-18T18:34:05Z
    name: game-config
    namespace: default
    resourceVersion: "407"-  
    selfLink: /api/v1/namespaces/default/configmaps/game-config
    uid: 30944725-d66e-11e5-8cd0-68f728db1985
```

ConfigMap from files

```
kubectl get configmaps \
game-config-2 \
-o yaml
```

```
kubectl create configmap \
game-config-2 \
--from-file=file1 \
--from-file=file2
```

```
apiVersion: v1
data:
  game.properties: |-  
    enemies=aliens  
    lives=3  
    ...  
  ui.properties: |-  
    color.good=purple  
    ...  
kind: ConfigMap
metadata:  
  creationTimestamp: 2016-02-18T18:52:05Z  
  name: game-config-2  
  namespace: default  
  resourceVersion: "516"-  
  selfLink: /api/v1/namespaces/default/configmaps/game-config-2  
  uid: b4952dc3-d670-11e5-8cd0-68f728db1985
```

ConfigMap options

--from-file=/path/to/directory

--from-file=/path/to/file1 (/path/to/file2)

Literal key=value: --from-literal=special.how=very

ConfigMap in PODs

- Populate Environment Variables

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: special-config
  namespace: default
data:
  special.how: very
  special.type: charm
```

OUTPUT

```
SPECIAL_LEVEL_KEY=very
SPECIAL_TYPE_KEY=charm
```

```
apiVersion: v1
kind: Pod
metadata:
  name: dapi-test-pod
spec:
  containers:
    - name: test-container
      image: busybox
      command: ["/bin/sh", "-c", "env"]
      env:
        - name: SPECIAL_LEVEL_KEY
          valueFrom:
            configMapKeyRef:
              name: special-config
              key: special.how
  restartPolicy: Never
```

ConfigMap Restrictions

- ConfigMaps must be created before they are consumed
- ConfigMaps can only be referenced by objects in the same namespace
- Quota for ConfigMap size not implemented yet
- Can only use ConfigMap for PODs created through API server.

Secrets



What secrets do applications have?

- Database credentials
- API credentials & endpoints (Twitter, Facebook etc.)
- Infrastructure API credentials (Google, Azure, AWS)
- Private keys (TLS, SSH)
- Many more!

Application Secrets

It is a bad idea to include these secrets in your code.

- Accidentally push up to GitHub with your code
- Push into your file storage and forget about
- Etc.

Application Secrets

There are bots crawling GitHub searching for secrets

Real life example:

Dev put keys out on GitHub, woke up next morning with a ton of emails and missed calls from Amazon

- 140 instances running under Dev's account.
- \$2,375 worth of Bitcoin mining

Create Secret

- Designed to hold all kinds of sensitive information
- Can be used by Pods (filesystem & environment variables) and the underlying kubelet when pulling images

```
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  password: mmyWfoidfluL==
  username: NyhdOKwB
```

Pod Secret

```
kind: Pod
metadata:
  name: secret-env-pod
spec:
  containers:
    - name: mycontainer
      image: redis
      env:
        - name: SECRET_USERNAME
          valueFrom:
            secretKeyRef:
              name: mysecret
              key: username
```

Volume Secret

```
spec:  
  containers  
    - name: mycontainer  
      image: redis  
      volumeMounts:  
        - name: "secrets"  
          mountPath: "/etc/my-secrets"  
          readOnly: true  
      volumes:  
        - name: "secrets"  
          secret:  
            secretName: "mysecret"
```

Labs: ConfigMap & Secrets



Helm: Package management



Helm sample application

The screenshot shows a web browser window titled "Guestbook" at the URL "dev.frontend.minikube.local/". The page displays a guestbook application with the title "Guestbook : 'MyPopRock Festival 2.0'" and the Globomantics logo.

The main content area shows a table of messages:

Name	Message
John	Very nice show !!

Below the table, there are pagination controls: "Items per page: 5" with a dropdown arrow, "1 - 1 of 1", and navigation arrows. A "Filter" input field is also present above the table.

The bottom section of the page is titled "Leave your feedback!" and contains form fields:

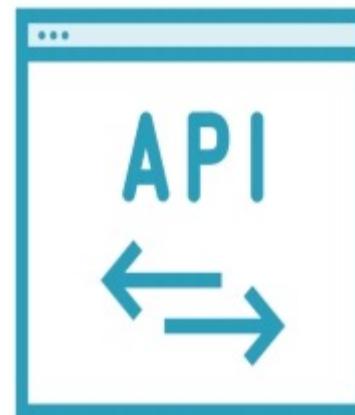
- "Your name *": Input field containing "Jane".
- "Your questbook message": Text area containing "Congrats to Globomantics DevOps !".
- A blue "Leave message" button.

Sample application components



Frontend

- ✓ ConfigMap
- ✓ Pod
- ✓ Service
- ✓ Ingress



Backend API

- ✓ Secret
- ✓ Pod
- ✓ Service



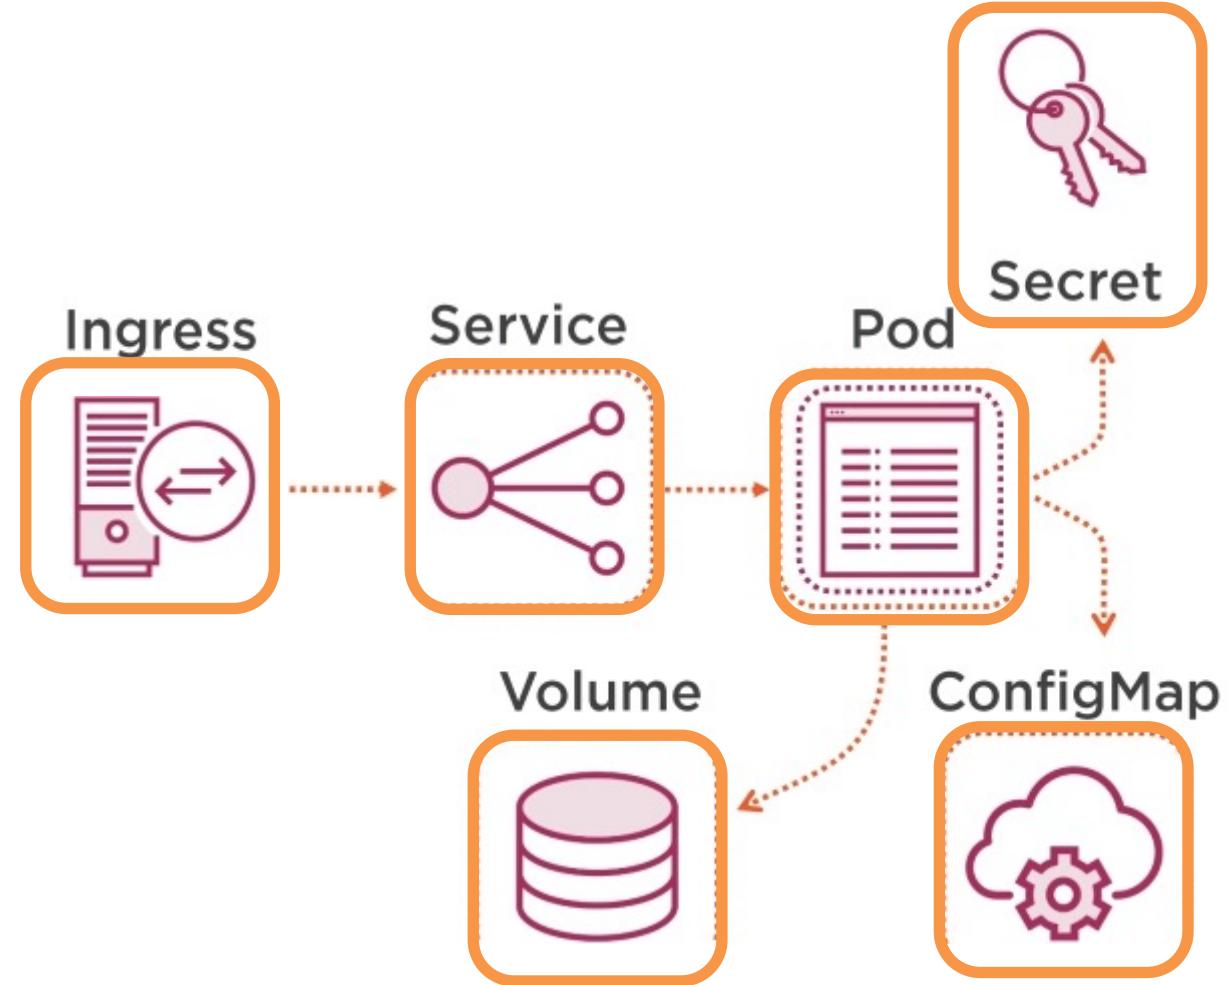
Database

- ✓ Secret
- ✓ PV
- ✓ PVC
- ✓ Pod
- ✓ Service



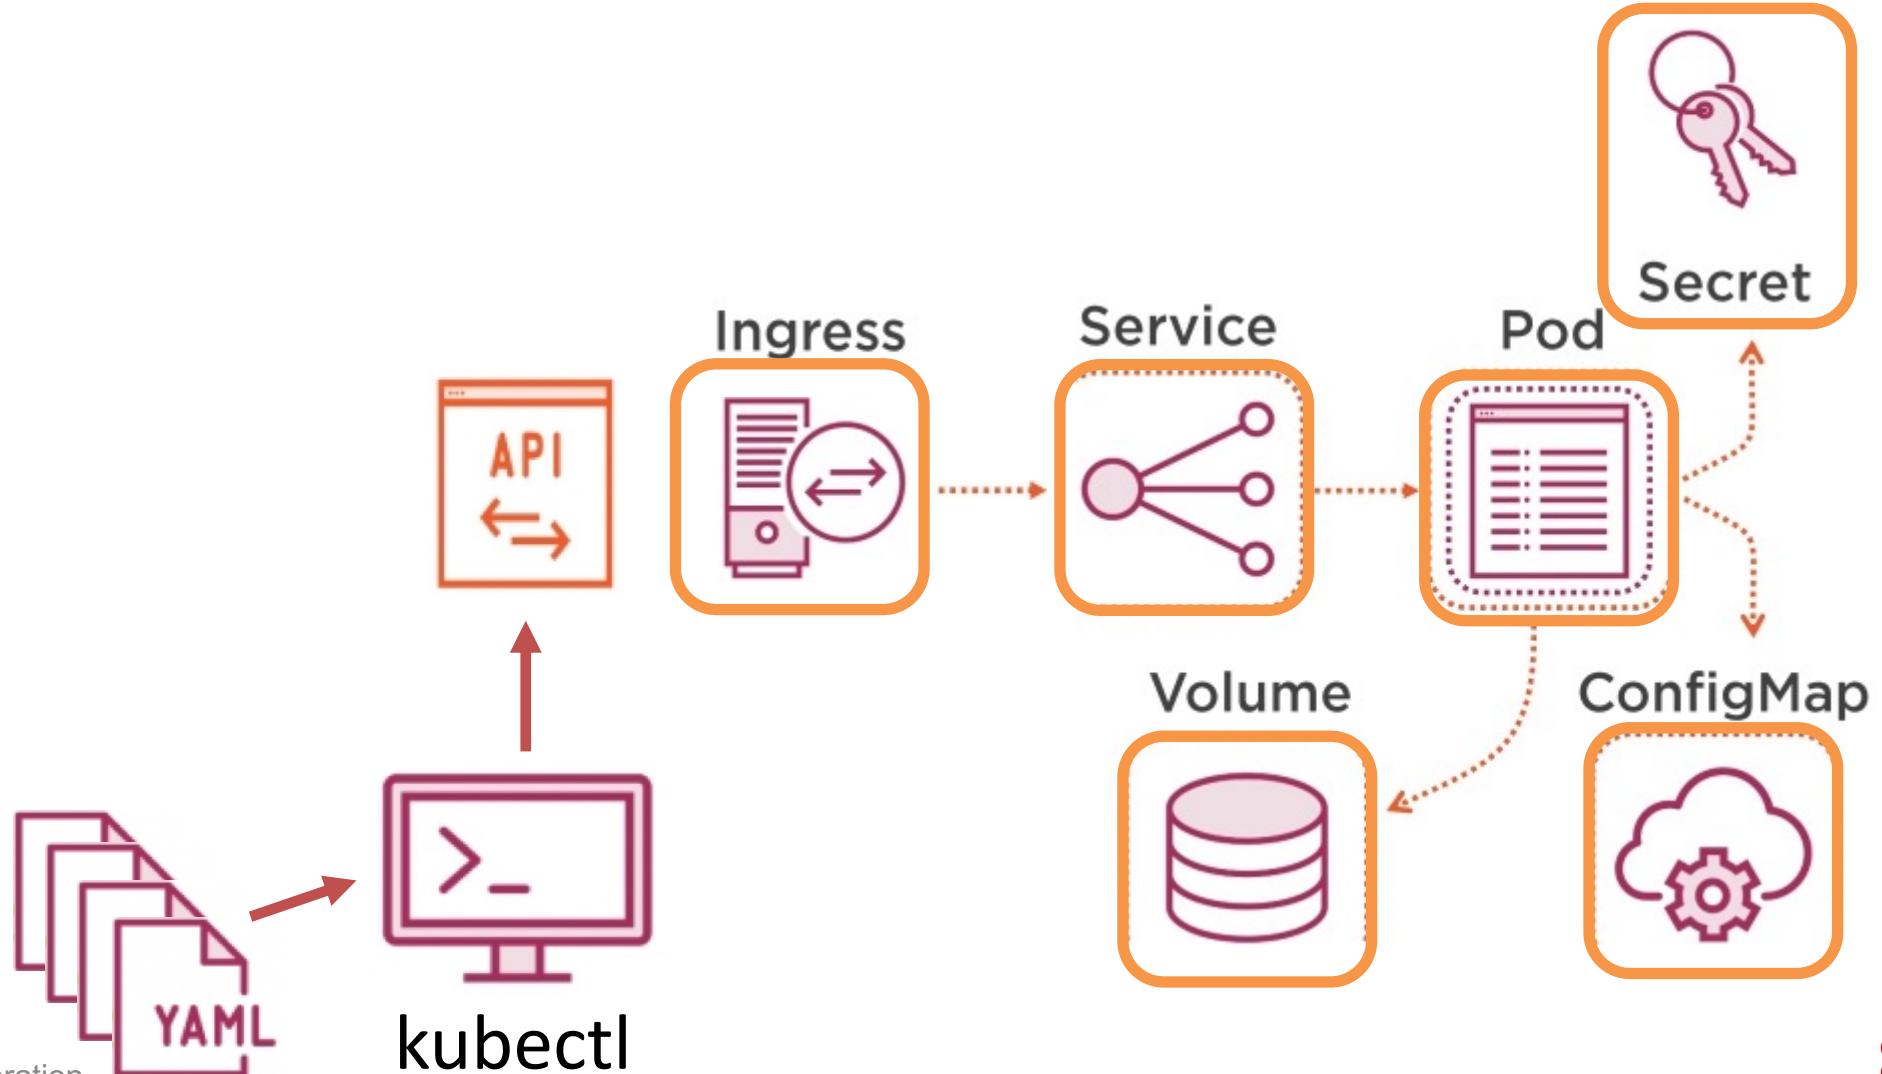
Native Kubernetes way

- Application
- Container
- Pod
- Service
- Ingress
- ConfigMap
- Secrets
- Volumes: PV, PVC, Storage



Native Kubernetes way

- Limitations
 - Packaging
 - Versioning



Guestbook: version 1

- ConfigMap
- Pod
- Service
- Ingress

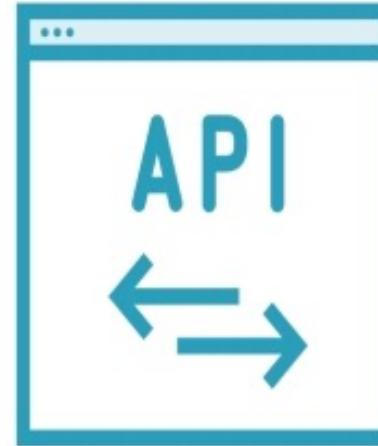


Frontend

Guestbook: version 2



Frontend



Backend



Database

- ConfigMap
- Pod
- Service
- Ingress

- Secret
- Pod
- Service

- Secret
- PV
- PVC
- Pod
- Service

Painful to manage

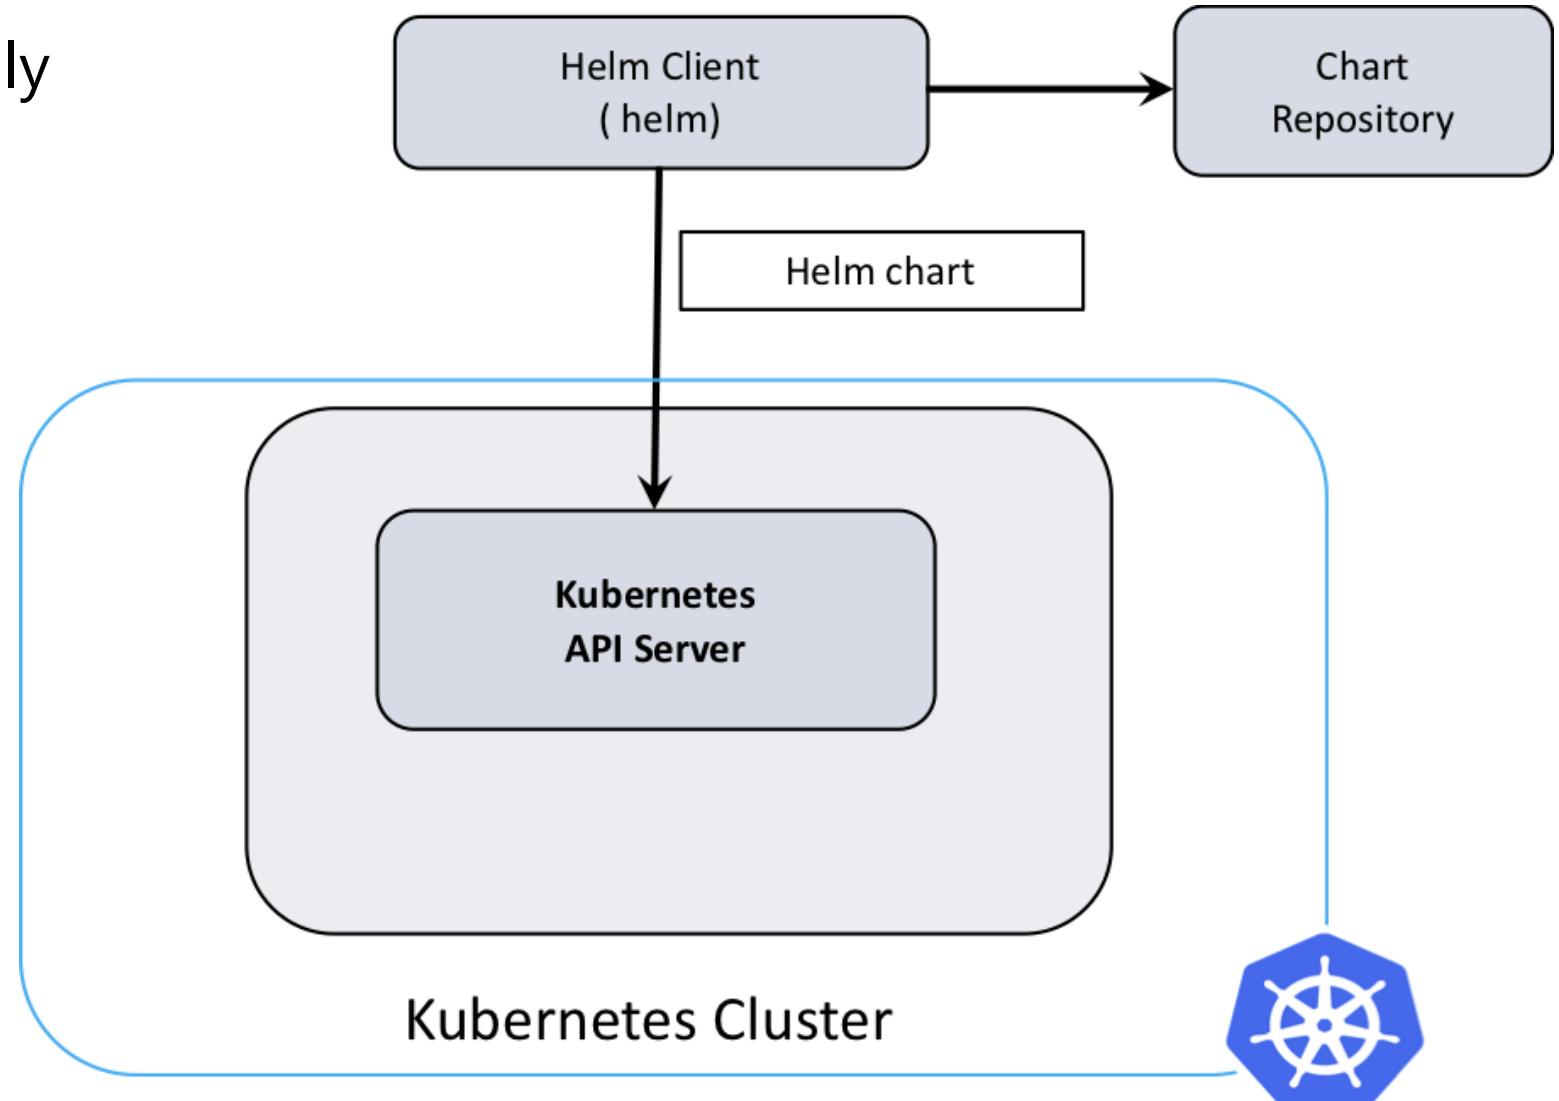


Helm features

	Package manager	Packages
System {	Apt Yum	deb rpm
Dev {	Maven Npm Pip	Jar, Ear, ... Node Modules Python packages
Kubernetes {	Helm	Charts

Helm architecture

- Helm:
 - develop charts locally
 - command-line



Helm features



Charts



Templates



Dependencies

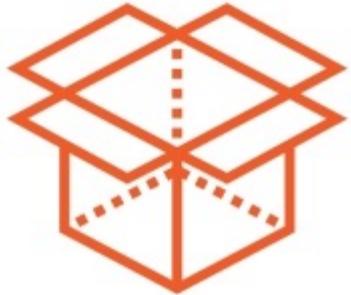


Repositories

Helm Chart Structure



Helm Chart structure



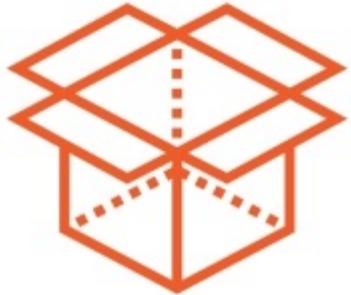
- nginx-demo
 - Chart.yaml
- Chart properties
 - name
 - version
 - more..

Helm Chart structure



- **nginx-demo**
 - Chart.yaml
 - README.md
- **Document chart**
 - Overrides
 - Maintainer
 - Instructions

Helm Chart structure



- nginx-demo
 - Chart.yaml
 - README.md
- templates
 - deployment.yaml
 - ingress.yaml
 - service.yaml
- Kubernetes object definitions
 - customizable YAML templates

Helm Chart structure



- `nginx-demo`
 - `Chart.yaml`
 - `README.md`
- `templates`
 - `deployment.yaml`
 - `ingress.yaml`
 - `service.yaml`
- `values.yaml`
- **Kubernetes object definitions**
 - customizable YAML templates
 - Provide default values.

Helm Chart structure



- nginx-demo
 - Chart.yaml
 - README.md
- templates
 - deployment.yaml
 - ingress.yaml
 - service.yaml
 - NOTES.txt
- values.yaml
- Provide helpful output after installation
 - How to access application
 - Create user/pass

Helm Chart structure



- nginx-demo
 - Chart.yaml
 - README.md
 - templates
 - deployment.yaml
 - ingress.yaml
 - service.yaml
 - NOTES.txt
 - tests
 - test-connection.yaml
 - values.yaml
- You can create tests to confirm Chart works as expected.

Helm Chart structure



- nginx-demo
 - Chart.yaml
 - README.md
 - requirements.yaml
 - templates
 - deployment.yaml
 - ingress.yaml
 - service.yaml
 - NOTES.txt
 - tests
 - test-connection.yaml
 - values.yaml
- Define sub-charts and dependencies

Helm Chart.yaml

Chart.yaml

```
apiVersion: v1
appVersion: "1.0"
description: A Helm chart for Kubernetes
name: nginx-demo
version: 0.1.0
```



Major, Minor, Patch (SemVer 2.0)

- App you are installing

Helm Chart.yaml

Chart.yaml

```
apiVersion: v1
appVersion: "1.0"
description: A Helm chart for Kubernetes
name: nginx-demo
version: 0.1.0
```



Major, Minor, Patch (SemVer 2.0)

- Version of Helm chart

Lab: Helm

